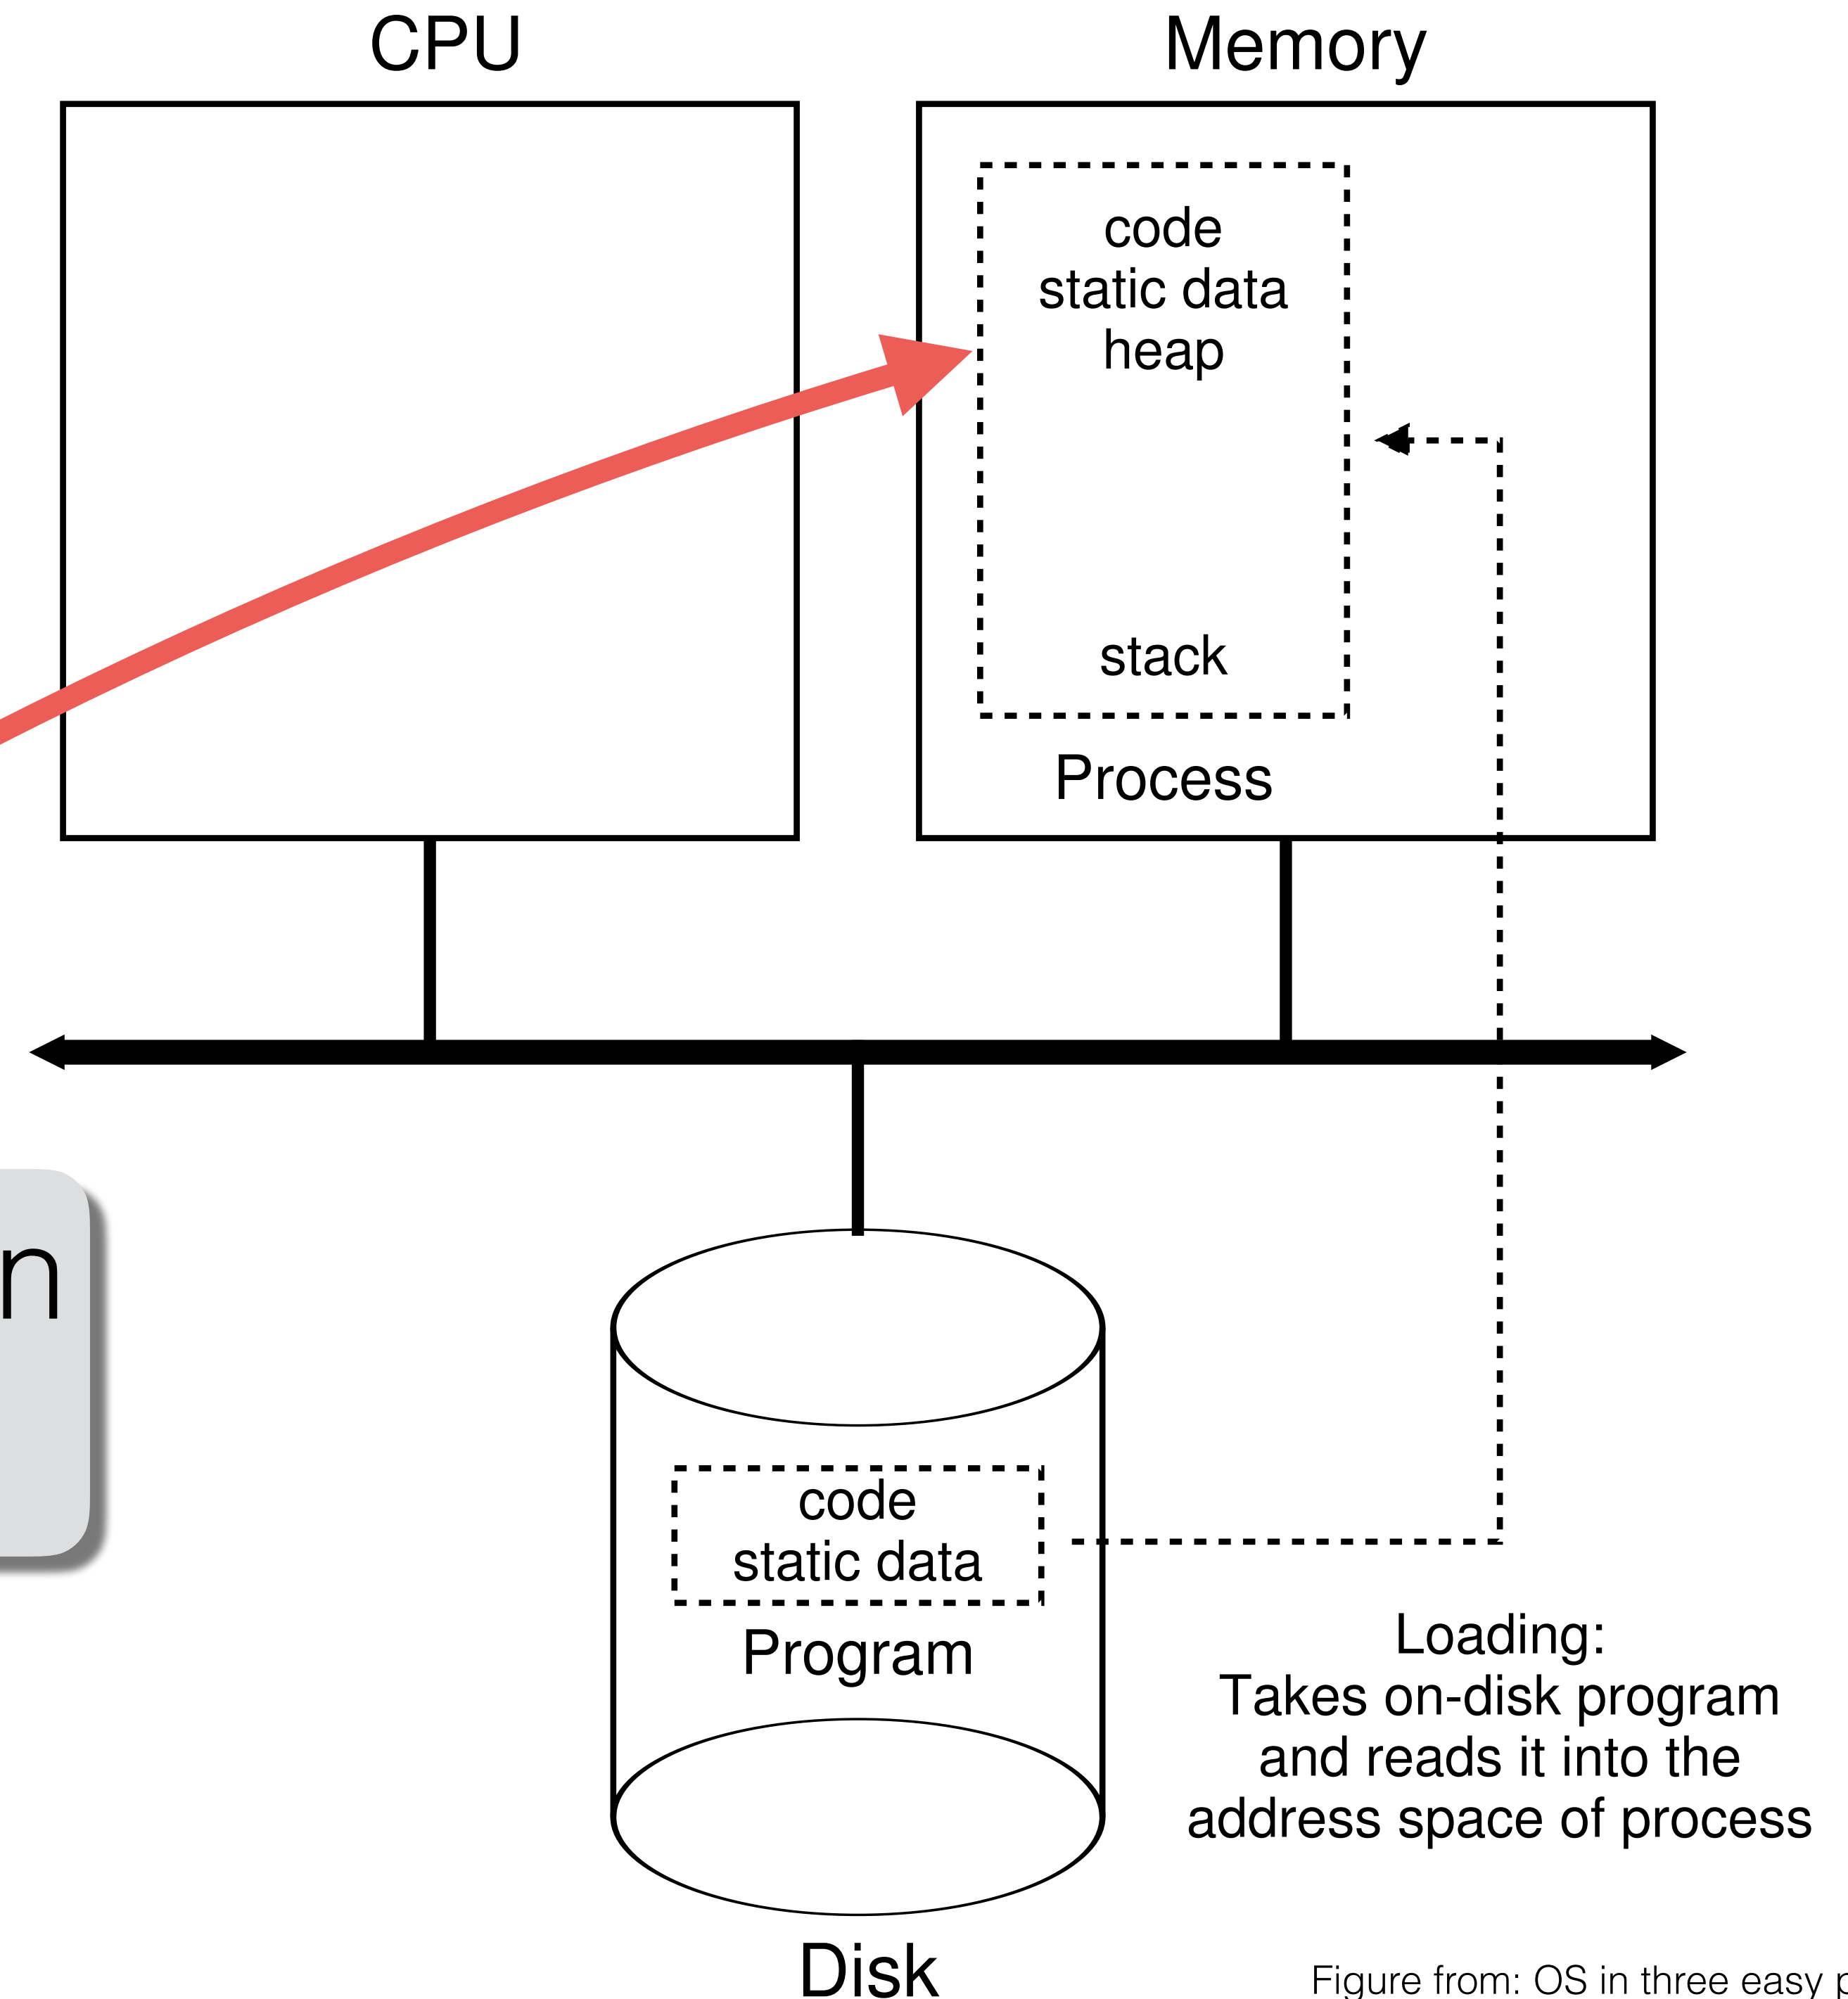


CSE4001: Operating Systems Concepts

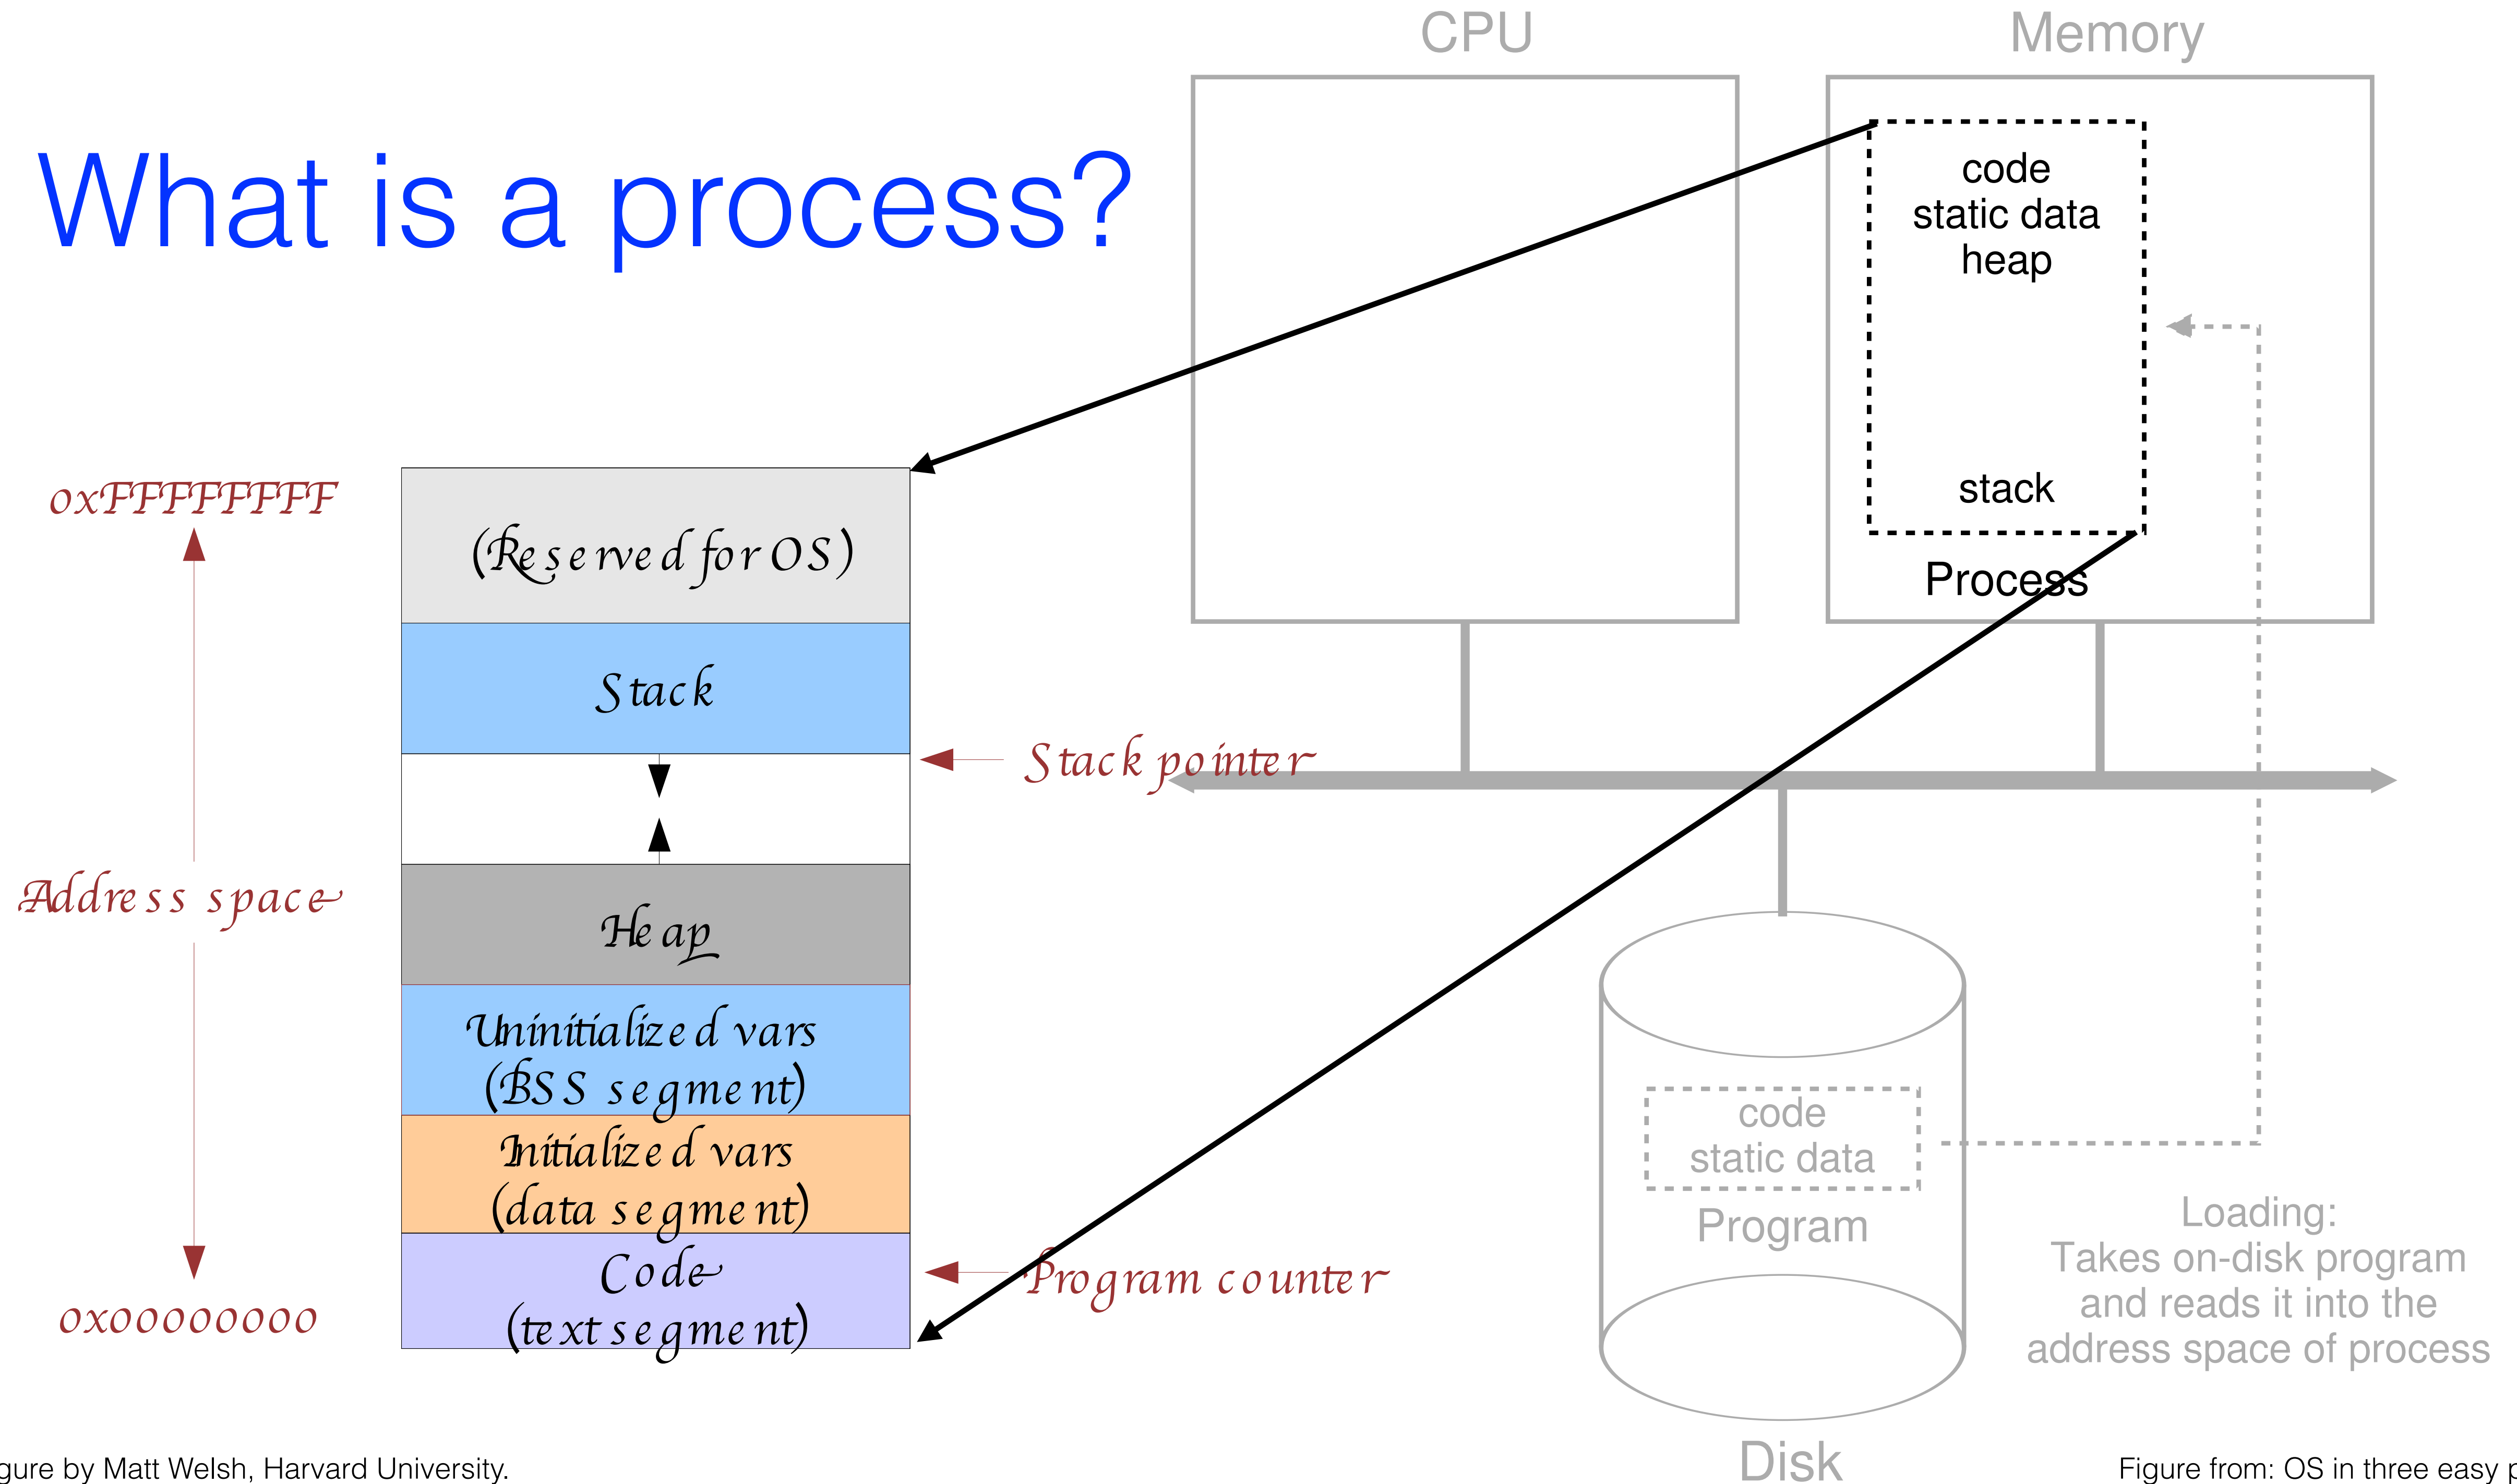
Processes and limited direct execution

What is a process?

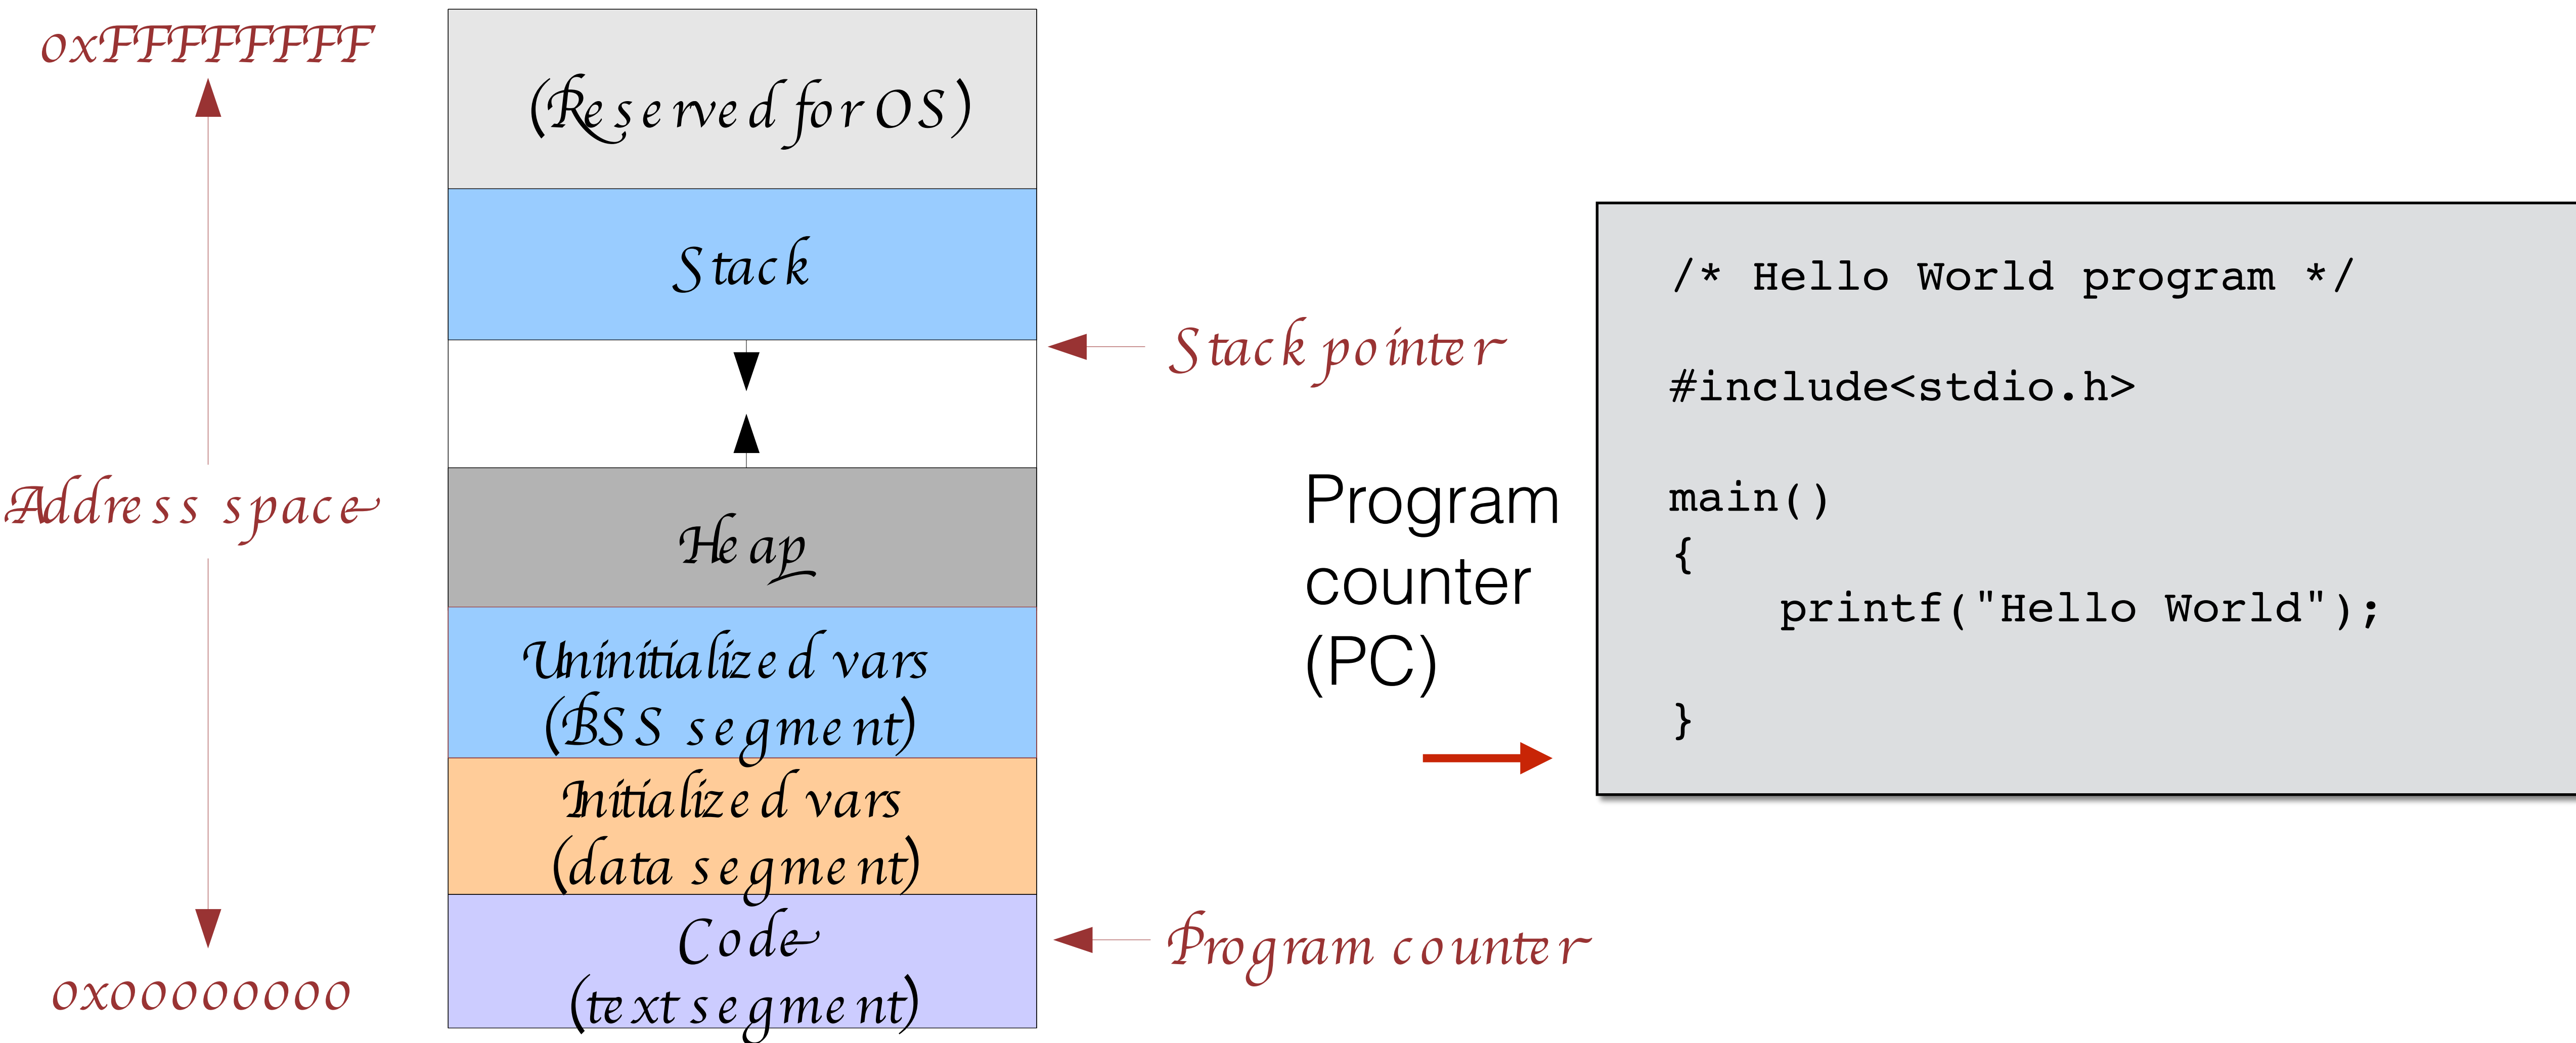
A process is an abstraction
of *a program in execution*.



What is a process?



What is a process?



The OS view of a process

- Process state (ready, running, blocked, ...)
- The **address space** (how many possible addresses)
- The **code** of the running program
- The **data** of the running program
- An execution **stack** encapsulating the state of procedure calls
- The **program counter (PC)** indicating the address of the next instruction.
- A set of general-purpose **registers** with current values
- A set of operating system **resources**
 - ◆ open files, network connections, signals, etc.
- CPU scheduling info: process **priority**
- Each process is identified by its **process ID (PID)**

All these information is stored in a construct called
Process Control Block (PCB)

Process Control Block (PCB)

The OS maintains a PCB for each process. It is a data structure with many fields.

Defined in:

`/include/linux/sched.h`

```
struct task_struct {
    volatile long state; Execution state
    unsigned long flags;
    int sigpending;
    mm_segment_t addr_limit;
    struct exec_domain *exec_domain;
    volatile long need_resched;
    unsigned long ptrace;
    int lock_depth;
    unsigned int cpu;
    int prio, static_prio;
    struct list_head run_list;
    prio_array_t *array;
    unsigned long sleep_avg;
    unsigned long last_run;
    unsigned long policy;
    unsigned long cpus_allowed;
    unsigned int time_slice, first_time_slice;
    atomic_t usage;
    struct list_head tasks;
    struct list_head ptrace_children;
    struct list_head ptrace_list;
    struct mm_struct *mm, *active_mm; Memory mgmt info
    struct linux_binfmt *binfmt;
    int exit_code, exit_signal;
    int pdeath_signal;
    unsigned long personality;
    int did_exec:1;
    unsigned task_dumpable:1;
    pid_t pid; Process ID
    pid_t pgrp;
    pid_t tty_old_pgrp;
    pid_t session;
    pid_t tgid;
    int leader;
    struct task_struct *real_parent;
    struct task_struct *parent;
    struct list_head children;
    struct list_head sibling;
    struct task_struct *group_leader;
    struct pid_link pids[PIDTYPE_MAX];
    wait_queue_head_t wait_chldexit;
    struct completion *vfork_done;
    int *set_child_tid;
    int *clear_child_tid;
    unsigned long rt_priority; Priority
```

```
    unsigned long it_real_value, it_prof_value, it_virt_value;
    unsigned long it_real_incr, it_prof_incr, it_virt_incr;
    struct timer_list real_timer;
    struct tms times;
    struct tms group_times; Accounting info
    unsigned long start_time;
    long per_cpu_utime[NR_CPUS], per_cpu_stime[NR_CPUS];
    unsigned long min_flt, maj_flt, nswap, cmin_flt, cmaj_flt,
    cnsnap;
    int swappable:1;
    uid_t uid, euid, suid, fsuid; User ID
    gid_t gid, egid, sgid, fsgid;
    int ngroups;
    gid_t groups[NGROUPS];
    kernel_cap_t cap_effective, cap_inheritable, cap_permitted;
    int keep_capabilities:1;
    struct user_struct *user;
    struct rlimit rlim[RLIM_NLIMITS];
    unsigned short used_math;
    char comm[16];
    int link_count, total_link_count;
    struct tty_struct *tty;
    unsigned int locks;
    struct sem_undo *semundo;
    struct sem_queue *semsleeping;
    struct thread_struct thread; CPU state
    struct fs_struct *fs;
    struct files_struct *files; Open files
    struct namespace *namespace;
    struct signal_struct *signal;
    struct sighand_struct *sighand;
    sigset_t blocked, real_blocked;
    struct sigpending pending;
    unsigned long sas_ss_sp;
    size_t sas_ss_size;
    int (*notifier)(void *priv);
    void *notifier_data;
    sigset_t *notifier_mask;
    void *tux_info;
    void (*tux_exit)(void);
    u32 parent_exec_id;
    u32 self_exec_id;
    spinlock_t alloc_lock;
    spinlock_t switch_lock;
    void *journal_info;
    unsigned long ptrace_message;
    siginfo_t *last_siginfo;
};
```

Example of simple PCB: The xv6 proc structure

```
// the registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    int eip;
    int esp;
    int ebx;
    int ecx;
    int edx;
    int esi;
    int edi;
    int ebp;
};
```

[CK+08] “The xv6 Operating System”

Russ Cox, Frans Kaashoek, Robert Morris, Nickolai Zeldovich

From: <http://pdos.csail.mit.edu/6.828/2008/index.html>

The coolest real and little OS in the world. Download and play with it to learn more about the details of how operating systems actually work.

Example of simple PCB: The xv6 proc structure

```
// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };

// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;                // Start of process memory
    uint sz;                  // Size of process memory
    char *kstack;             // Bottom of kernel stack
                              // for this process
    enum proc_state state;    // Process state
    int pid;                  // Process ID
    struct proc *parent;      // Parent process
    void *chan;               // If non-zero, sleeping on chan
    int killed;               // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;         // Current directory
    struct context context;    // Switch here to run process
    struct trapframe *tf;     // Trap frame for the
                              // current interrupt
};
```


PCB in OS161

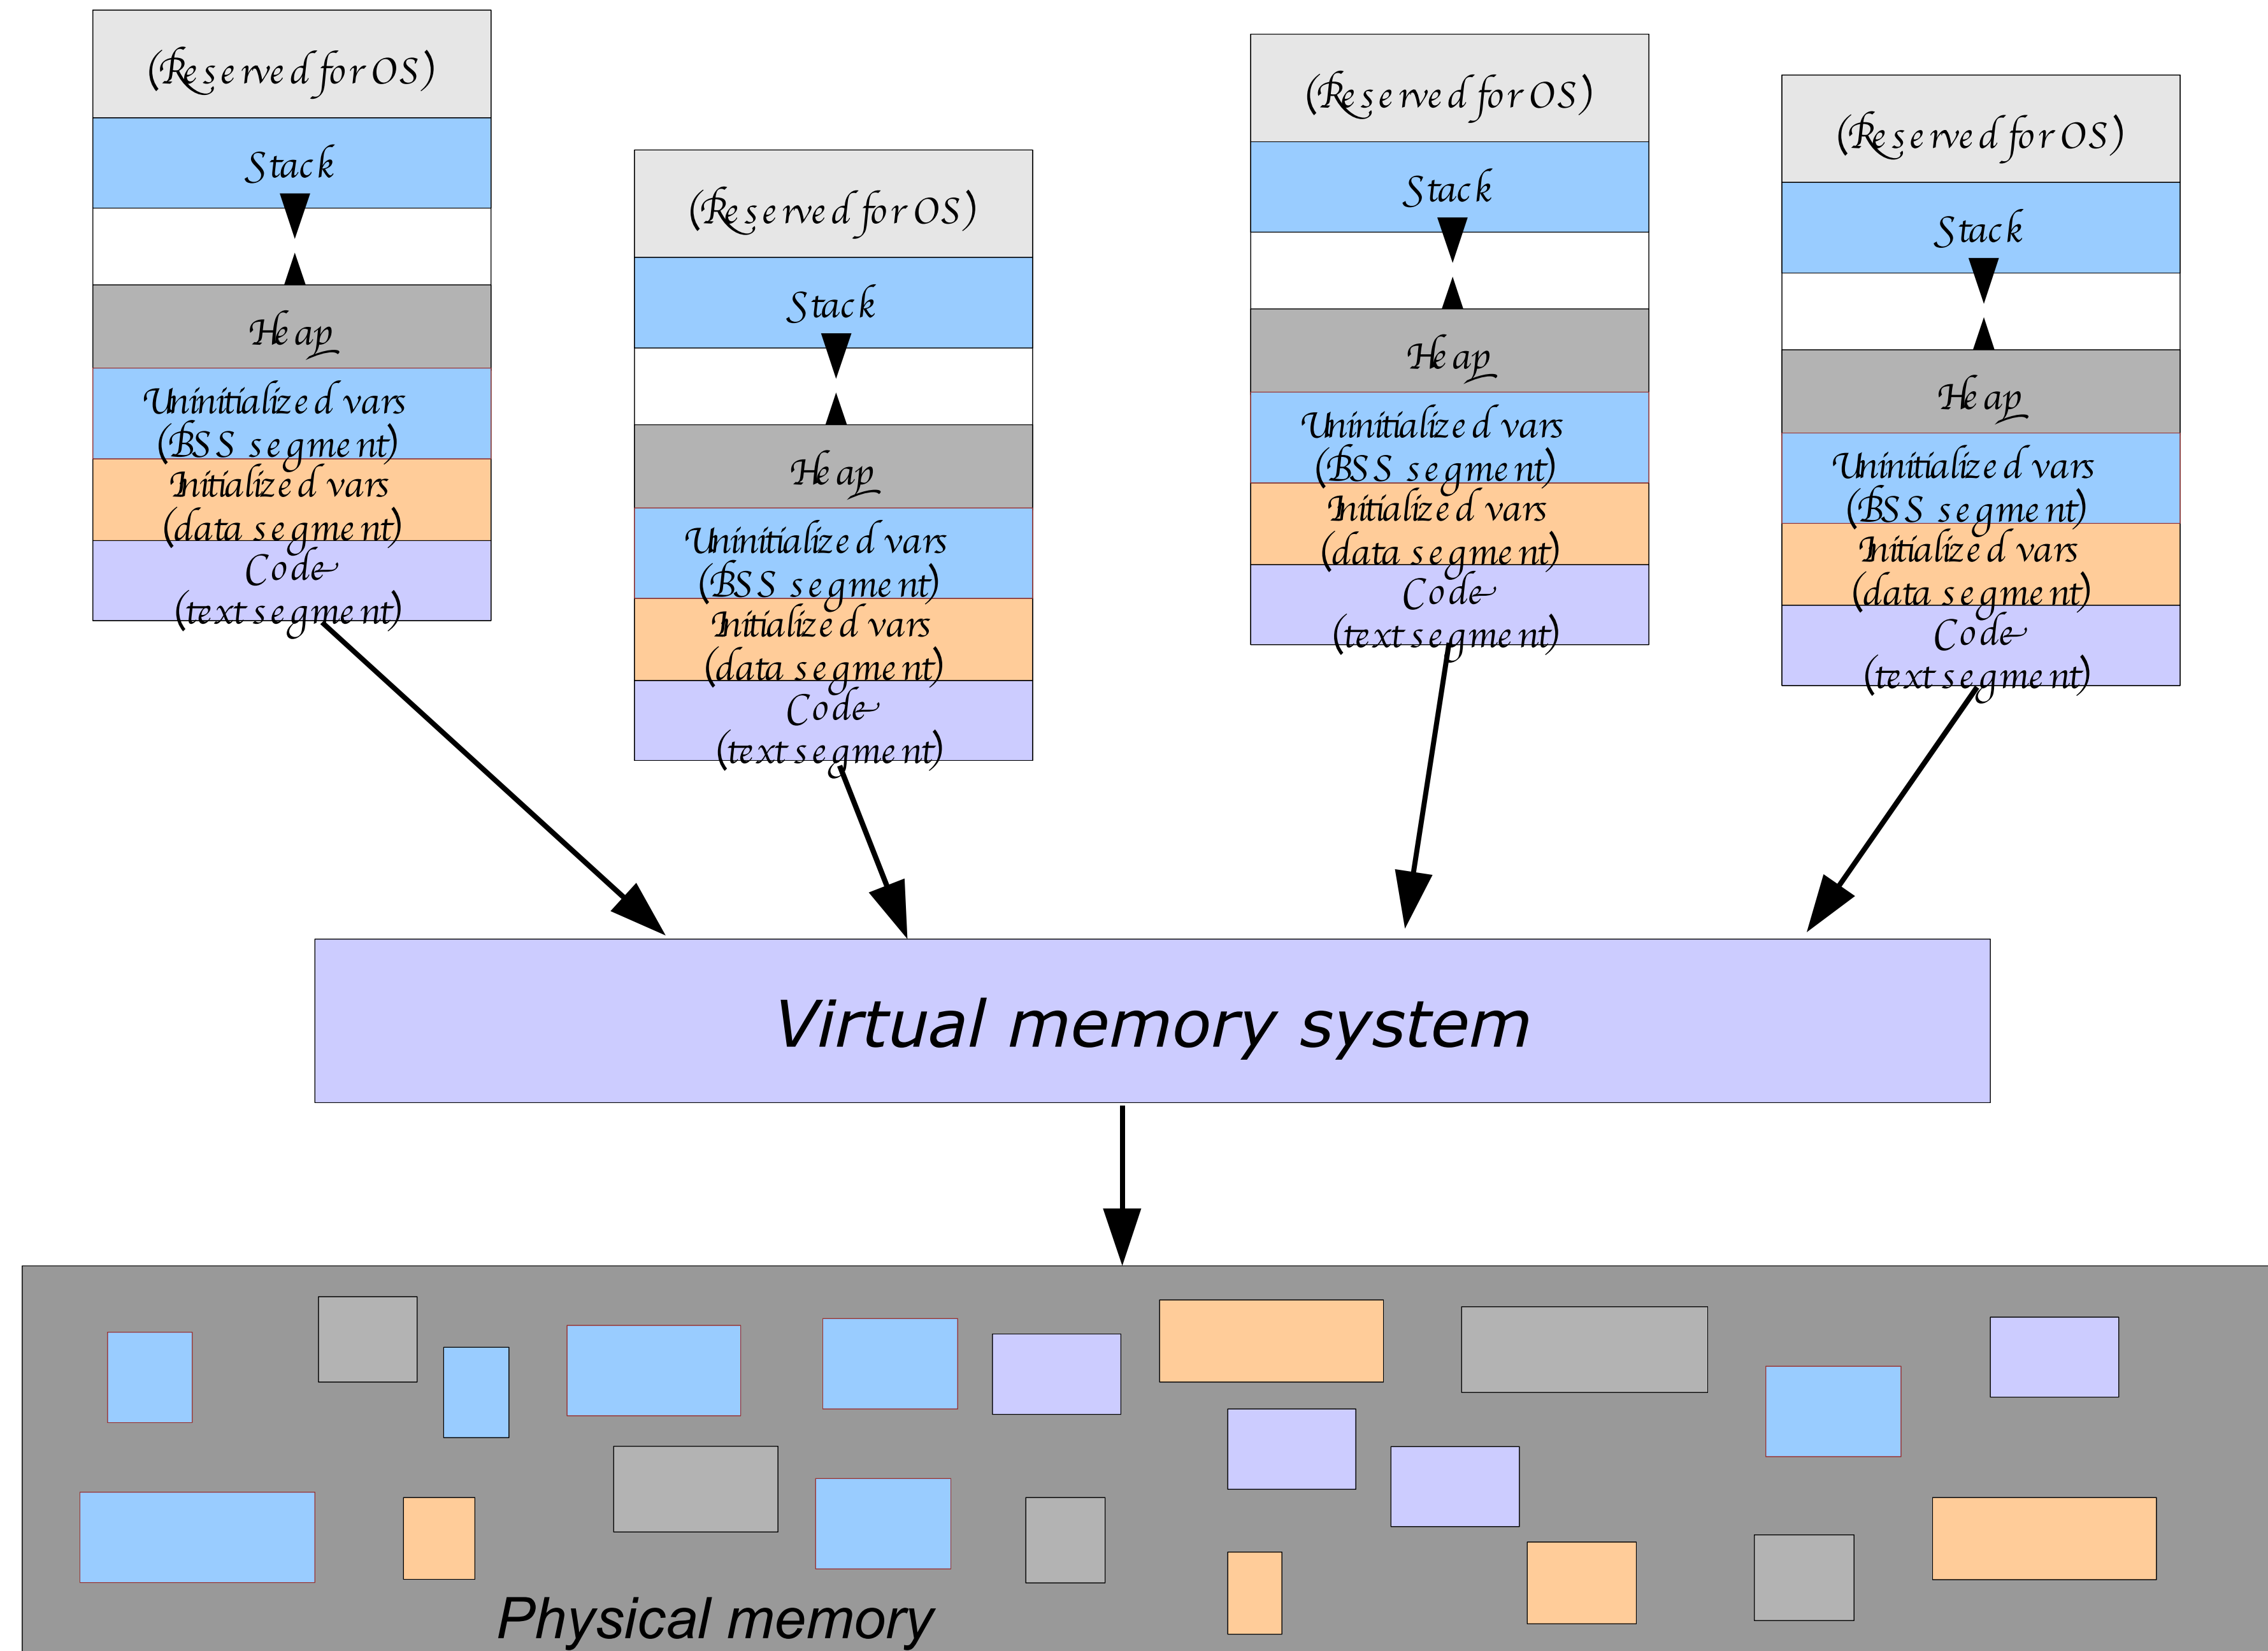
kern/src/kern/
include/thread.h

```
60  /* States a thread can be in. */
61  typedef enum {
62      S_RUN,          /* running */
63      S_READY,        /* ready to run */
64      S_SLEEP,        /* sleeping */
65      S_ZOMBIE,       /* zombie; exited but not yet deleted */
66  } threadstate_t;
67
68  /* Thread structure. */
69  struct thread {
70      /*
71       * These go up front so they're easy to get to even if the
72       * debugger is messed up.
73       */
74      char *t_name;          /* Name of this thread */
75      const char *t_wchan_name; /* Name of wait channel, if sleeping */
76      threadstate_t t_state; /* State this thread is in */
77
78      /*
79       * Thread subsystem internal fields.
80       */
81      struct thread_machdep t_machdep; /* Any machine-dependent goo */
82      struct threadlistnode t_listnode; /* Link for run/sleep/zombie lists */
83      void *t_stack;                  /* Kernel-level stack */
84      struct switchframe *t_context; /* Saved register context (on stack) */
85      struct cpu *t_cpu;              /* CPU thread runs on */

```

Running multiple processes concurrently

- Thousands of processes may be running
- Users do not need to worry about CPU availability
- Modern operating systems use a technique called **time sharing**.



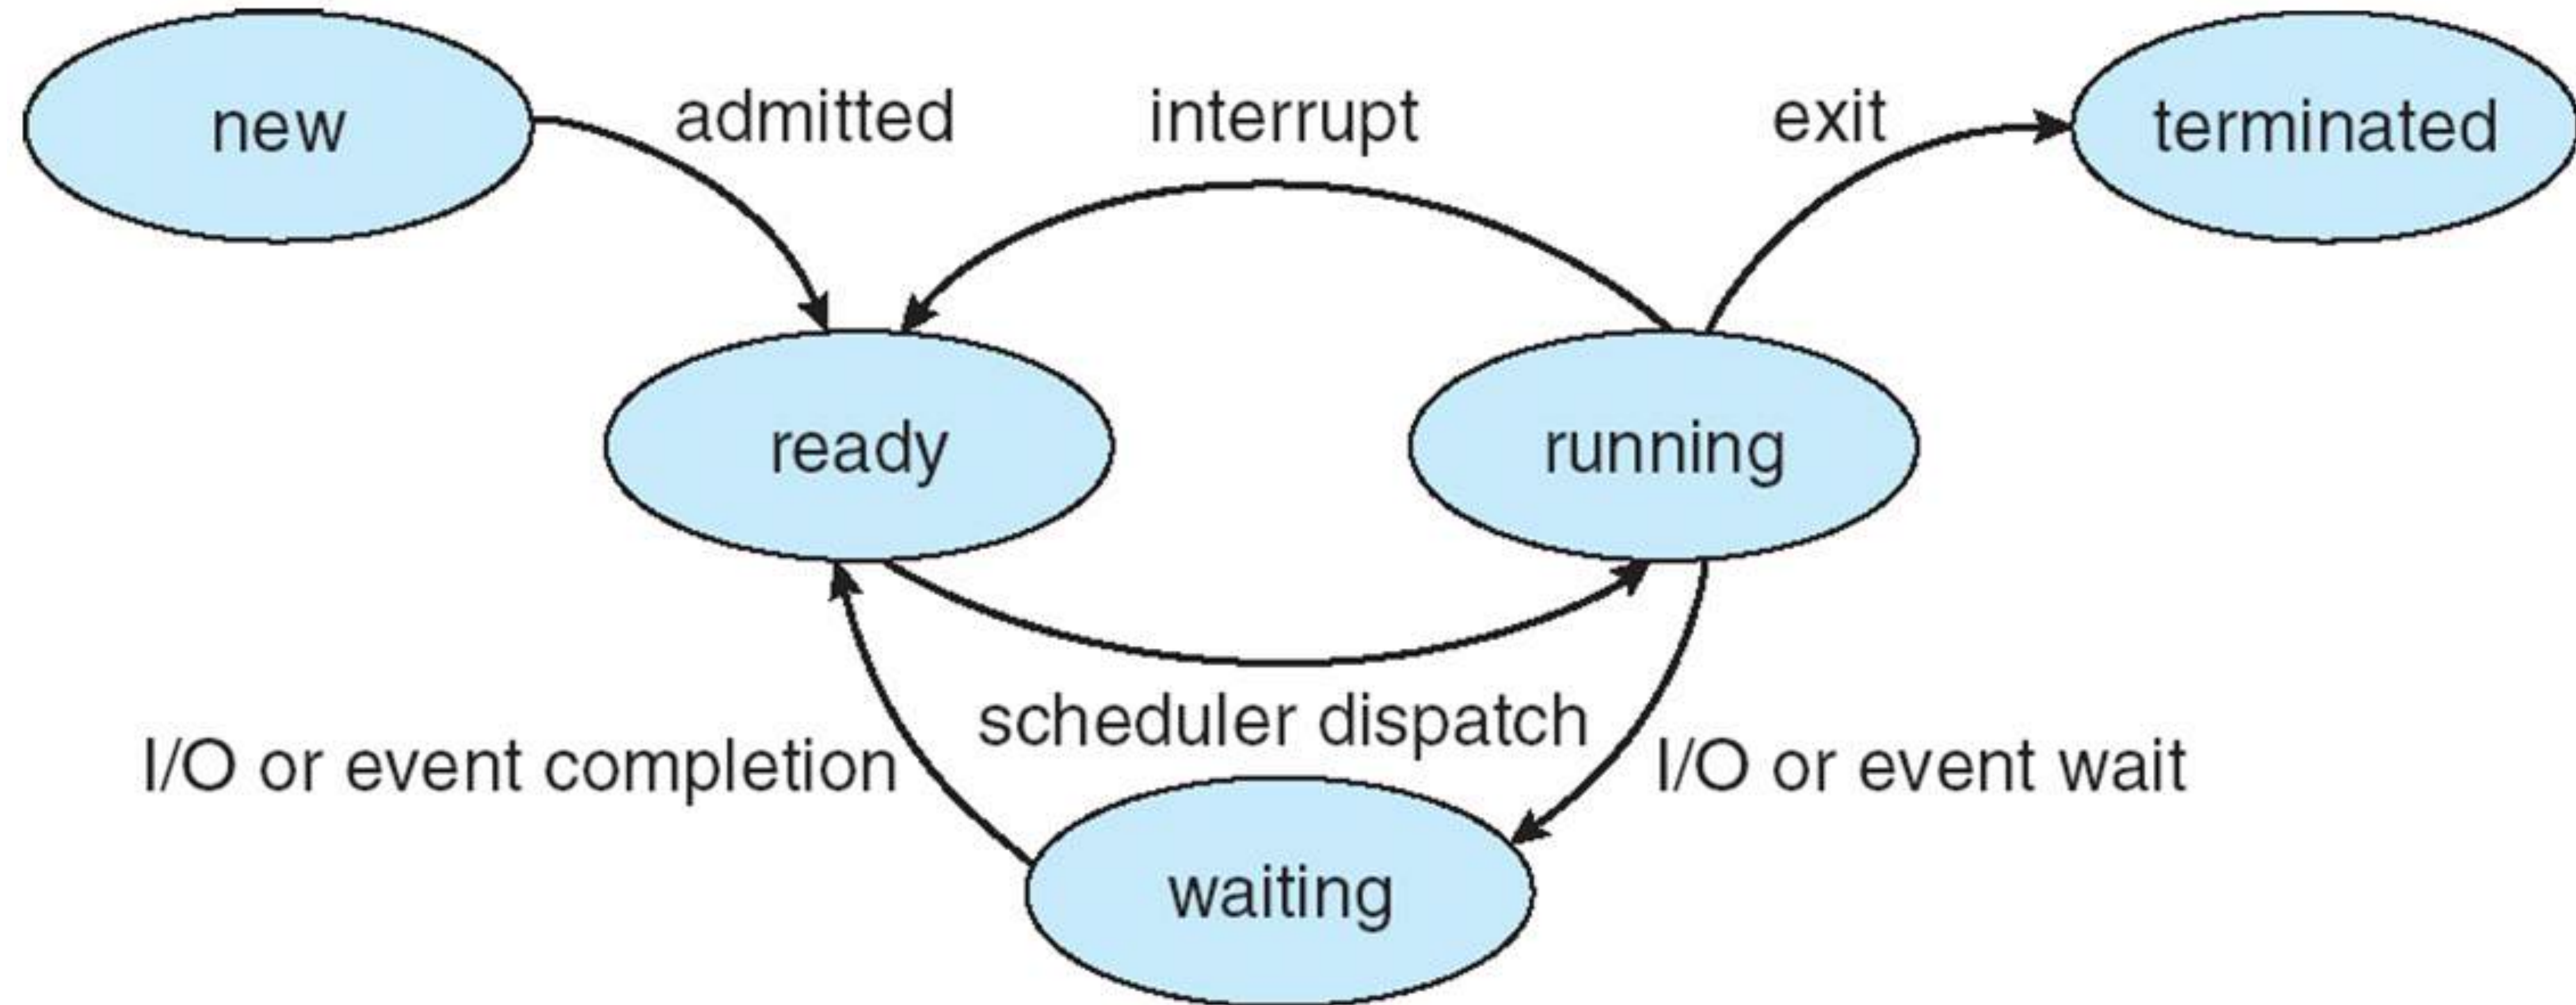
Process Example

A **process** is an instance of a program being executed

- Use “ps” to list processes on UNIX systems

PID	TTY	STAT	TIME	COMMAND
842	tty1	S	0:00	-bash
867	tty1	S	0:00	xinit
873	tty1	S	0:00	fvwm2
887	tty1	S	0:00	xload
888	tty1	S	0:02	/usr/local/j2sdk1.4.0/bin/java ApmView 896 243
1881	tty1	S	0:00	rxvt -fn fixed -cr red -fg white -bg #586570 -geometr
1883	pts/2	S	0:00	bash
1910	pts/0	S	0:00	/bin/sh /home/mdw/bin/ooffice arch.sxi
1911	pts/0	S	1:20	/usr/local/OpenOffice.org1.1.0/program/soffice.bin ar
1937	tty1	S	0:00	/bin/sh /home/mdw/bin/set-wlan-OFF
2310	pts/2	R	0:00	ps -Umdw -x

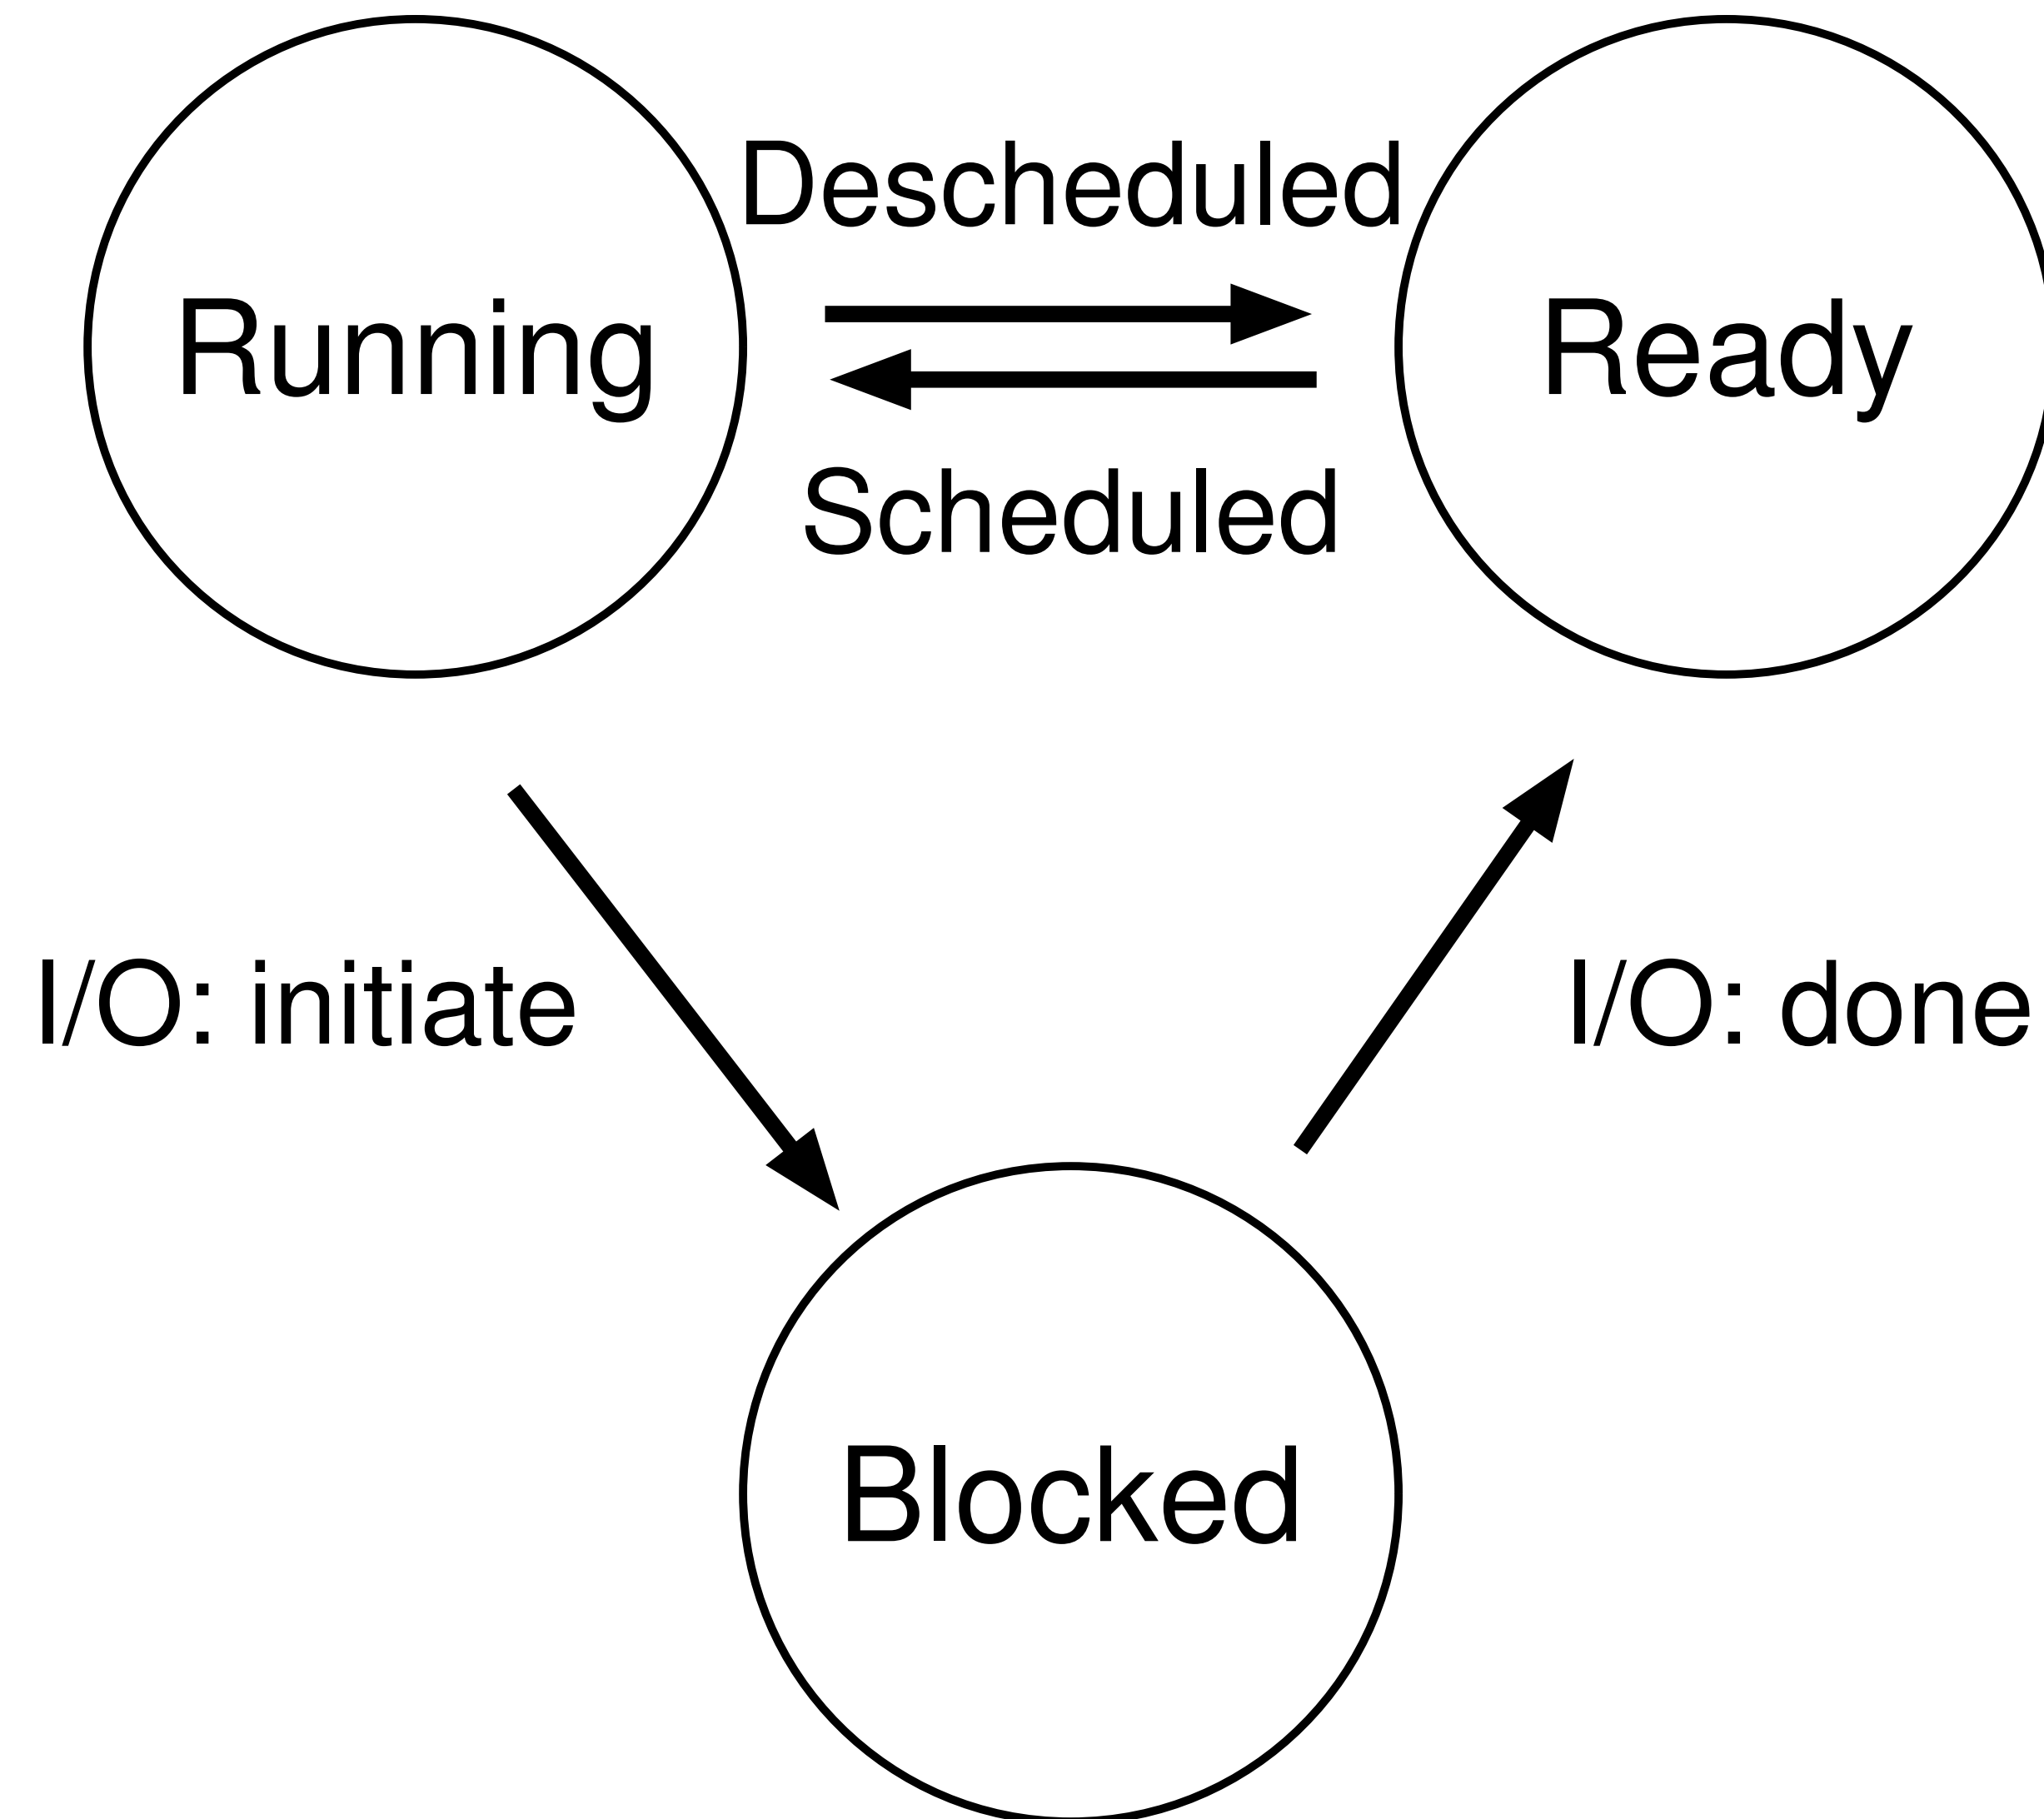
Life cycle of a process



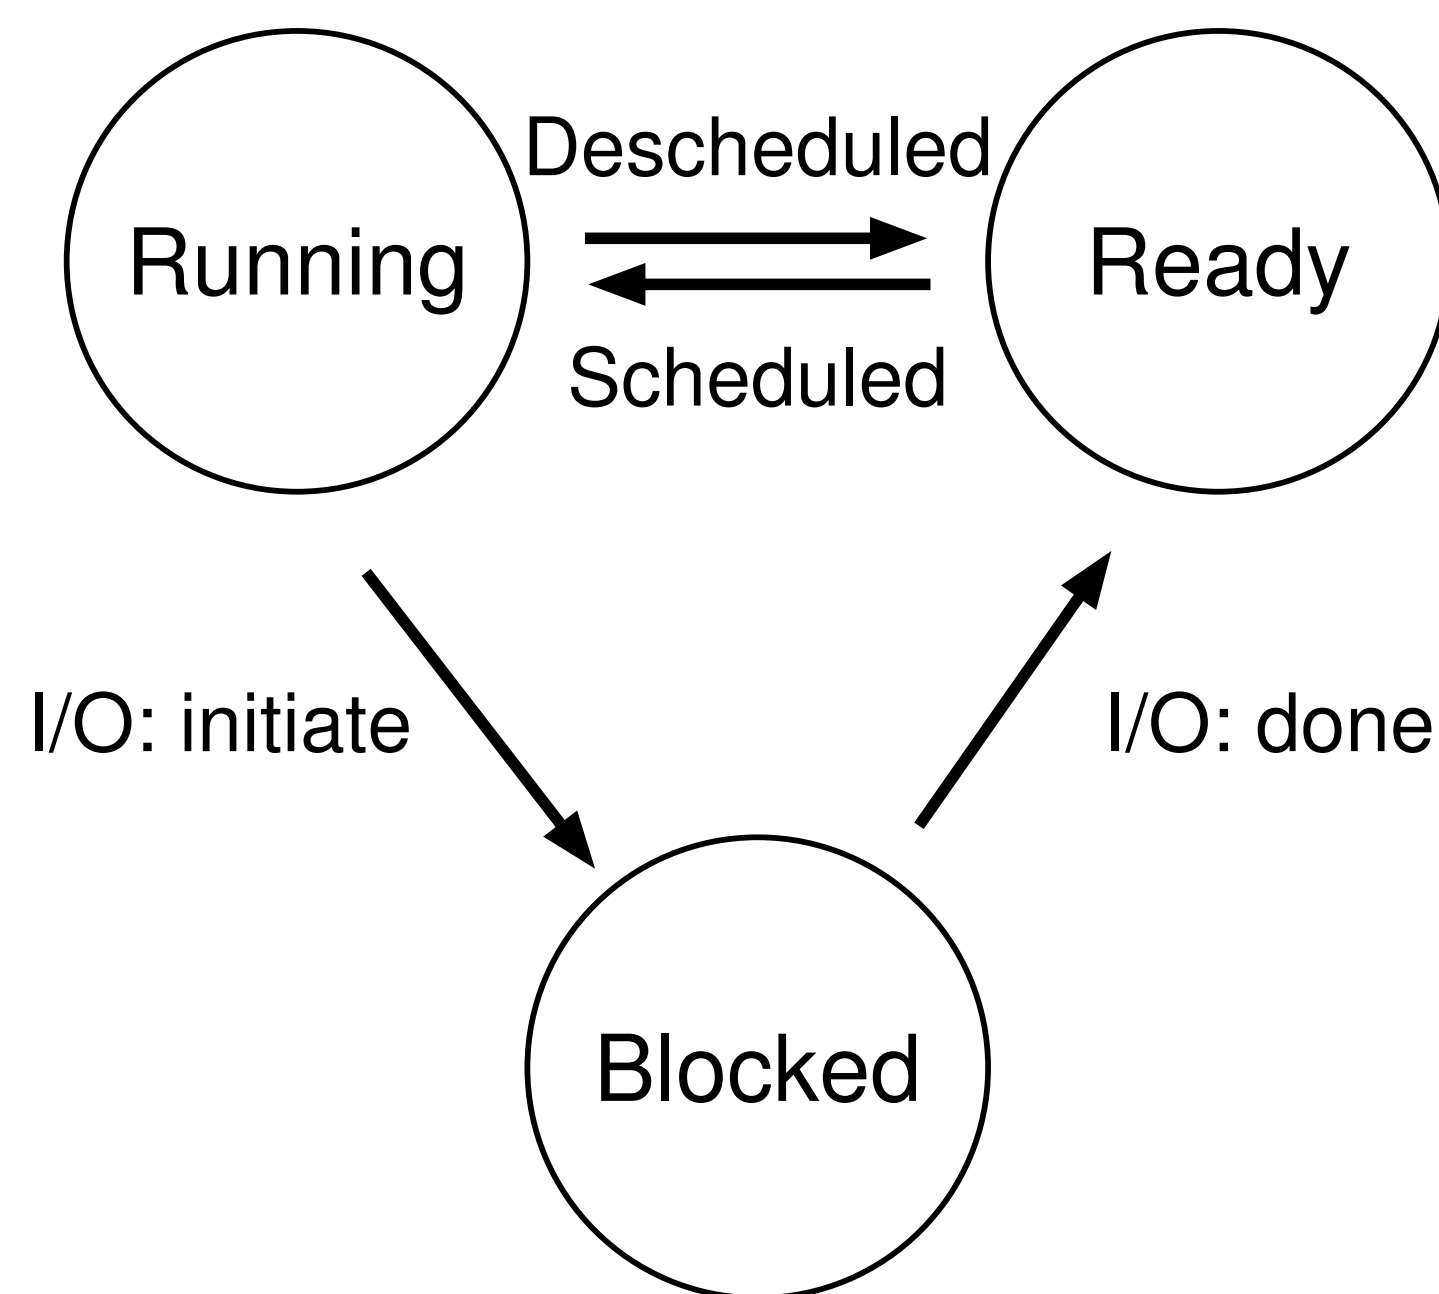
States of a process:

- **new**: The process is being created
- **running**: Instructions are being executed
- **waiting**: The process is waiting for some event to occur
- **ready**: The process is waiting to be assigned to a processor
- **terminated**: The process has finished execution

Life cycle of a process




Example: two running processes, no I/O



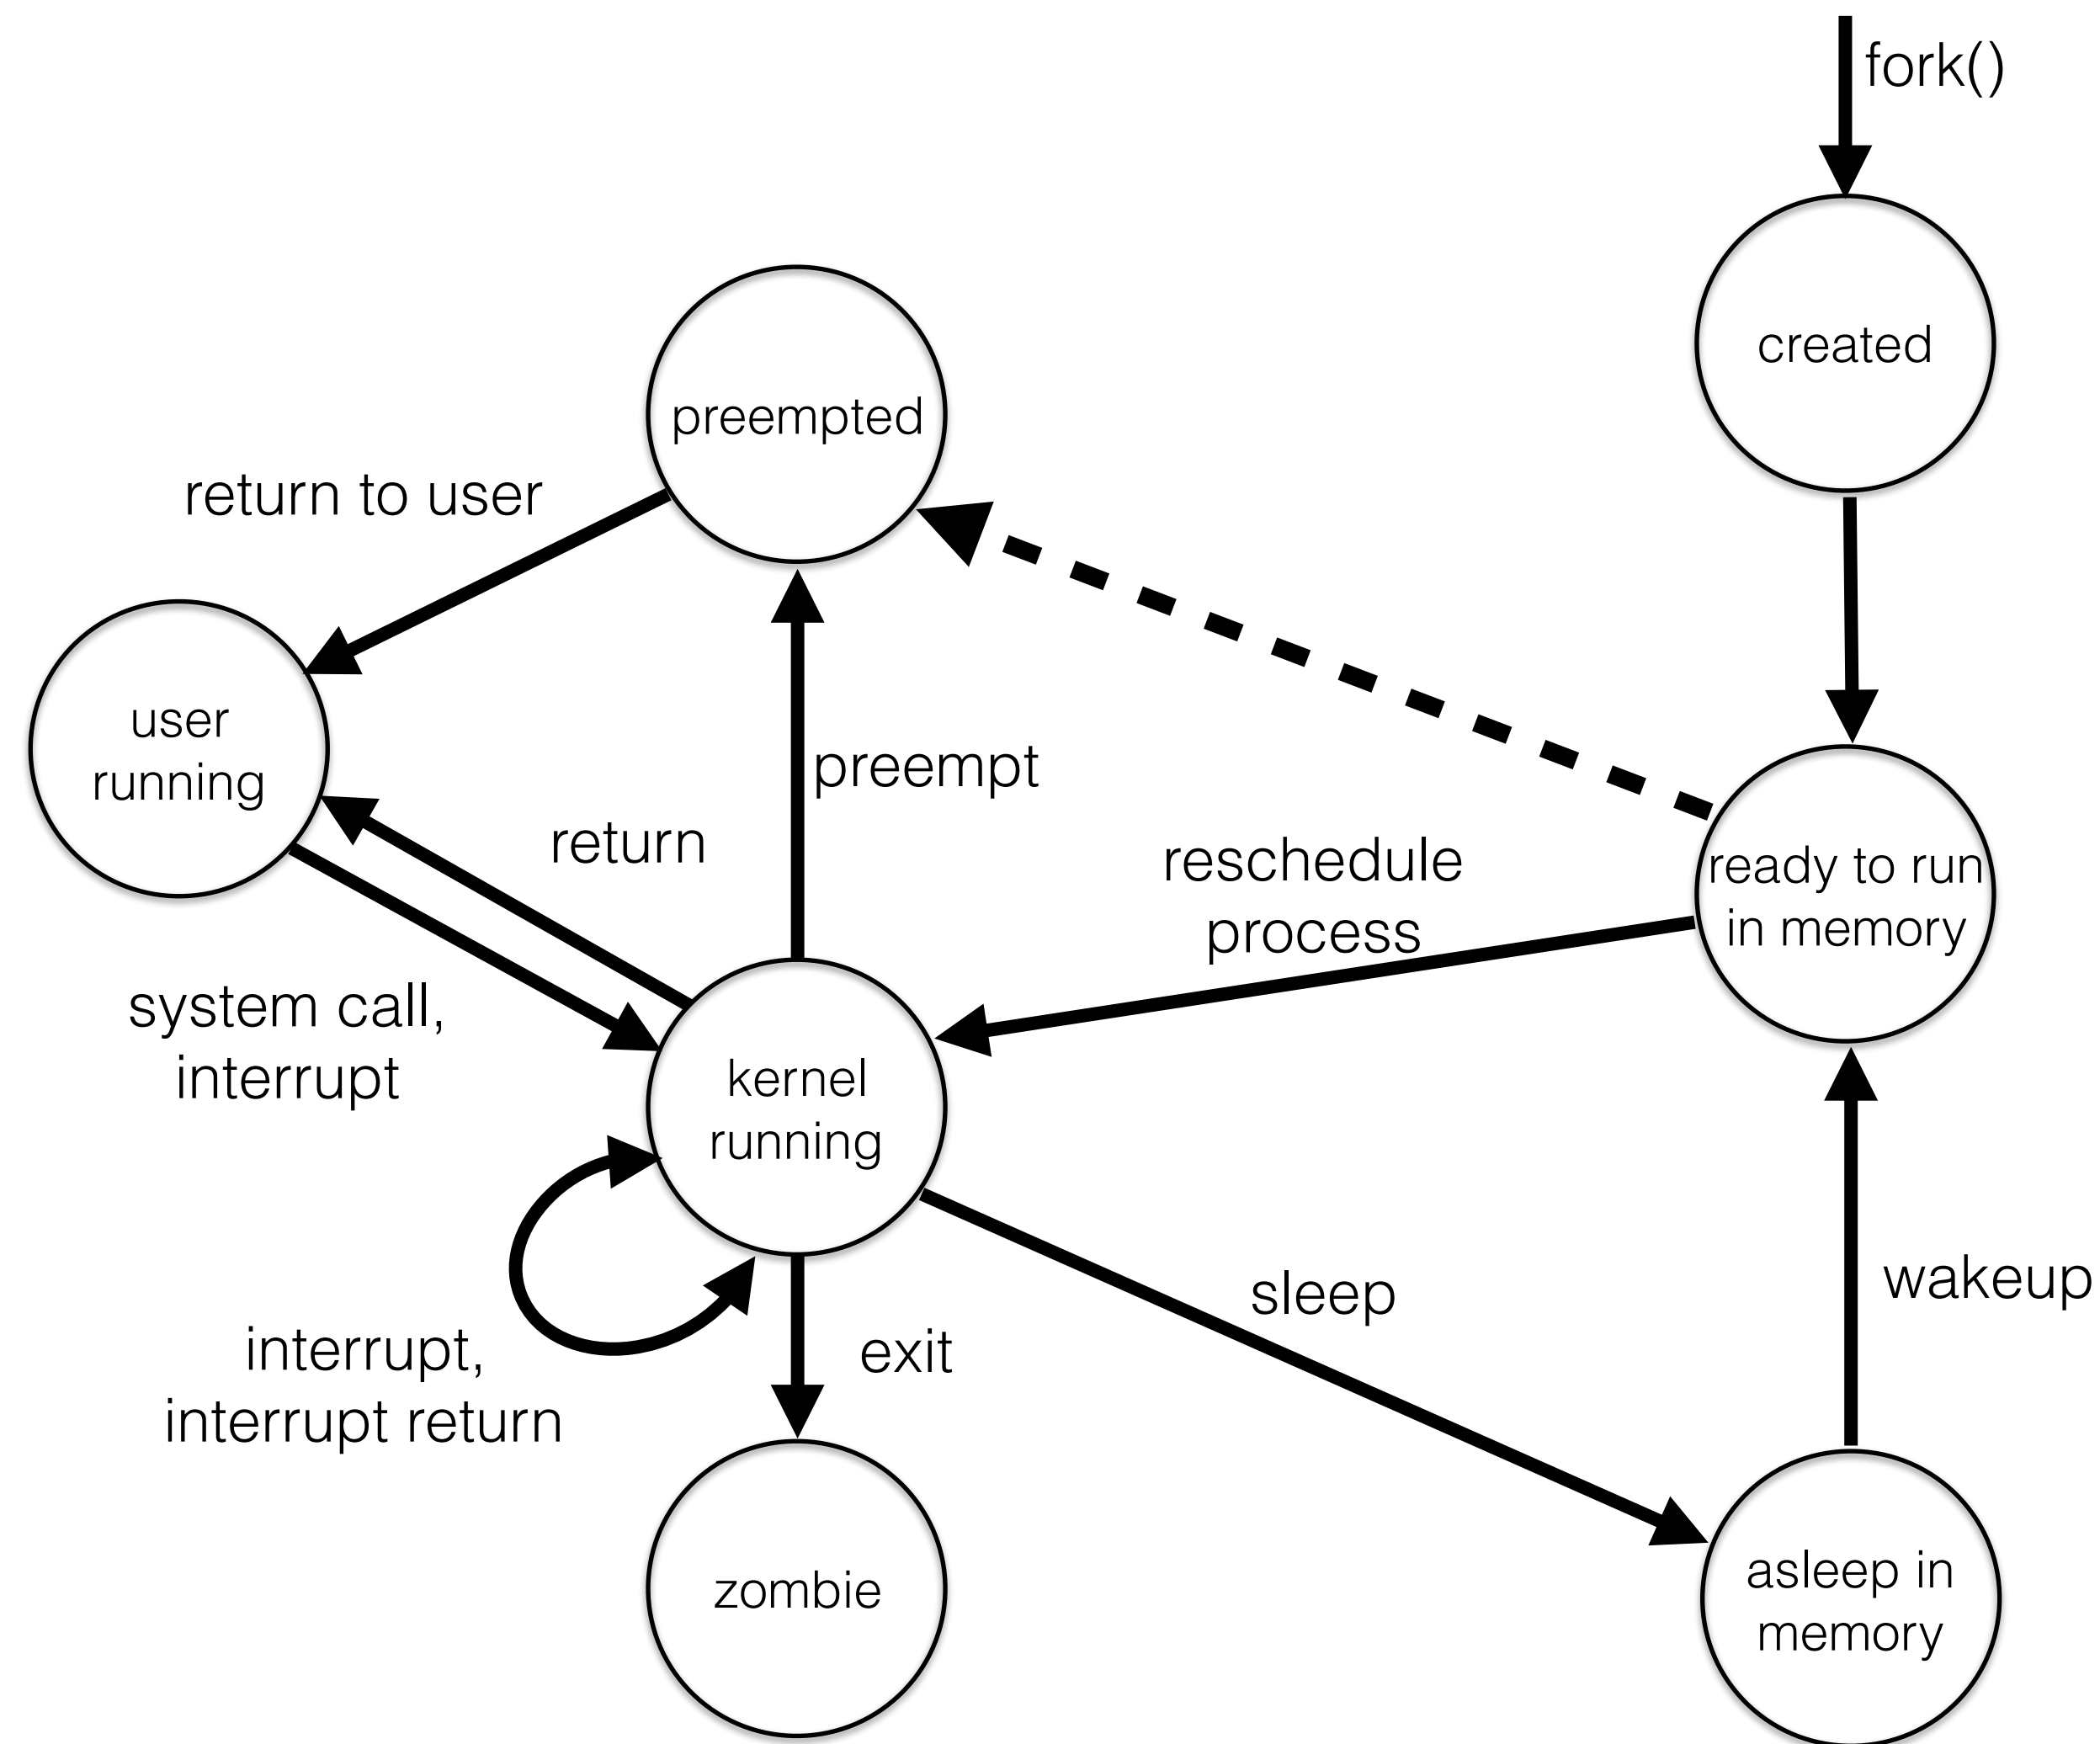
Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	
4	Running	Ready	Process ₀ now done
5	—	Running	
6	—	Running	
7	—	Running	
8	—	Running	Process ₁ now done

Example: two running processes, with I/O



Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	Process ₀ initiates I/O
4	Blocked	Running	Process ₀ is blocked, so Process ₁ runs
5	Blocked	Running	
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	Process ₁ now done
9	Running	–	
10	Running	–	Process ₀ now done

Process States (Unix)

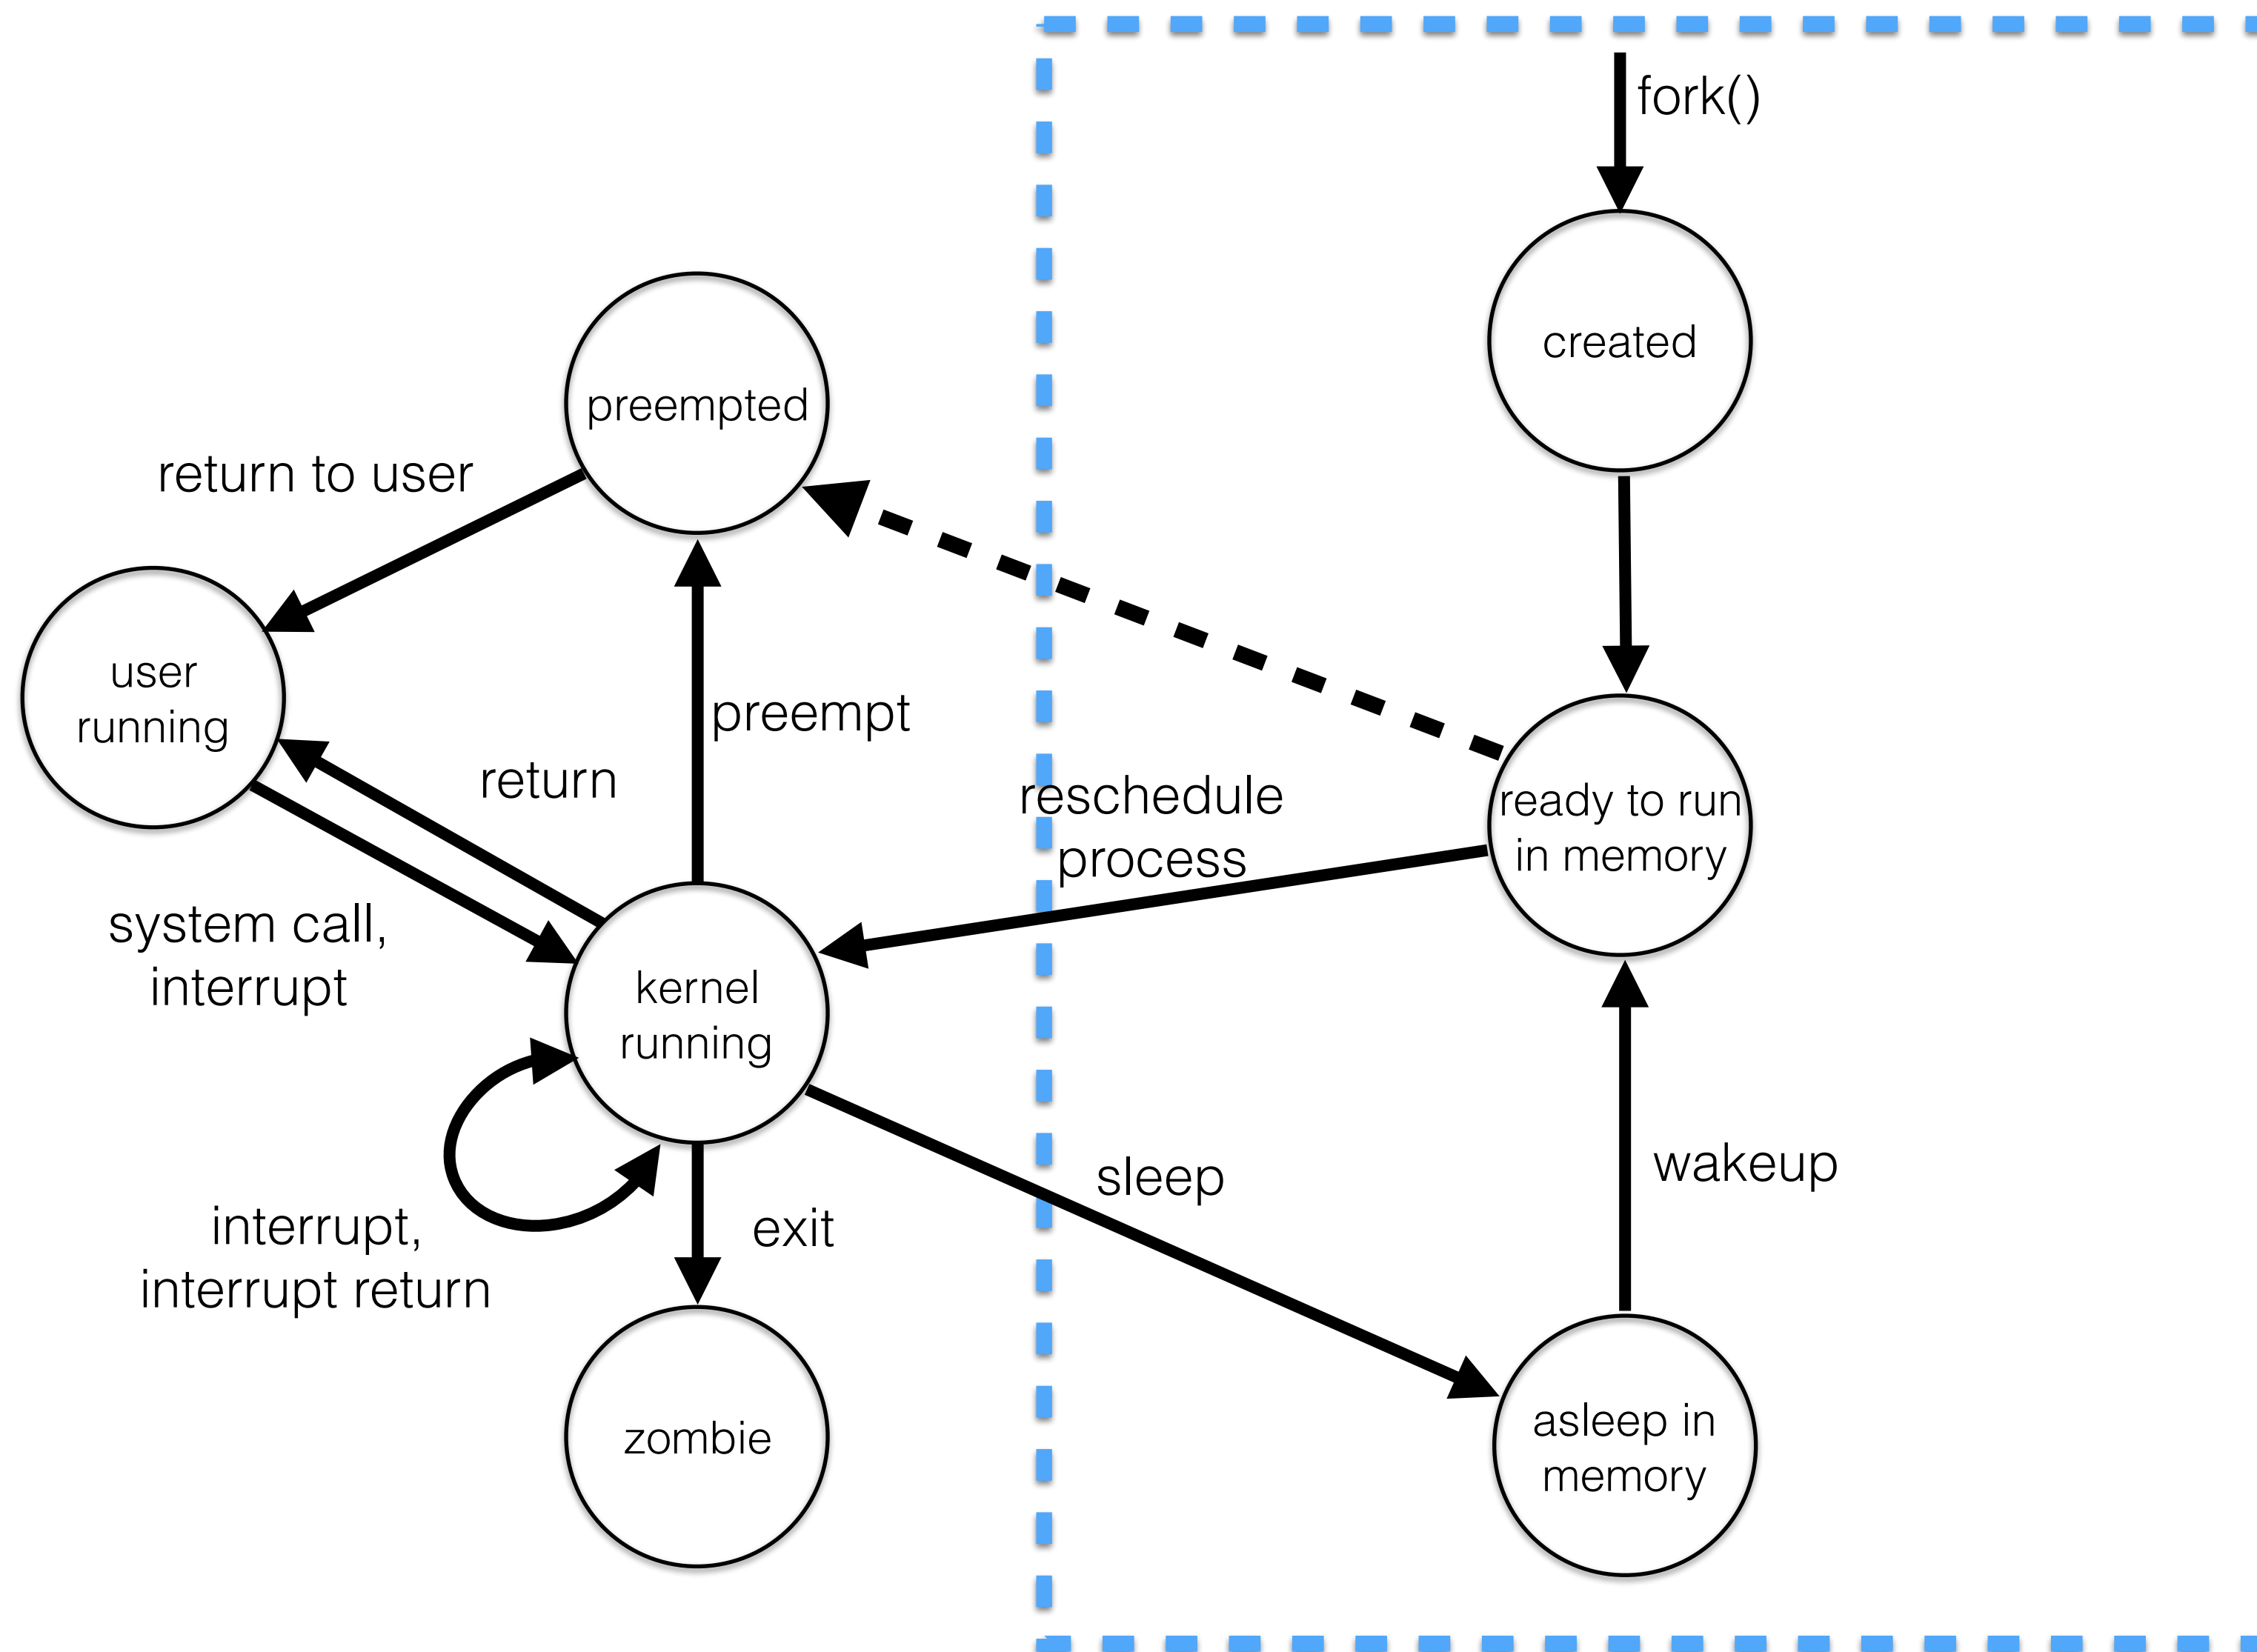


Created: Process is newly created but it is not ready to run yet.

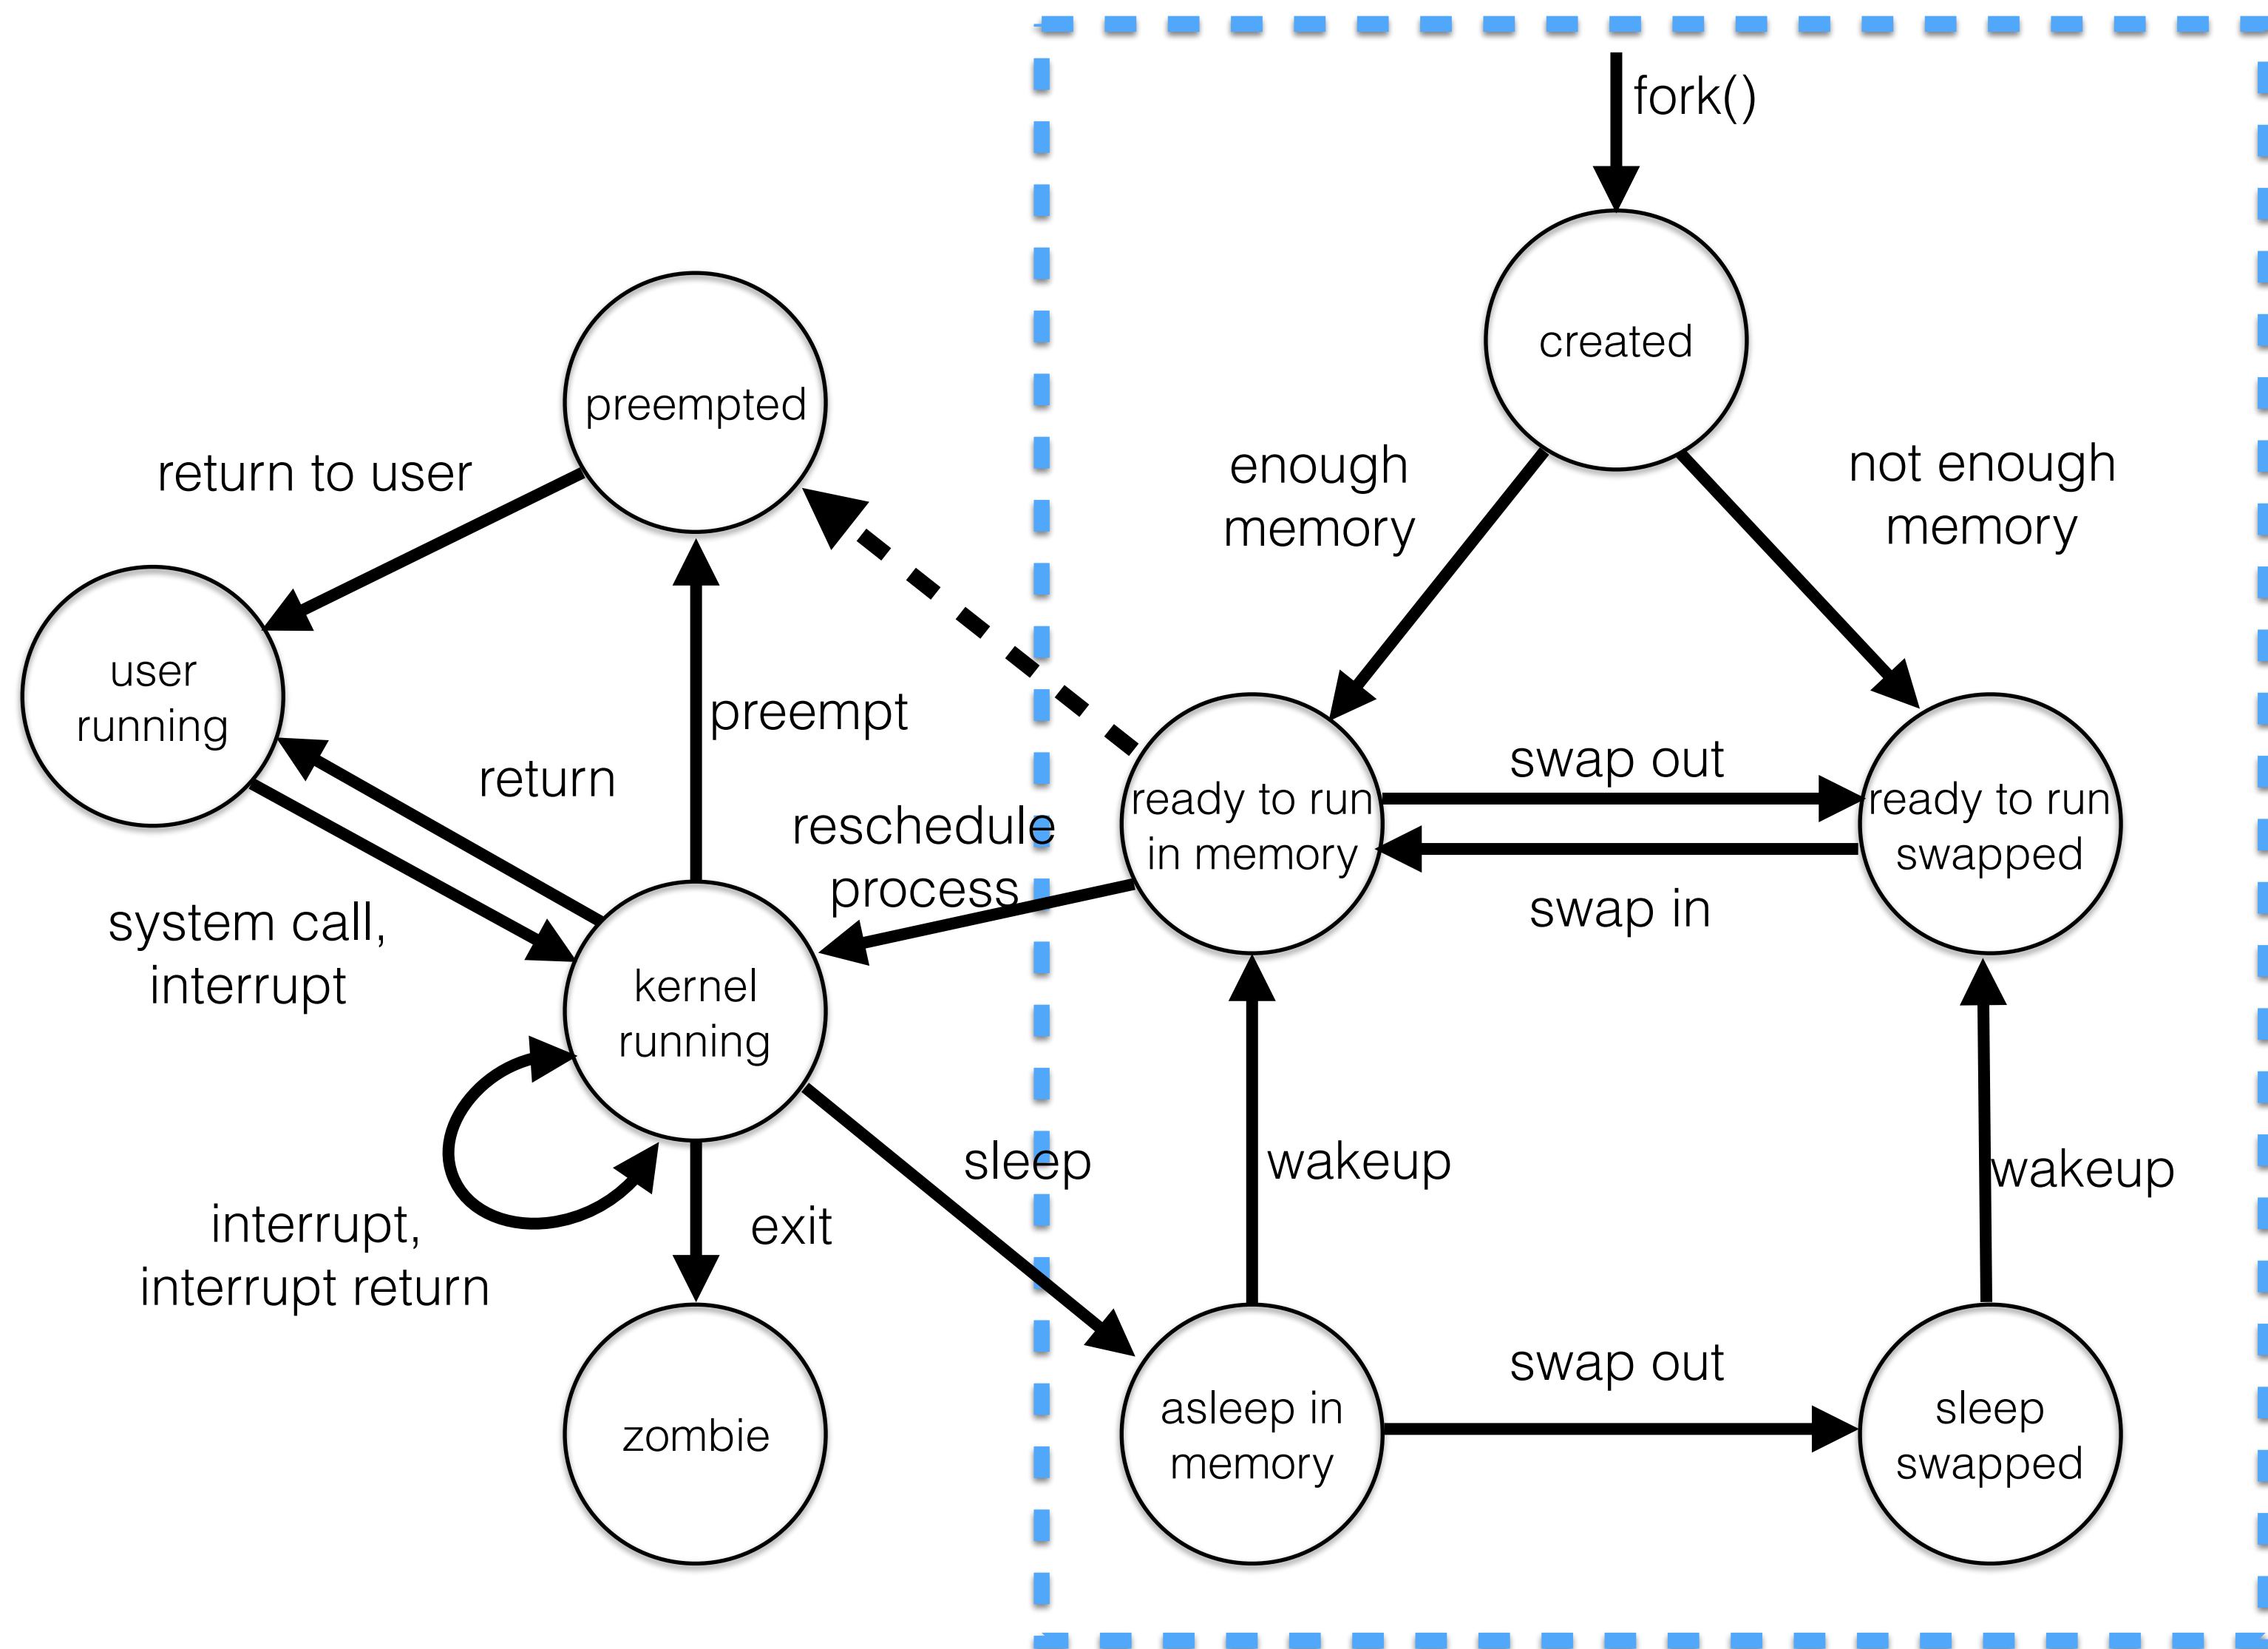
Preempted: Process is returning from kernel to user mode, but the kernel preempts it and does a process switch to schedule another process.

Zombie: Process is no longer exists, but it leaves a record for its parent process to collect.

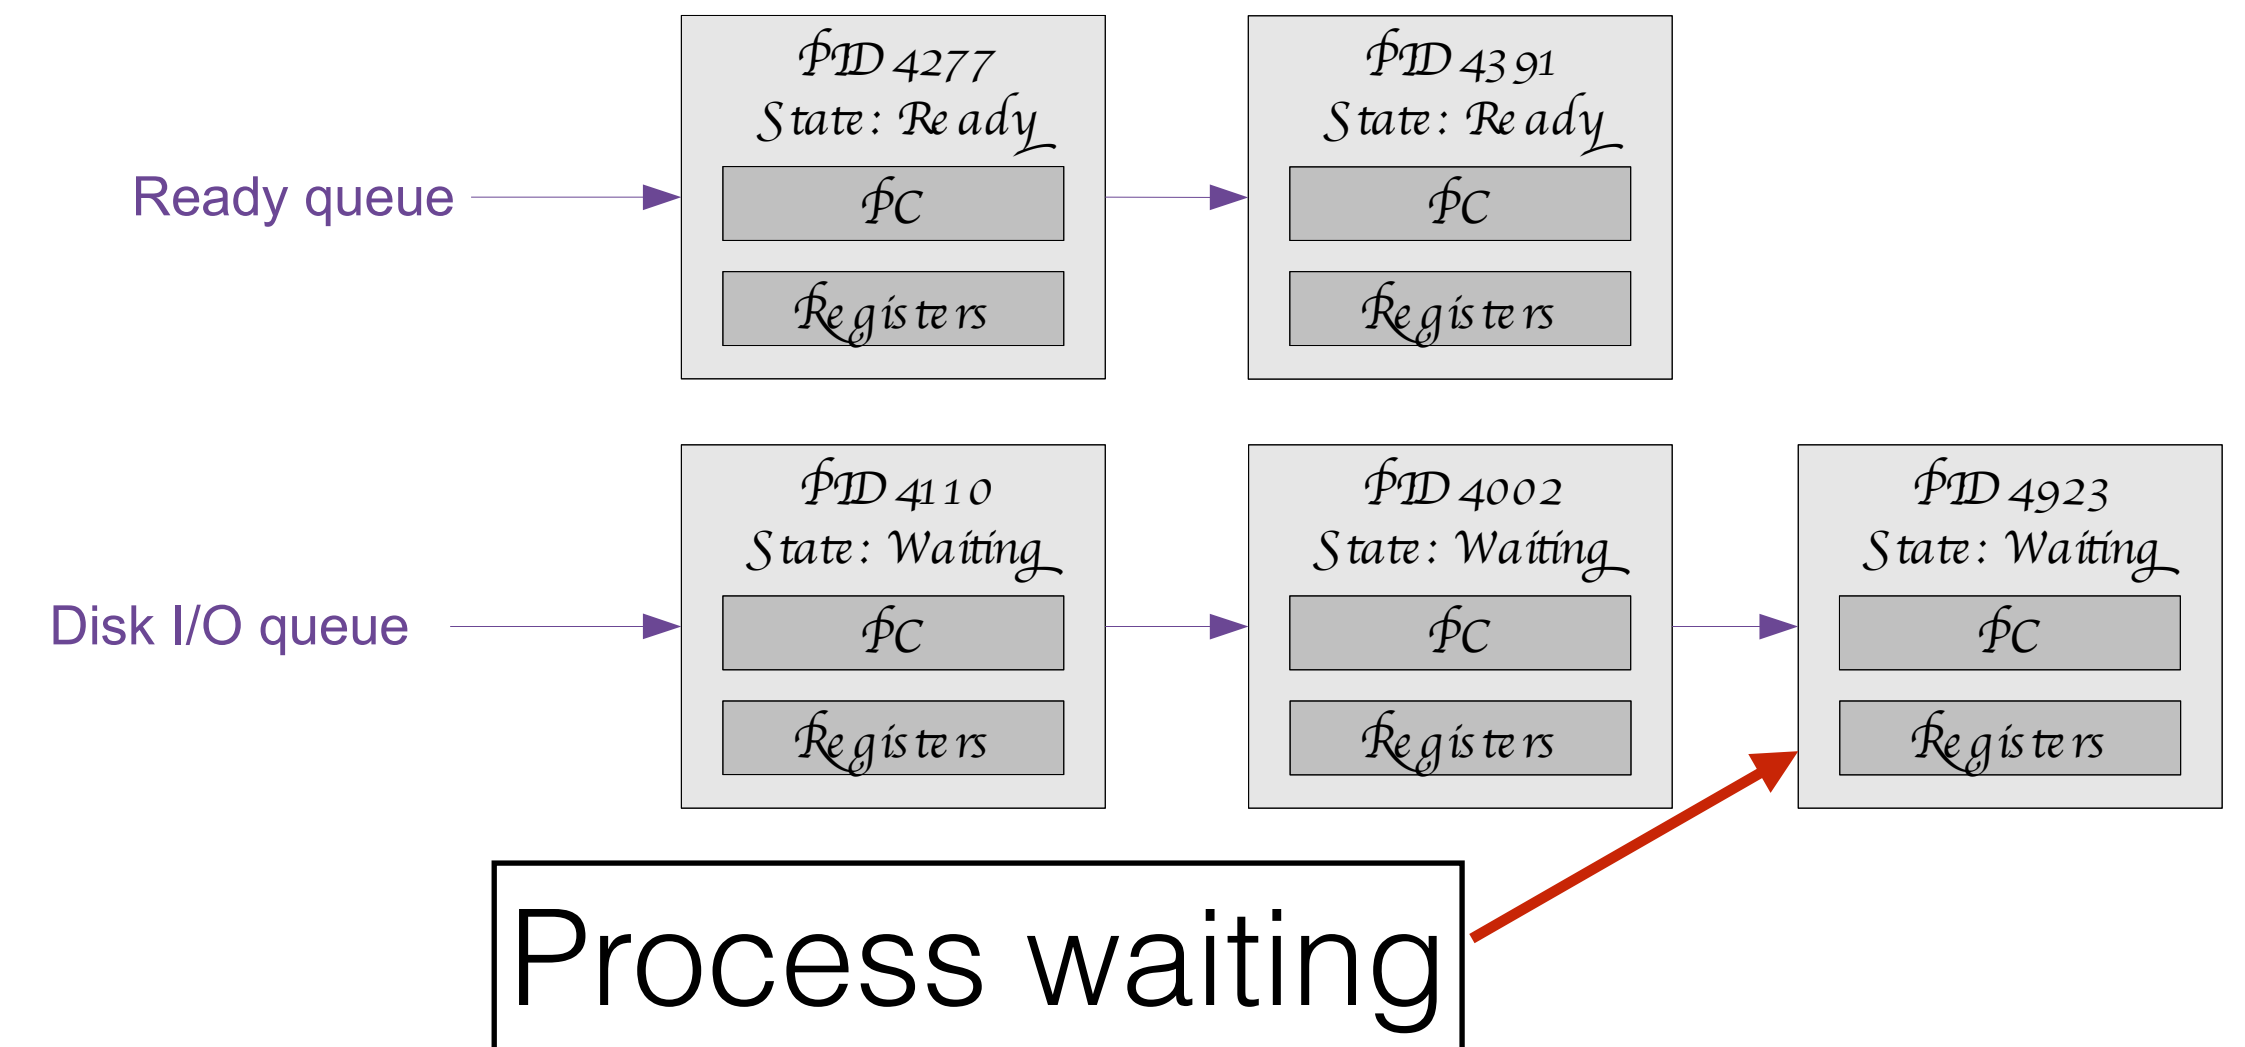

Process States (Unix)



Process States (Unix)

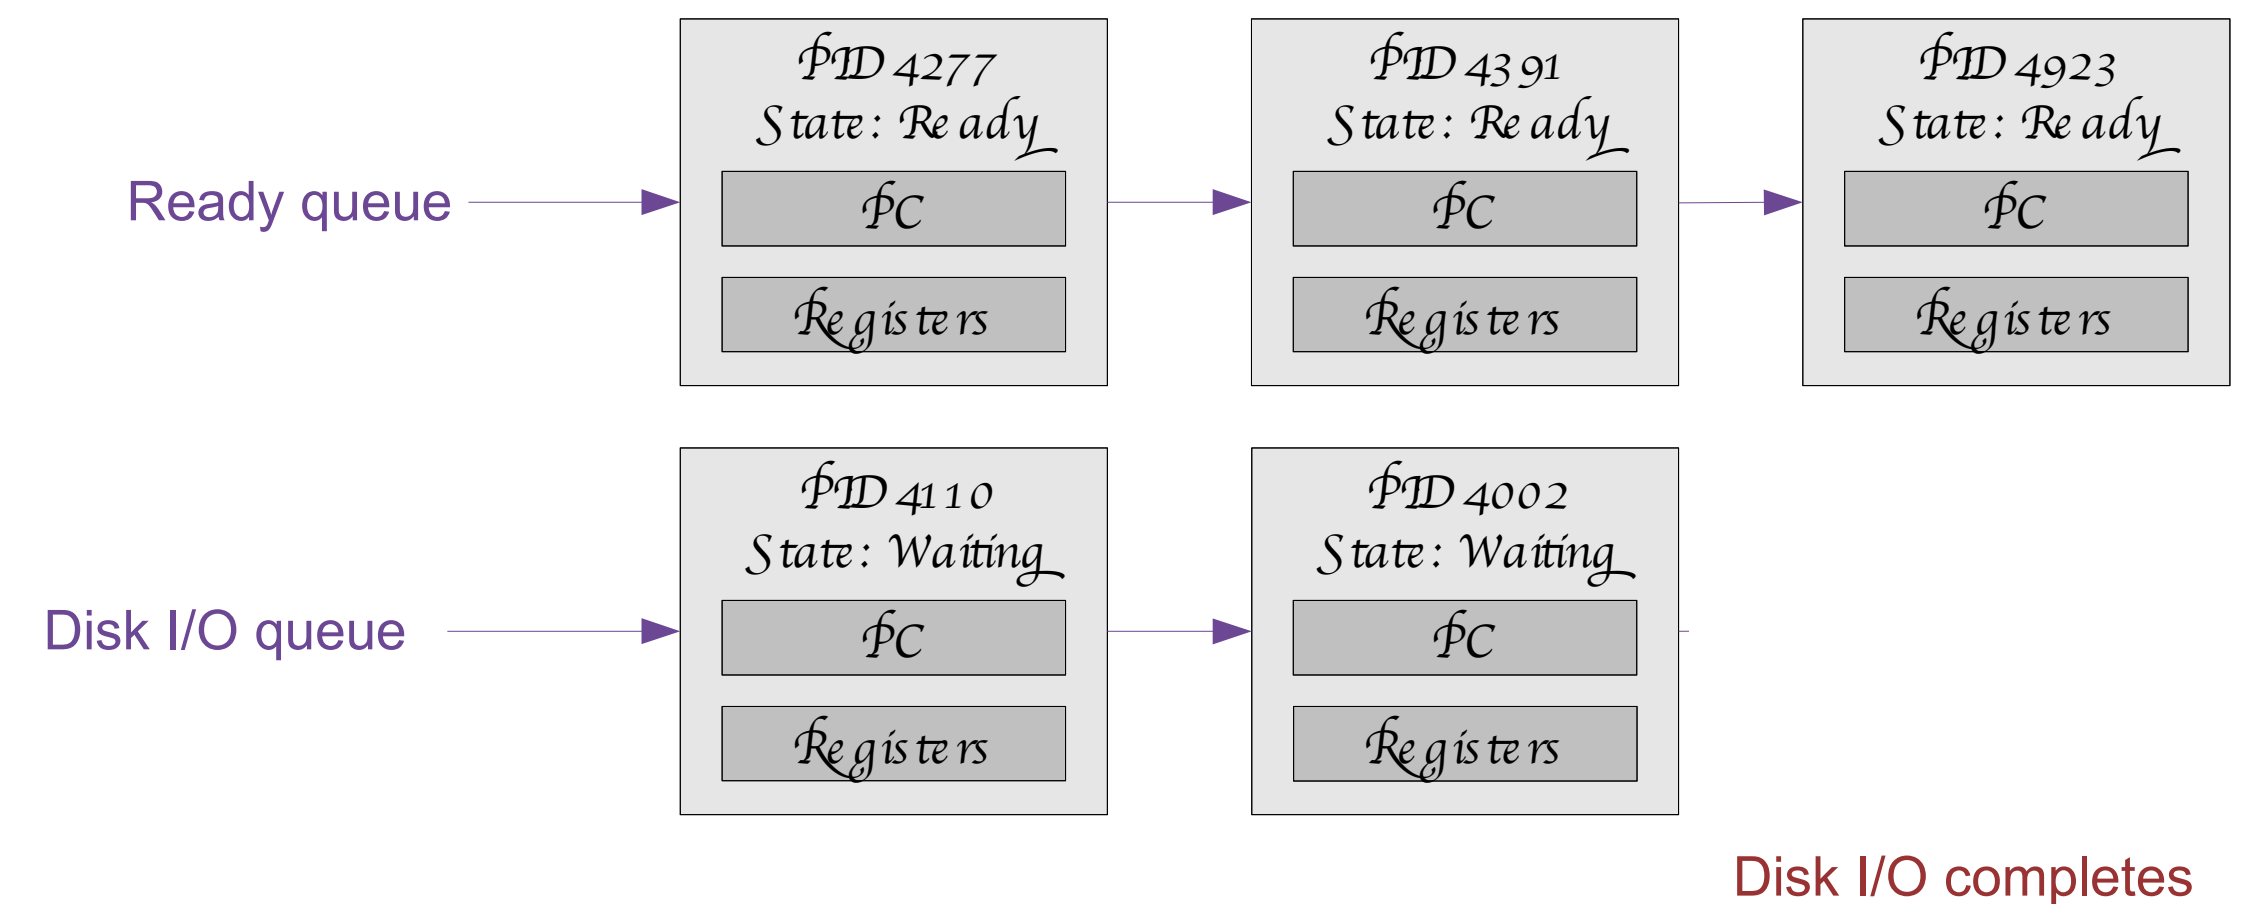
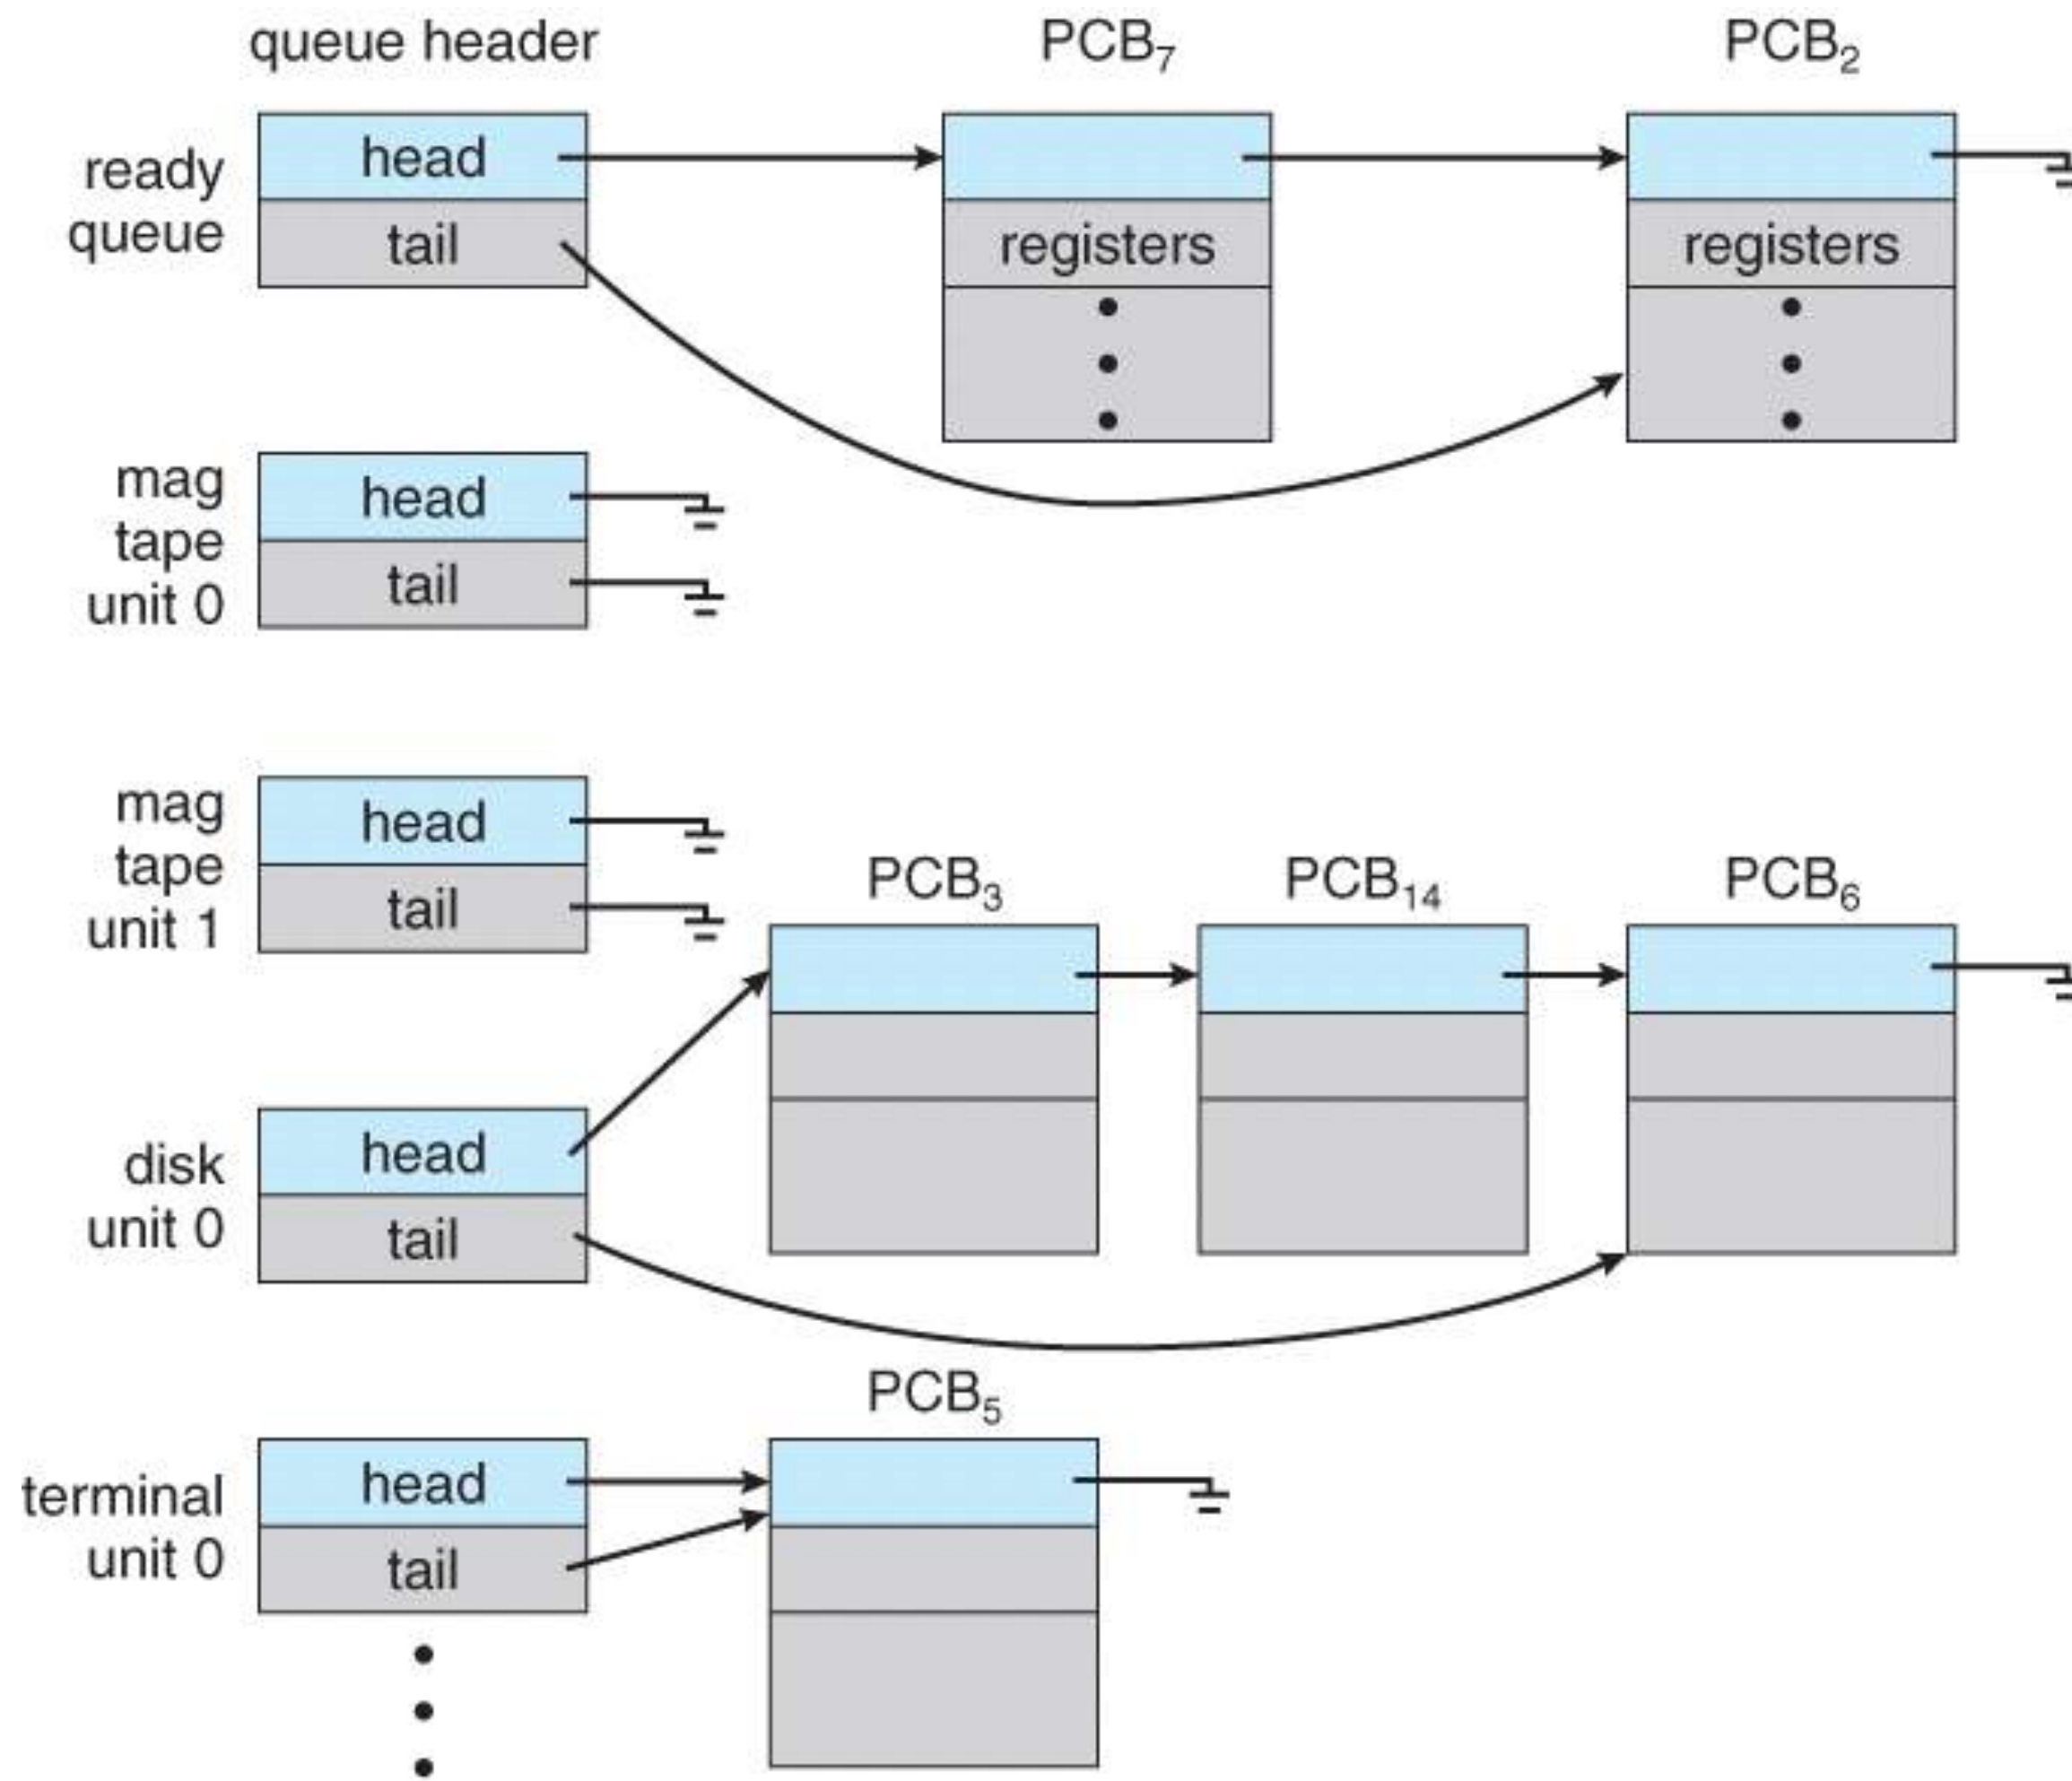


Ready queue and various I/O queues



- OS maintains a set of queues
- Each PCB is queued on a state queue based on the process' current state.
- As processes change states, PCBs are unlinked from one queue and linked into another.

Ready queue and various I/O queues



- OS maintains a set of queues
- Each PCB is queued on a state queue based on the process' current state.
- As processes change states, PCBs are unlinked from one queue and linked into another.

Main Question:

How can OS **regain control** of the CPU
from a process so that it can
switch to another process?

LA LA LAND

Approach 1: Cooperative Processes

- OS trusts processes will cooperate and give up control of CPU. For example, process can periodically calls system call `yield()`.
- Process gives up control when it causes a trap.

Approach 2: Non-Cooperative Processes

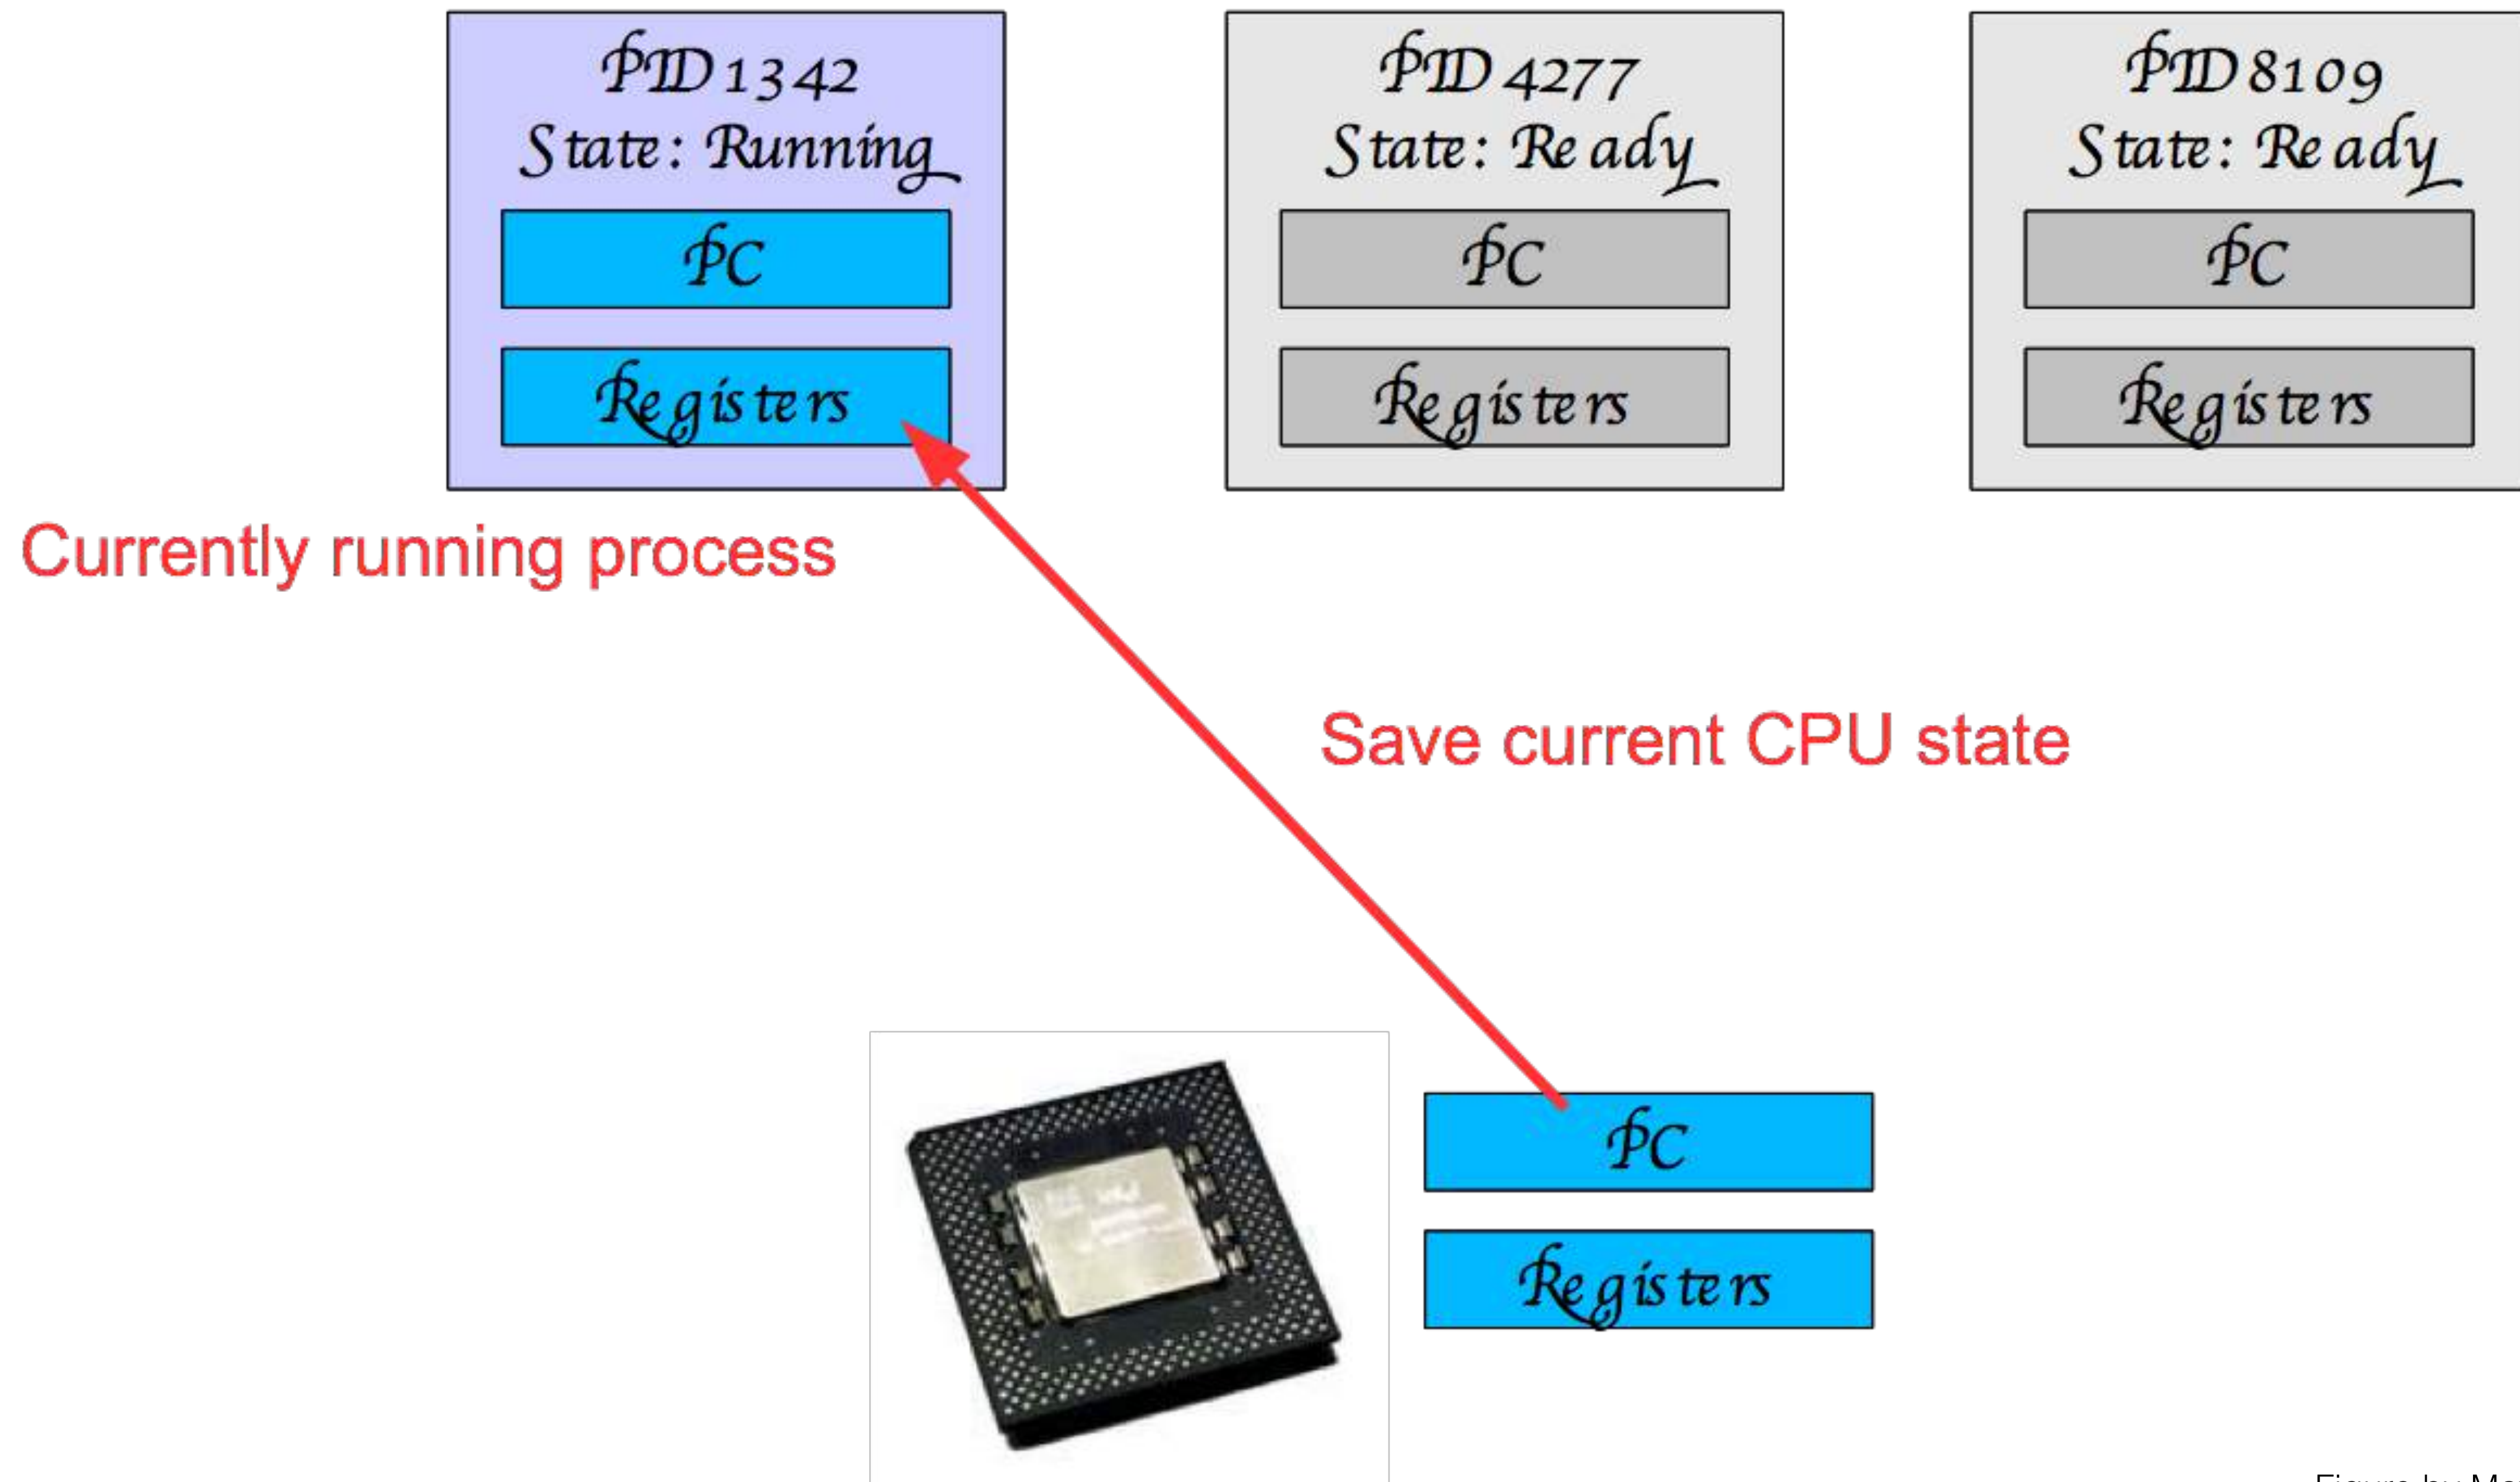
- OS takes control periodically (e.g., timer interrupt).
- Timer can be programmed to raise an interrupt periodically.
- When interrupt is raised, OS *Interrupt Handler* runs, and OS regains control.






Now, OS has control. How to switch to another process?

- OS decides the process to which to switch (i.e., scheduler decides).
- OS executes a piece of assembly code (i.e., context switch).
- **Context switch:**
 1. Save register values of current process to kernel stack.
 2. Restore register values of the next process from its kernel stack.


CPU switch from process to process



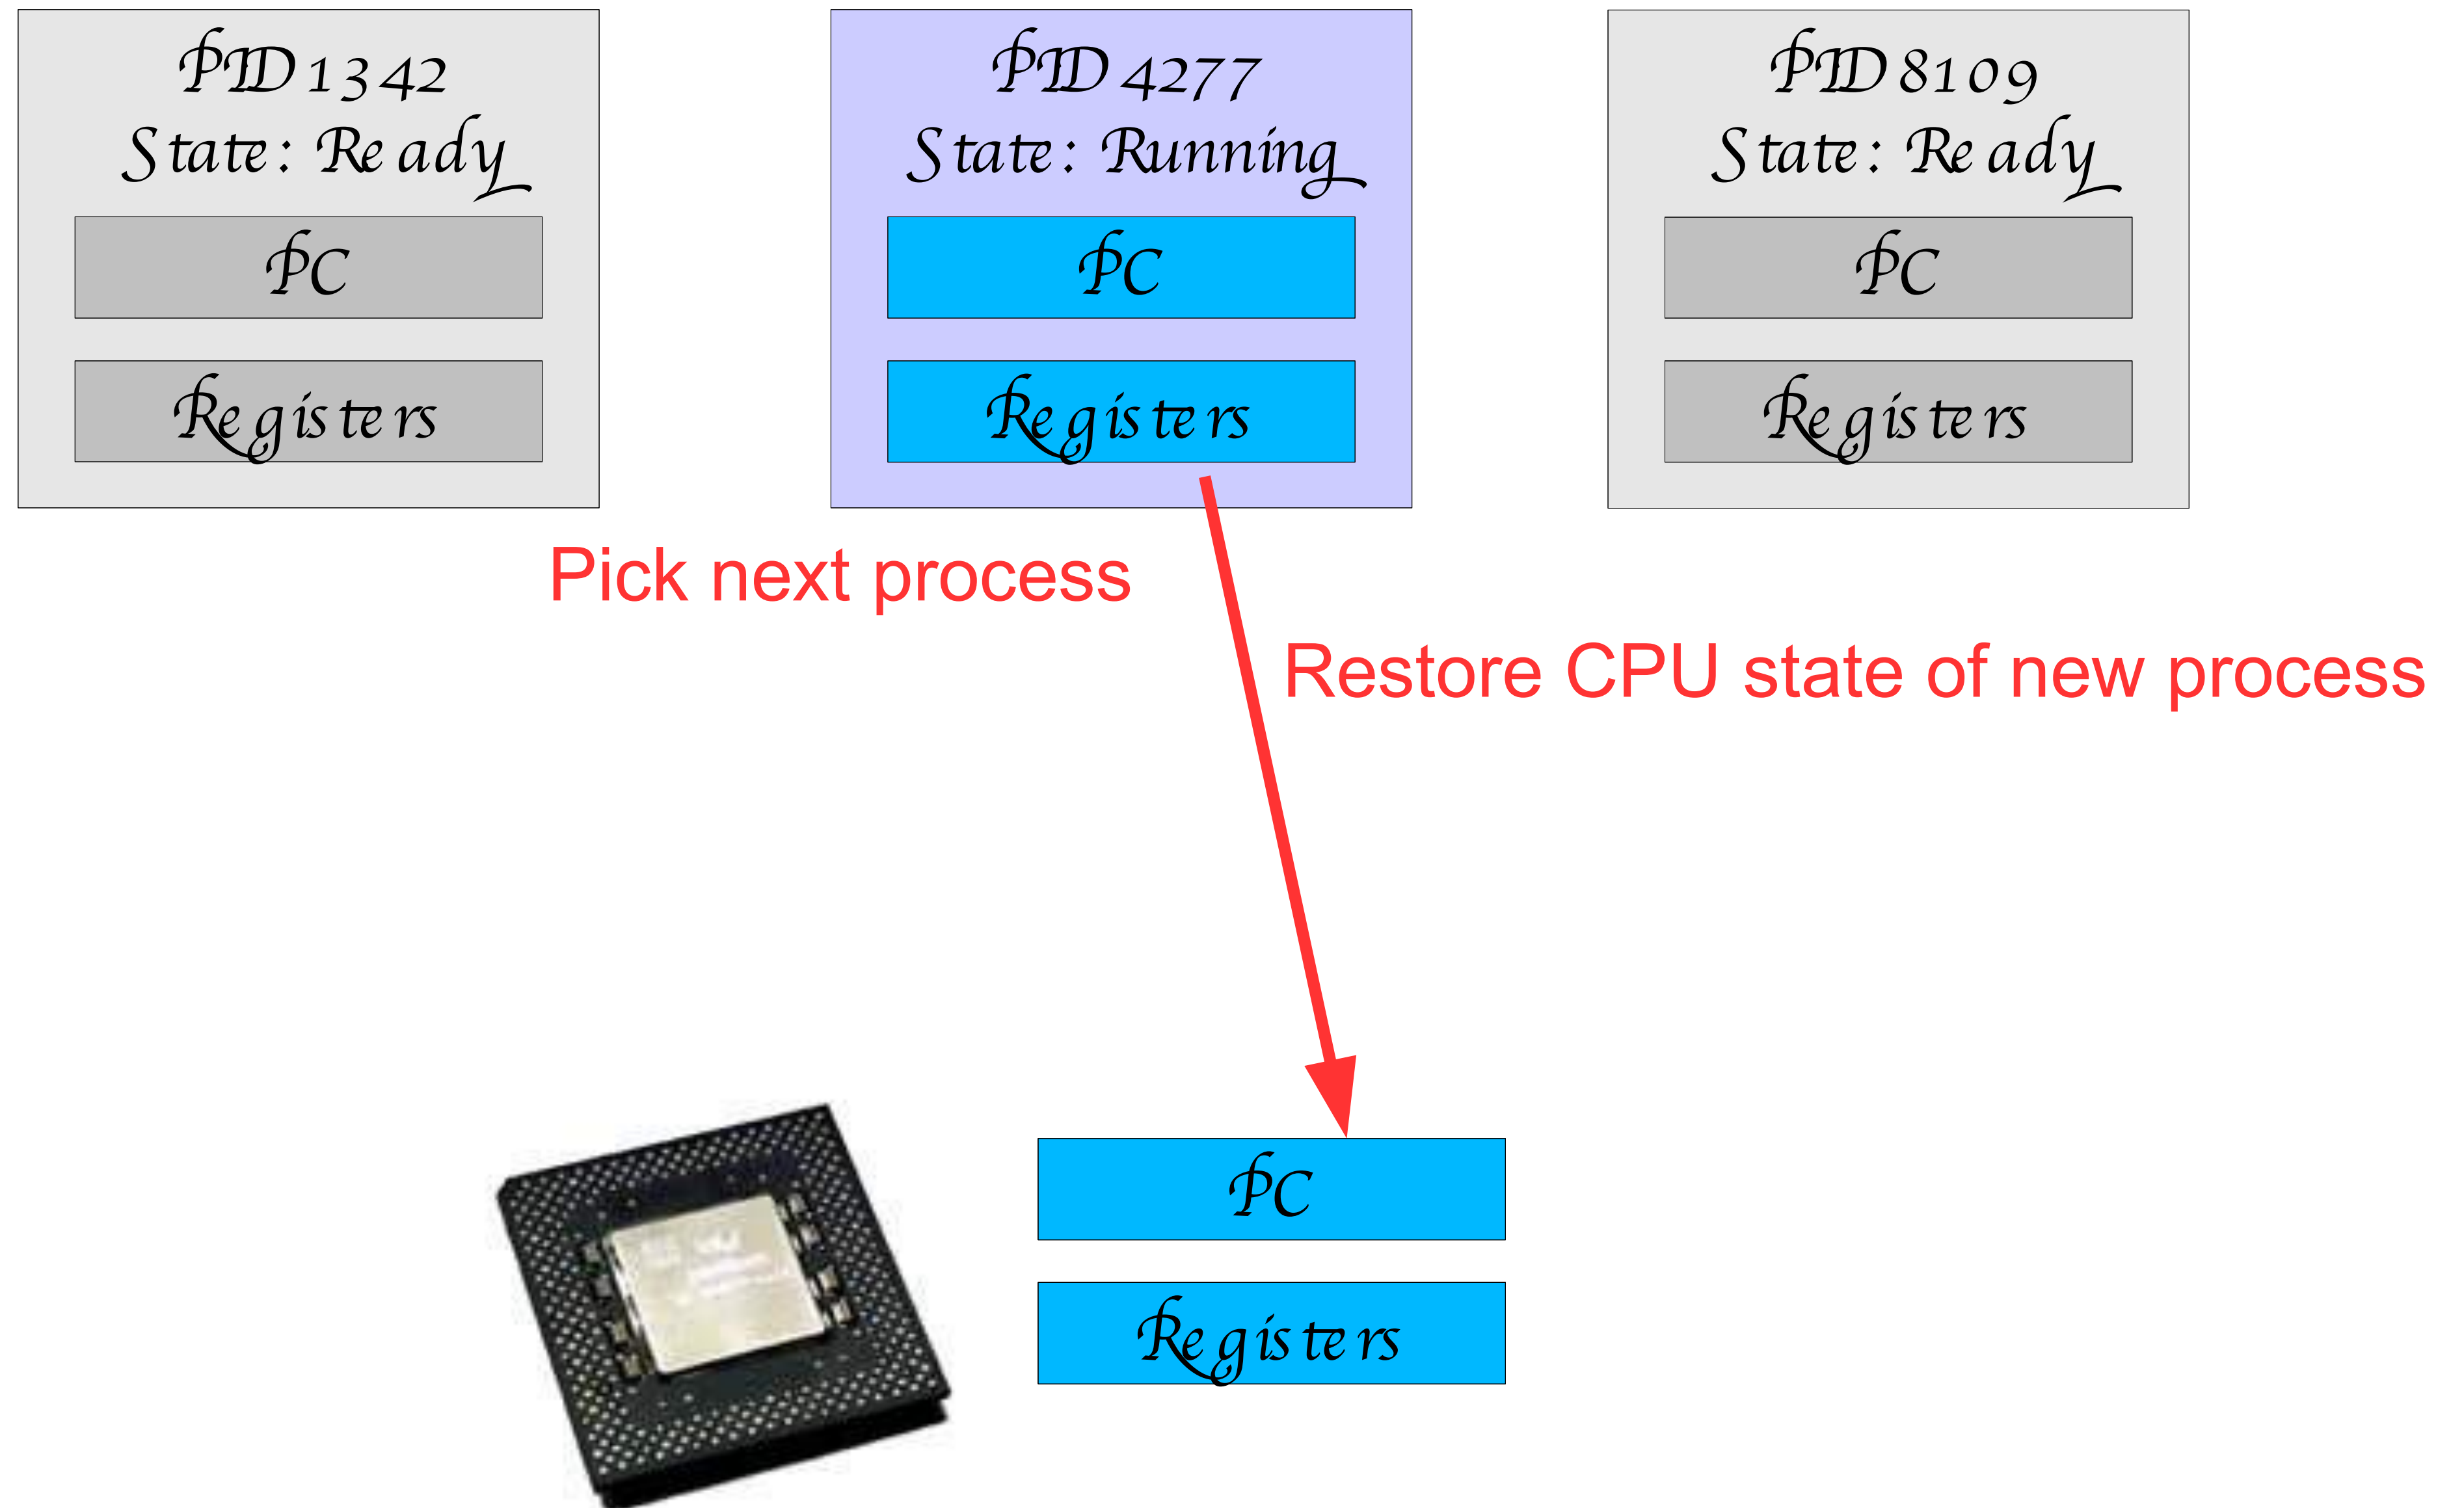
CPU switch from process to process



Suspend process



CPU switch from process to process



Context switch in OS/161

kern/thread/thread.c

```
593  * The current thread is queued appropriately and its state is changed
594  * to NEWSTATE; another thread to run is selected and switched to.
595  *
596  * If NEWSTATE is S_SLEEP, the thread is queued on the wait channel
597  * WC. Otherwise WC should be NULL.
598  */
599  static
600  void
601  thread_switch(threadstate_t newstate, struct wchan *wc)
602  {
603      struct thread *cur, *next;
604      int spl;
605
606      DEBUGASSERT(curcpu->c_curthread == curthread);
607      DEBUGASSERT(curthread->t_cpu == curcpu->c_self);
608
609      /* Explicitly disable interrupts on this processor */
610      spl = splhigh();
611
612      cur = curthread;
613
```



Context switch in OS/161

kern/thread/thread.c

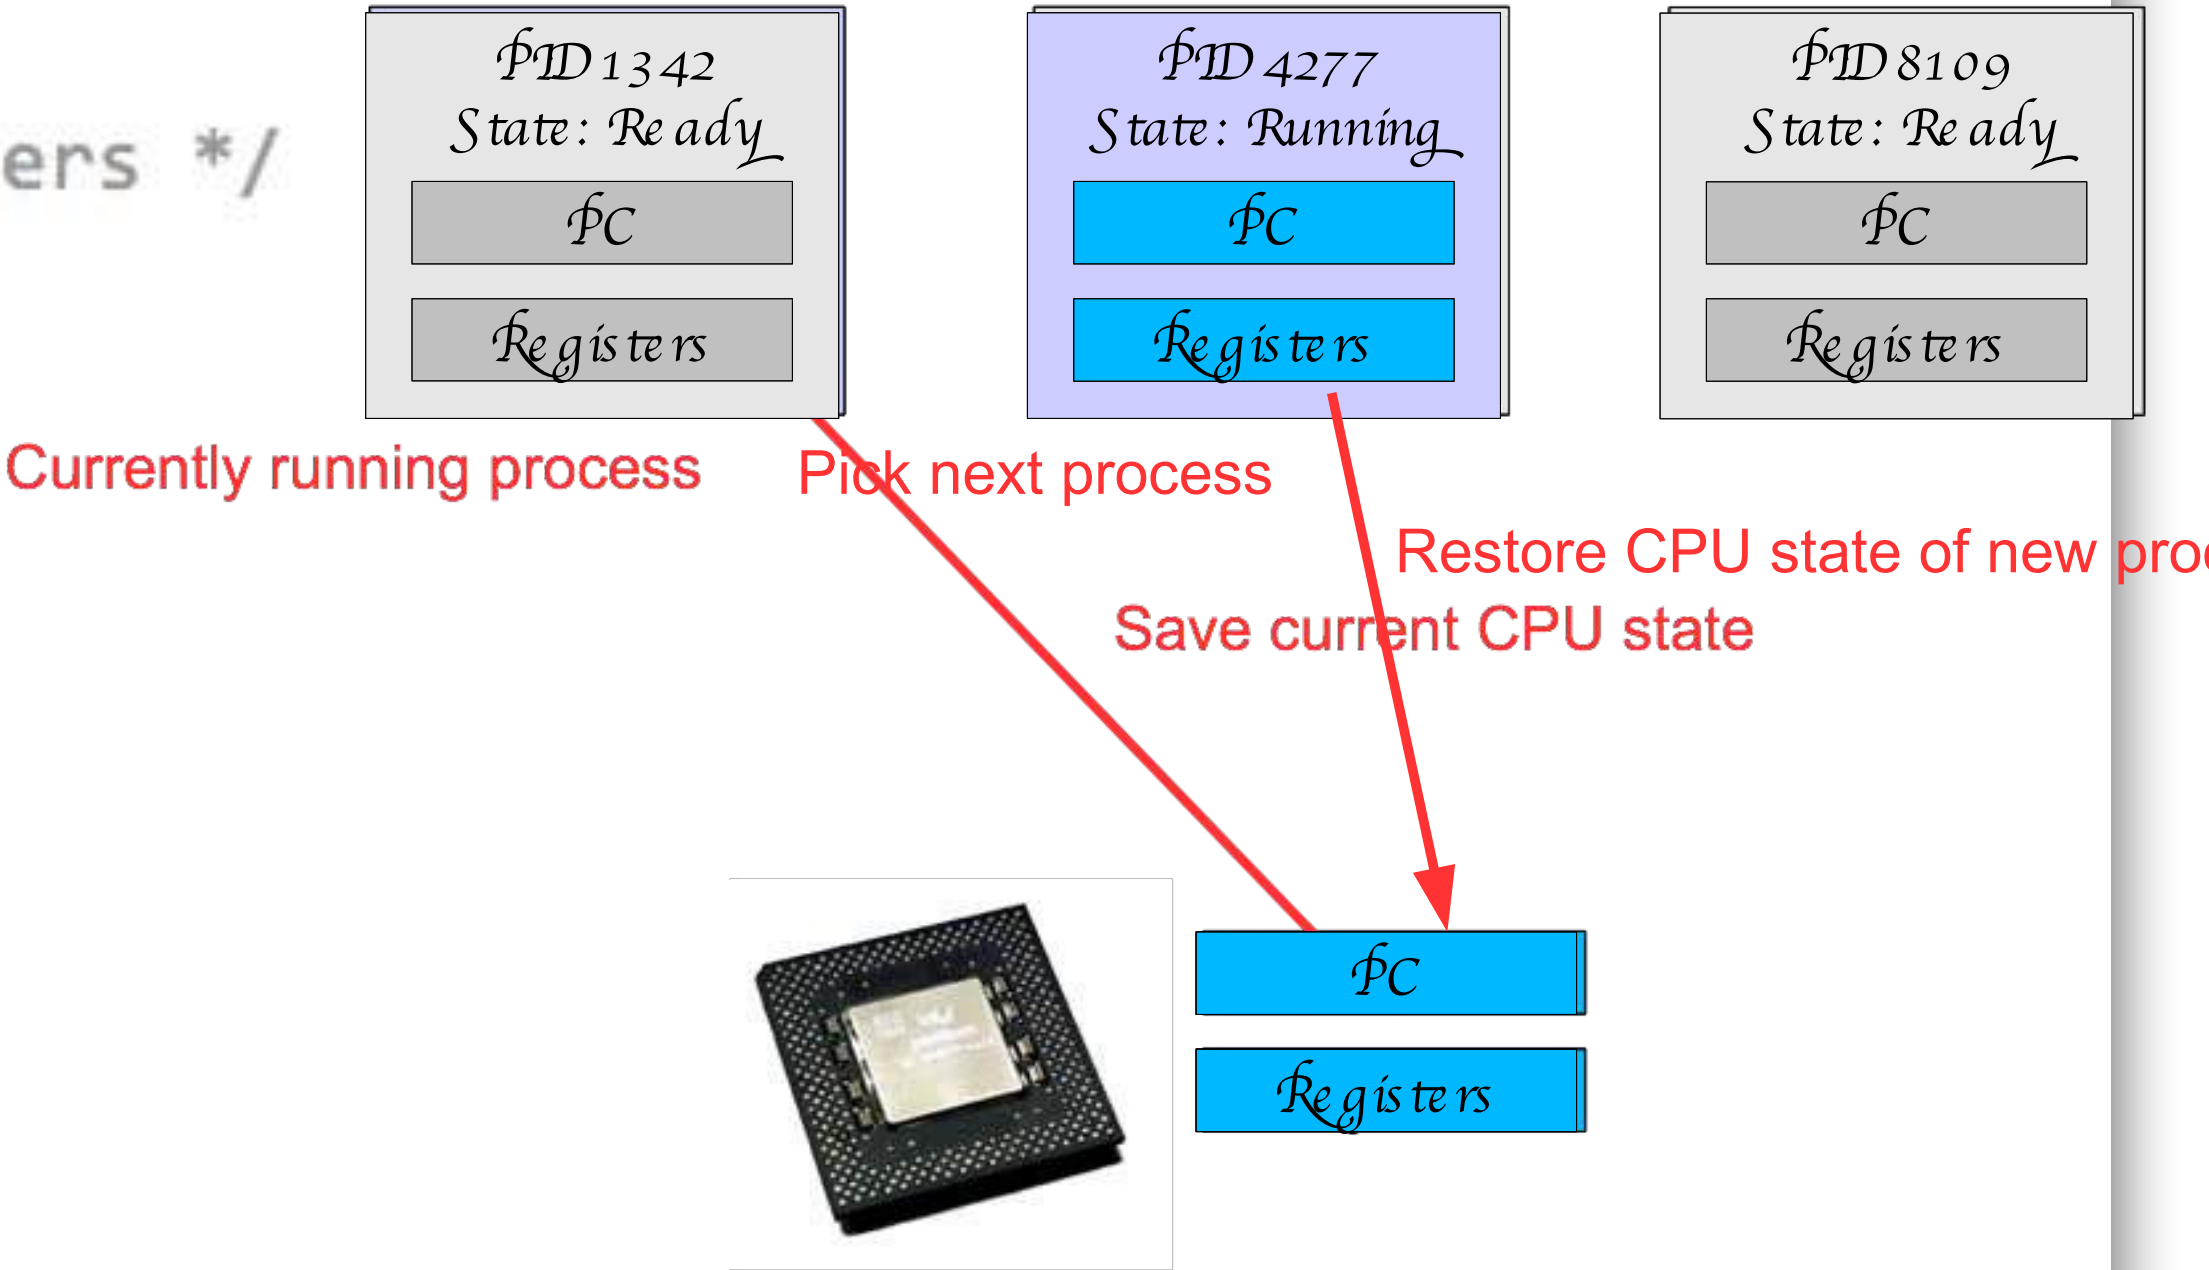
```
694 } while (next == NULL);
695 curcpu->c_isidle = false;
696
697 /*
698  * Note that curcpu->c_curthread may be the same variable as
699  * curthread and it may not be, depending on how curthread and
700  * curcpu are defined by the MD code. We'll assign both and
701  * assume the compiler will optimize one away if they're the
702  * same.
703  */
704 curcpu->c_curthread = next;
705 curthread = next;
706
707 /* do the switch (in assembler in switch.S) */
708 switchframe_switch(&cur->t_context, &next->t_context);
709
710 /*
711  * When we get to this point we are either running in the next
712  * thread, or have come back to the same thread again,
713  * depending on how you look at it. That is,
```


Context switch in OS/161

src/kern/arch/mips/
thread/switch.S

SW saves a word from a register into RAM.

```
61      /* Allocate stack space for saving 10 registers. 10*4 = 40 */
62      addi sp, sp, -40
63
64      /* Save the registers */
65      SW    ra, 36(sp)
66      SW    gp, 32(sp)
67      SW    s8, 28(sp)
68      SW    s6, 24(sp)
69      SW    s5, 20(sp)
70      SW    s4, 16(sp)
71      SW    s3, 12(sp)
72      SW    s2, 8(sp)
73      SW    s1, 4(sp)
74      SW    s0, 0(sp)
75
76      /* Store the old stack pointer in the old thread */
77      SW    sp, 0(a0)
78
79      /* Get the new stack pointer from the new thread */
80      lw    sp, 0(a1)
81      nop          /* delay slot for load */
82
```

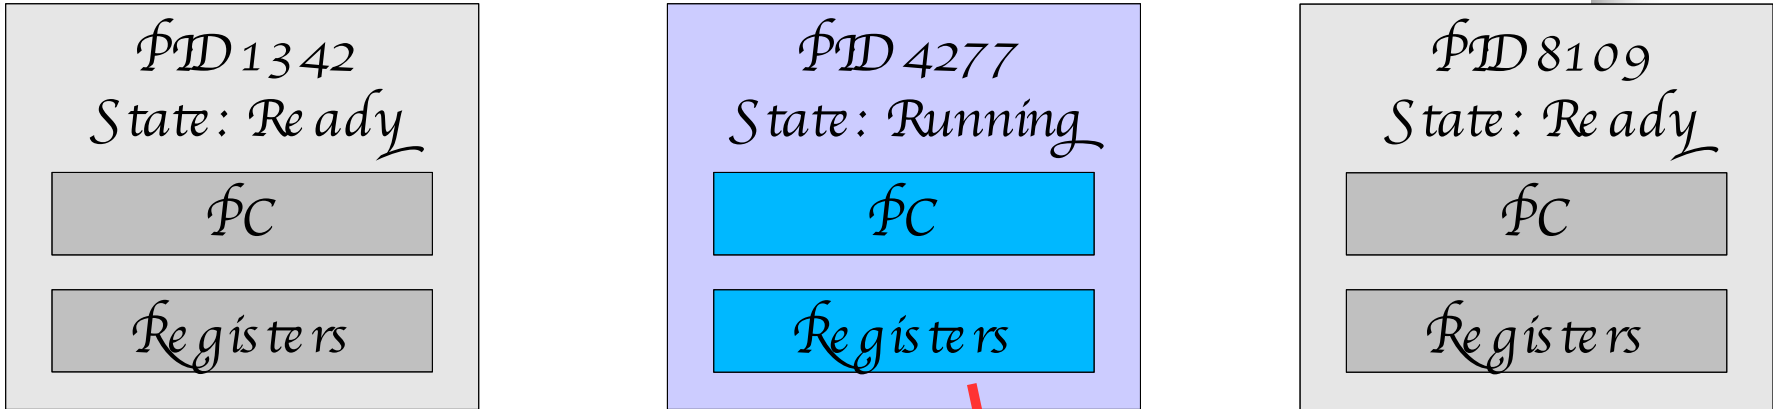


Context switch in OS/161

src/kern/arch/mips/
thread/switch.S

lw loads a word memory
into a register.

```
83      /* Now, restore the registers */
84      lw      s0, 0(sp)
85      lw      s1, 4(sp)
86      lw      s2, 8(sp)
87      lw      s3, 12(sp)
88      lw      s4, 16(sp)
89      lw      s5, 20(sp)
90      lw      s6, 24(sp)
91      lw      s8, 28(sp)
92      lw      gp, 32(sp)
93      lw      ra, 36(sp)
94      nop
95
96      /* and return. */
97      j ra
98      addi sp, sp, 40      /* in delay slot */
99      .end switchframe_switch
```



Pick next process

Restore CPU state of new process



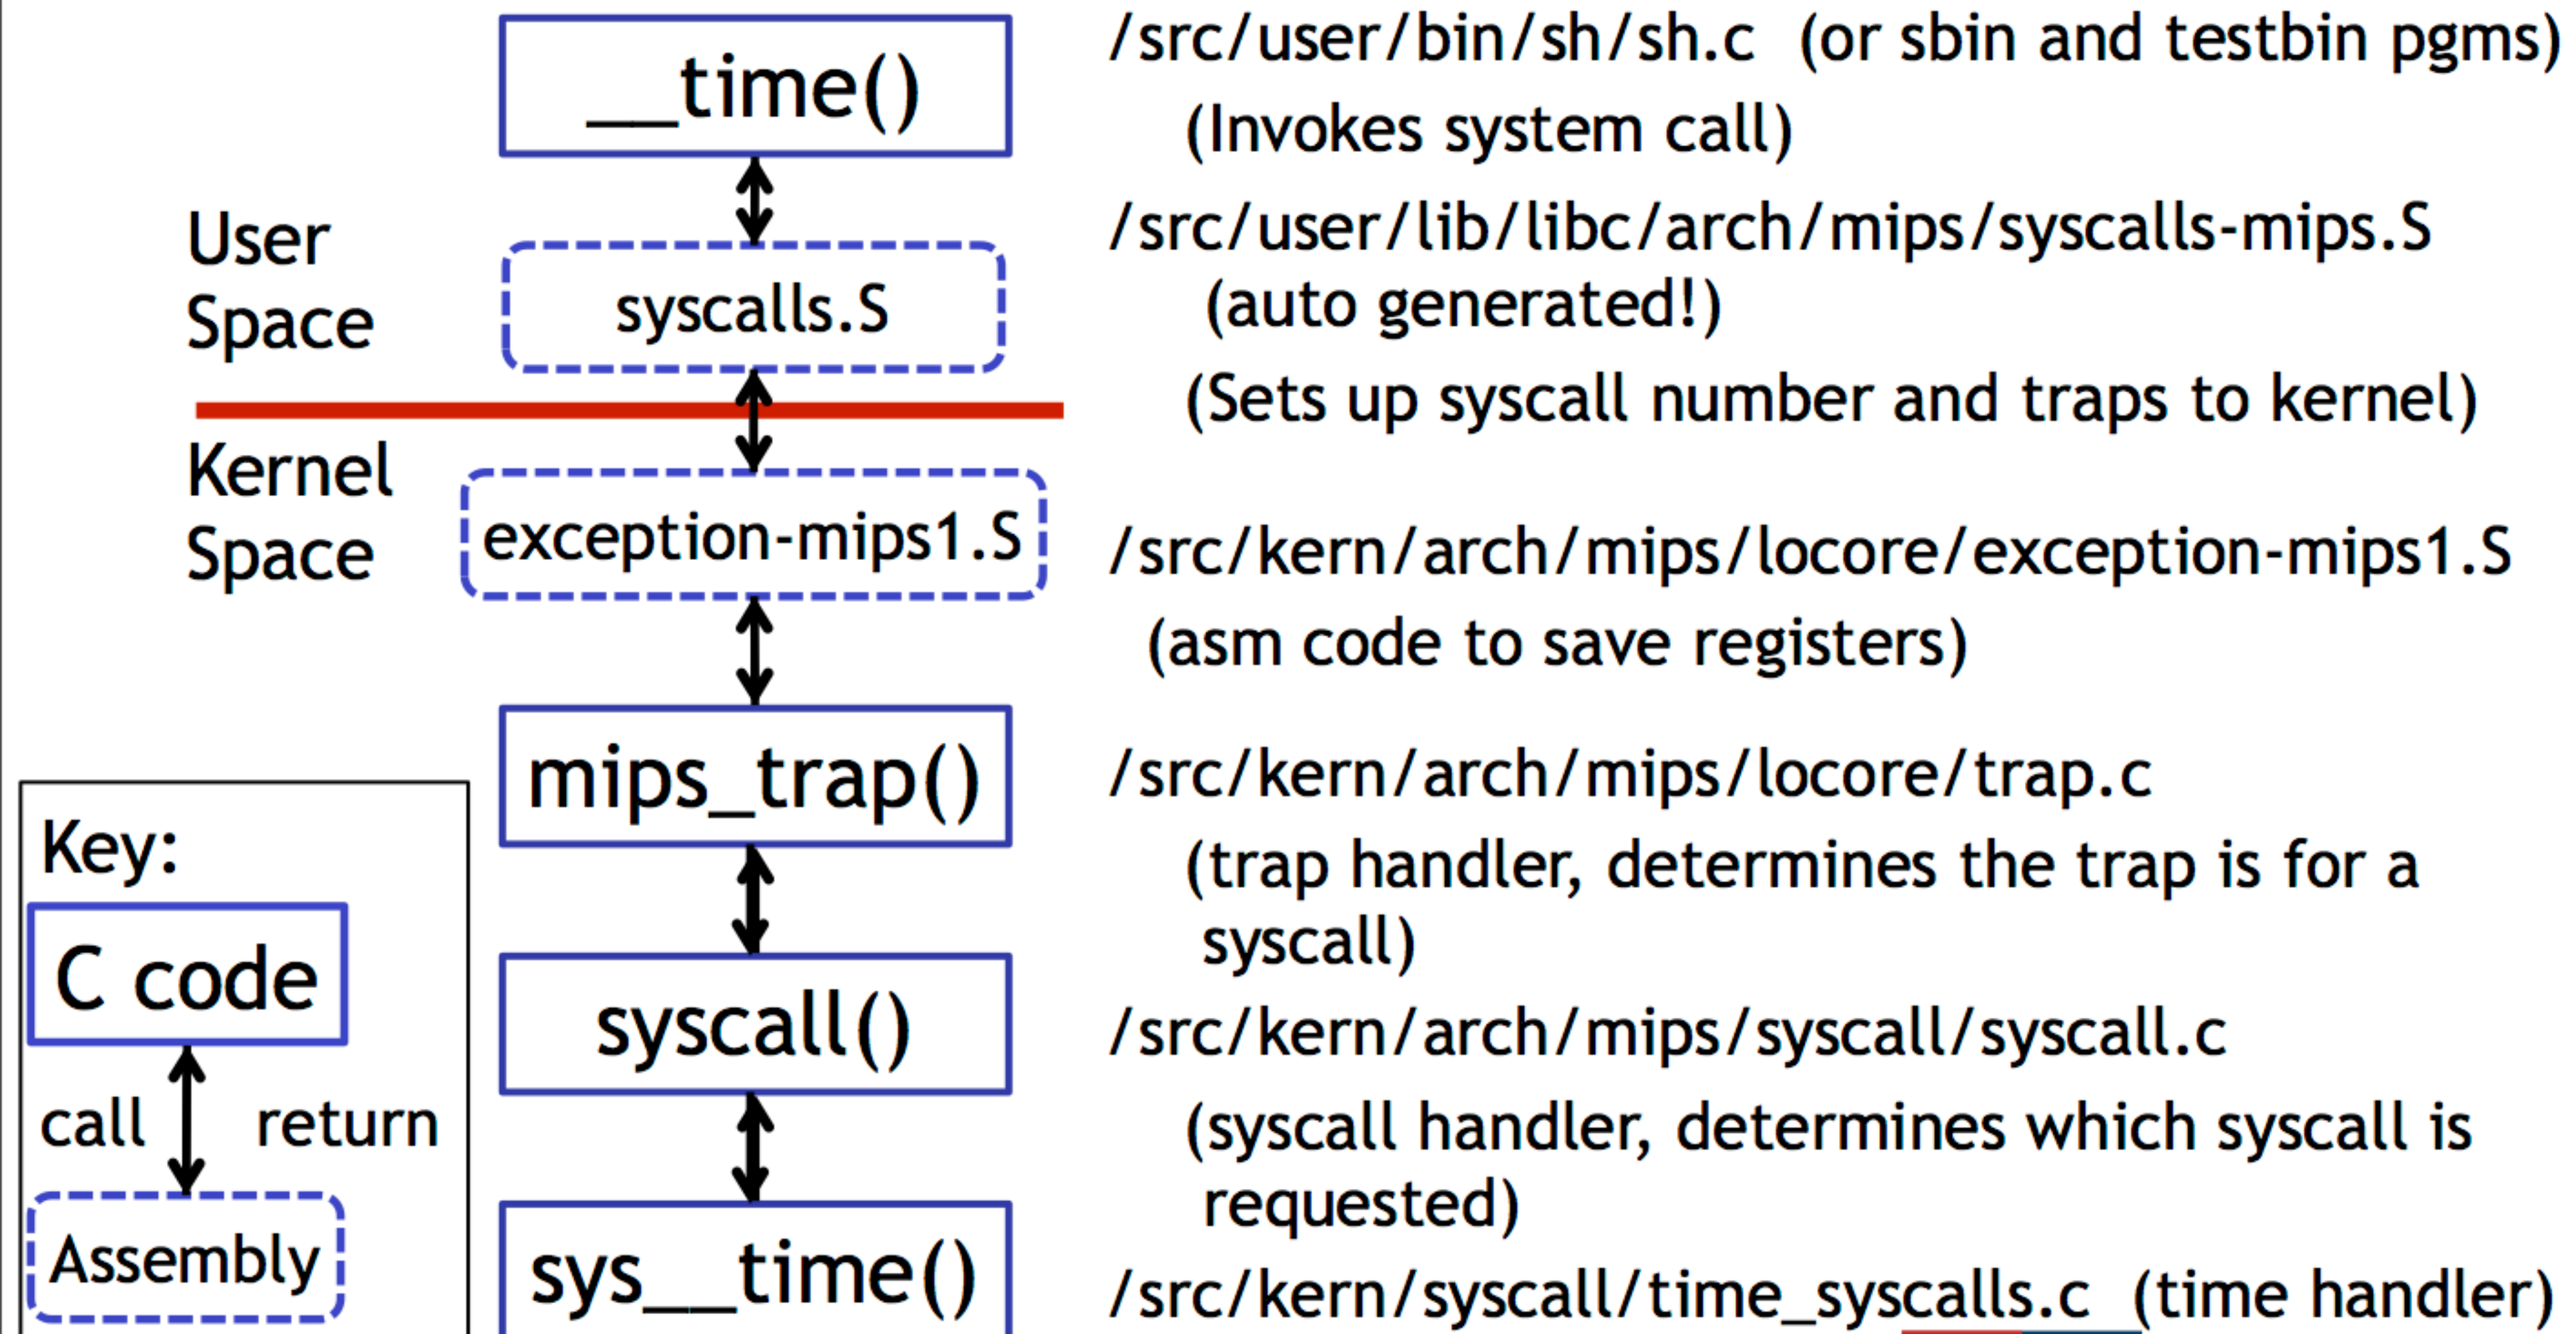
Summary

The challenge: **efficiently** virtualize CPU with **control**.

Mechanism: **time sharing** with **limited direct execution**.

- Process can perform restrict operations without messing around with hardware
 - ◆ System calls
- OS can switch from one process to another
 - ◆ Cooperative vs noncooperative
 - ◆ Timer interrupt
 - ◆ Context switch

OS161 System Call Example: time



Creating a New System Call in OS161

- Define the new system call code in `src/kern/include/kern/syscall.h`
- Define the prototypes for the new syscall
 - Kernel space: `src/kern/include/syscall.h`
 - User space: `src/user/include/unistd.h`
- Write the source code for it in `src/kern/syscall/new_syscall.c`
 - Be sure to include this new file in `src/kern/conf/conf.kern!` (so it's included in the build path)
- If necessary, define any new error codes in `src/kern/include/kern/errno.h`
- Add a case in the handler switch statement in `src/kern/arch/mips/syscall/syscall.c`
- Create a test program in `src/testbin`
- Rebuild both kernel and user level programs