

OS 161 - Adding System Calls

(Detailed Version)

Overview

1. Kernel-level steps
2. User-level steps
3. Testing the new system call

Kernel-level steps

1. Add the prototype of the system-call function to the header file: `kern/include/syscall.h`
2. The kernel-level implementation (e.g., `newsyscall.c`) goes into `kern/syscall/`
3. Add a new ID number for the system call. The new entry goes in the file `kern/include/kern/syscall.h`
4. Add a new branch in the switch-case statement in: `kern/arch/mips/syscall/syscall.c`
5. Add *file entry* definition for `syscall/newsyscall.c` in `kern/conf/conf.kern`

User-level steps

1. Add the user-level prototype of the system call to: `user/include/unistd.h`
2. Add the user-level test function. For this, create a new subdirectory directory `user/testbin/testnewsyscall/` and inside it add the test function (e.g., `testnewsyscall.c`).
3. Create a *Makefile* inside this subdirectory for building the test function. You can use one of the subdirectories as a template.
4. Add an entry to the new function to the top-level *Makefile* in `user/testbin`

Testing the new system call

1. Re-build the kernel
2. Start the new kernel (i.e., run `sys161 kernel` in the root directory)
3. At the OS161 prompt, use the *p option* (from OS161 menu) to run the test program, i.e., `p testbin/testnewsyscall`

1. **Kernel-level steps**


2. User-level steps

3. Testing the new system call

1 Prototype of the system call

- Add the prototype of the system call to the header file: `kern/include/syscall.h`
- In the end of the file, you will find prototypes for `sys_reboot()` and `sys__time()`.

```
53
54  /*
55   * Prototypes for IN-KERNEL entry points for system call
56   * implementations.
57   */
58  int sys_reboot(int code);
59  int sys__time(userptr_t user_seconds, userptr_t user_nanoseconds);
60
61  #endif /* _SYSCALL_H_ */
62
```



2 Kernel-level implementation

- The kernel-level implementation goes into `kern/syscall`. This directory contains an example of a system call, i.e., `time_syscalls.c`.
- Here, create a program called `simple_syscall.c`, and implement your system call in it.

```
int sys_helloworld(void){  
    return kprintf("Hello World!\n");  
}
```

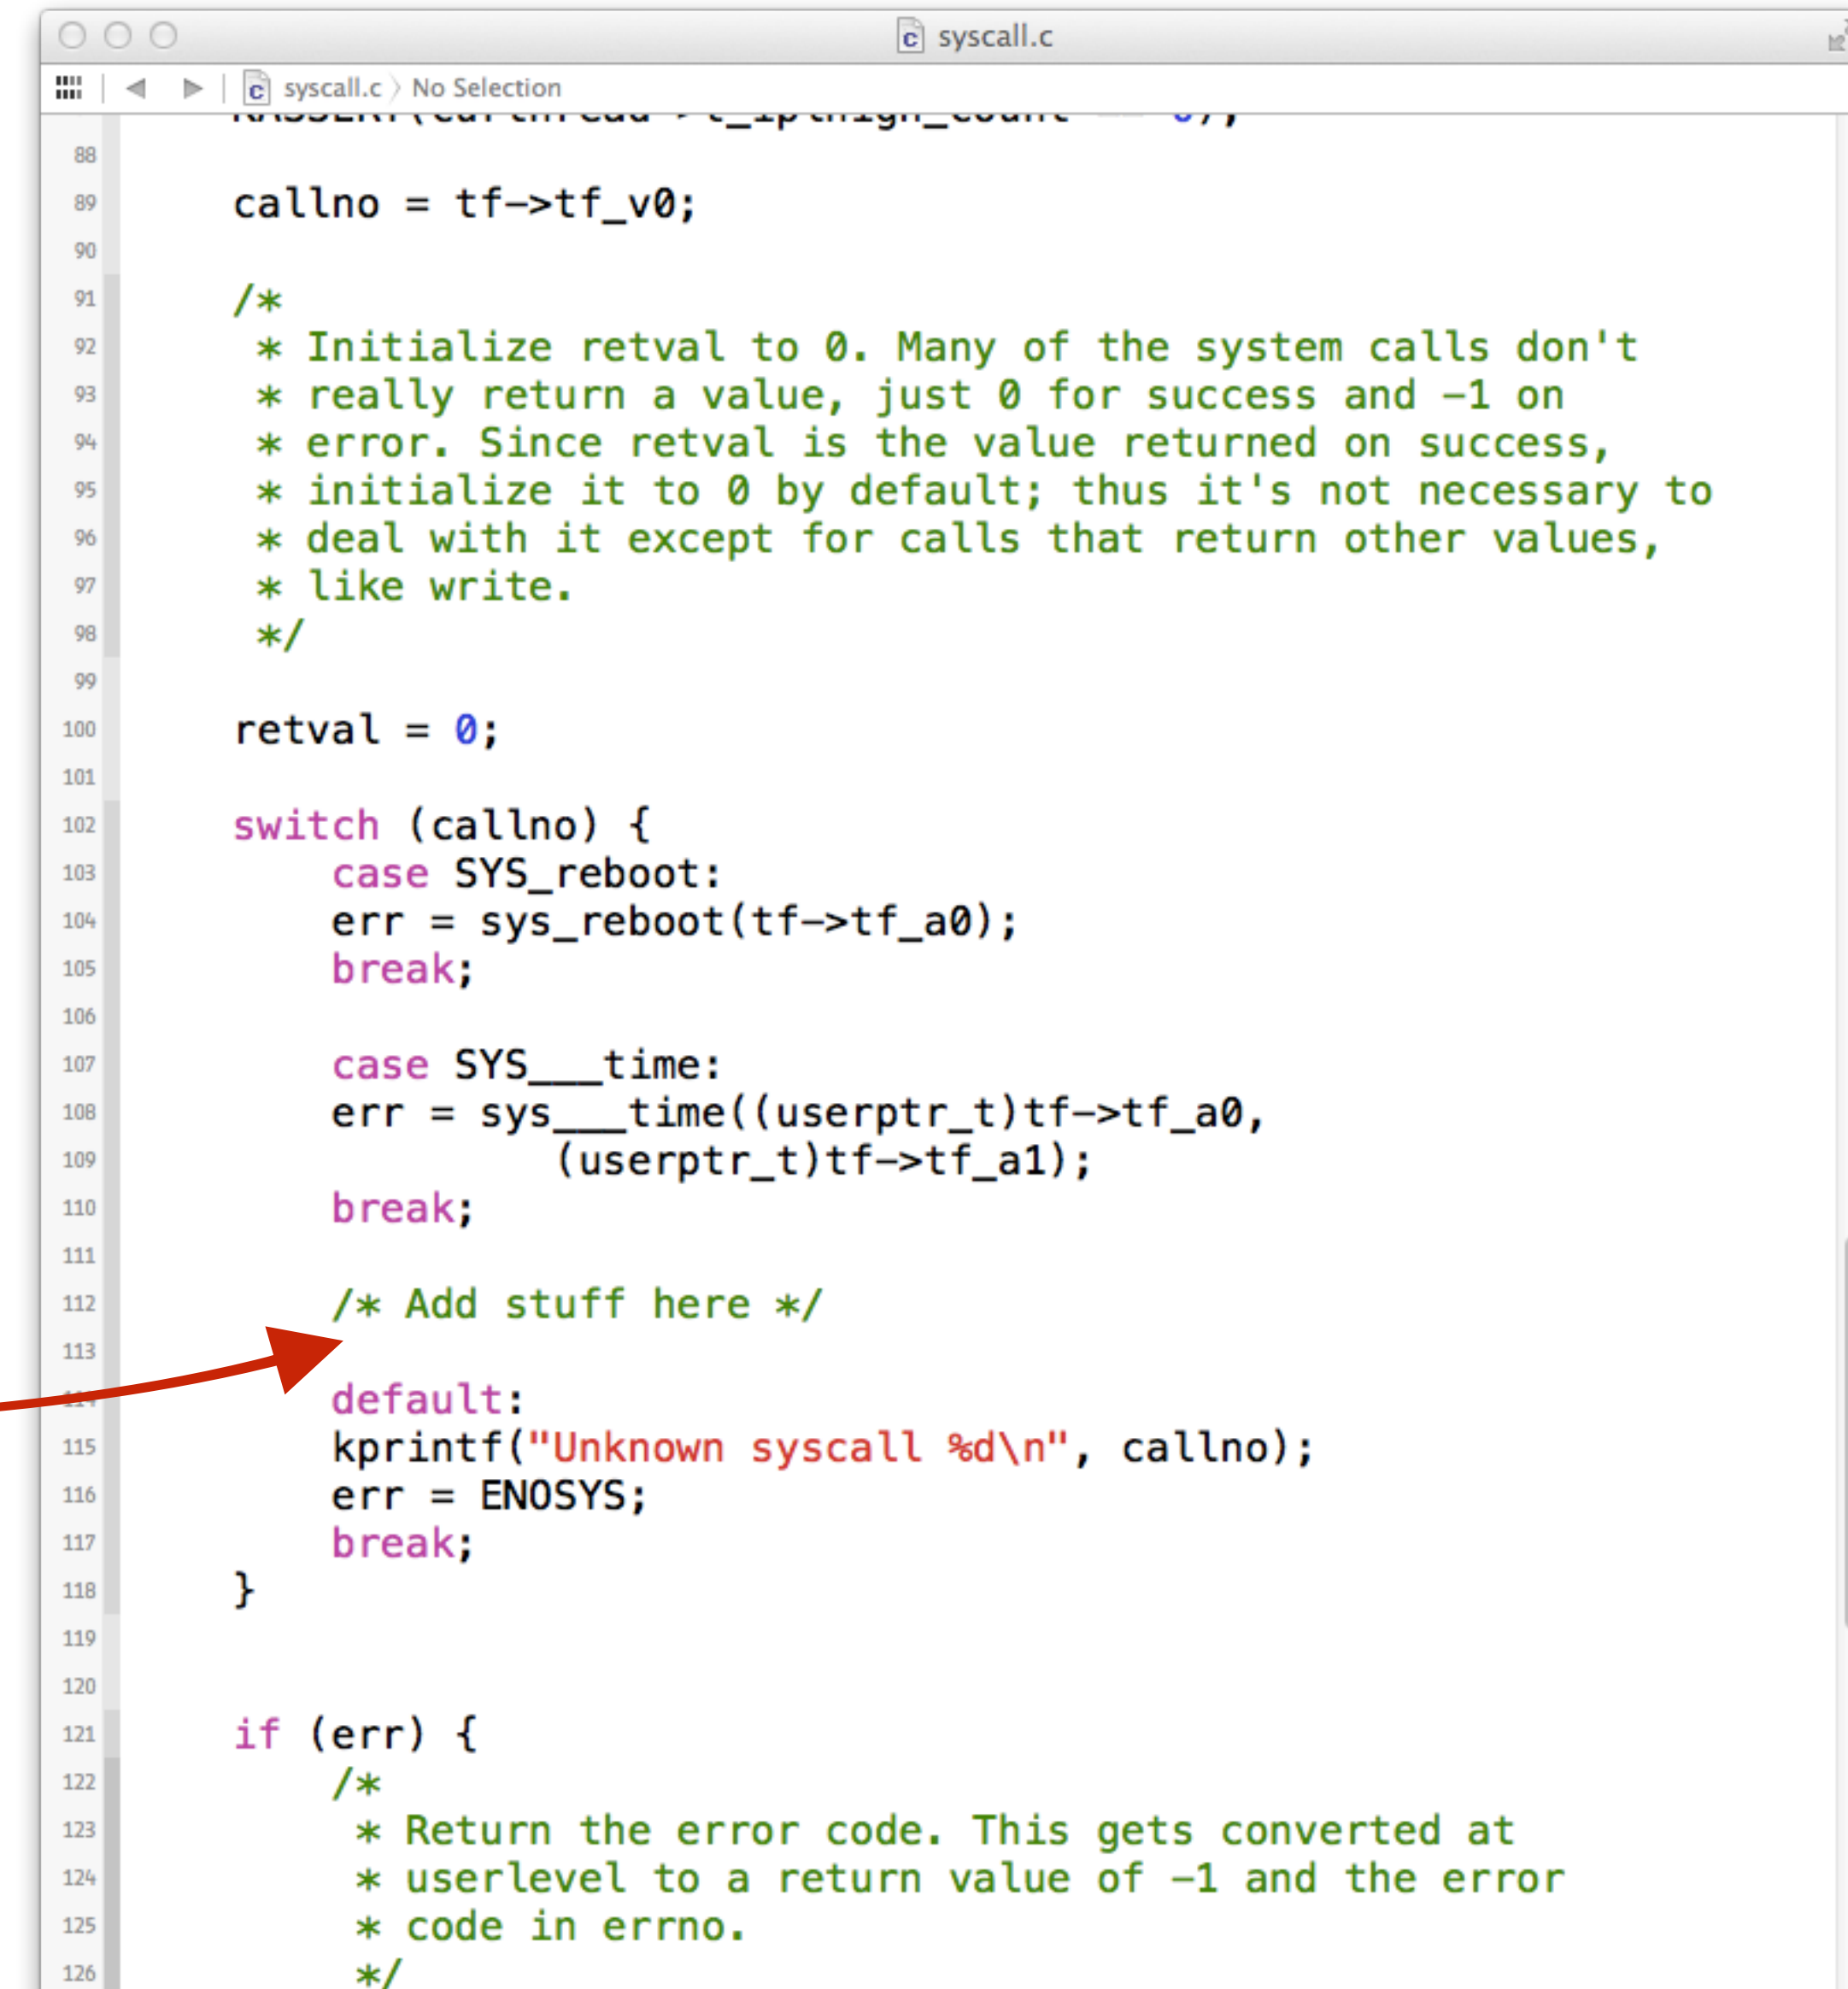

3 Create the ID number for the new system call

- The OS needs to know the ID number of the system call
- Add a new entry to the file `kern/include/kern/syscall.h`

```
98 // #define SYS_setpgid      41
99 // #define SYS_getsid       42
100 // #define SYS_setsid       43
101 //                               (userlevel debugging)
102 // #define SYS_ptrace        44
103
104 //                               -- File-handle-related --
105
106
107 #define SYS_open             45
108 #define SYS_pipe             46
109 #define SYS_dup              47
110 #define SYS_dup2             48
```

4 Add a new branch in the switch-case statement in:
`kern/arch/mips/syscall/syscall.c`

```
case SYS_helloworld:  
    err = sys_helloworld();  
    break;
```



```
88  
89     callno = tf->tf_v0;  
90  
91     /*  
92      * Initialize retval to 0. Many of the system calls don't  
93      * really return a value, just 0 for success and -1 on  
94      * error. Since retval is the value returned on success,  
95      * initialize it to 0 by default; thus it's not necessary to  
96      * deal with it except for calls that return other values,  
97      * like write.  
98      */  
99  
100    retval = 0;  
101  
102    switch (callno) {  
103        case SYS_reboot:  
104            err = sys_reboot(tf->tf_a0);  
105            break;  
106  
107        case SYS___time:  
108            err = sys___time((userptr_t)tf->tf_a0,  
109                            (userptr_t)tf->tf_a1);  
110            break;  
111  
112        /* Add stuff here */  
113        default:  
114            kprintf("Unknown syscall %d\n", callno);  
115            err = ENOSYS;  
116            break;  
117    }  
118  
119  
120  
121    if (err) {  
122        /*  
123         * Return the error code. This gets converted at  
124         * userlevel to a return value of -1 and the error  
125         * code in errno.  
126         */
```

5 Add file-entry definition to config.kern

```
358
359 file      vfs/devnull.c
360
361 #
362 # System call layer
363 # (You will probably want to add stuff here while doing the basic system
364 # calls assignment.)
365 #
366
367 file      syscall/loadelf.c
368 file      syscall/runprogram.c
369 file      syscall/time_syscalls.c
370
371 #
372 # Startup and initialization
373 #
374
375 file      startup/main.c
376 file      startup/menu.c
377
378 #####
379 #                                     #
380 #             Filesystems            #
381 #                                     #
382 #####
```

1. Kernel-level steps
- 2. User-level steps**
3. Testing the new system call

1. Add the user-level prototype of the system call to:
`user/include/unistd.h`

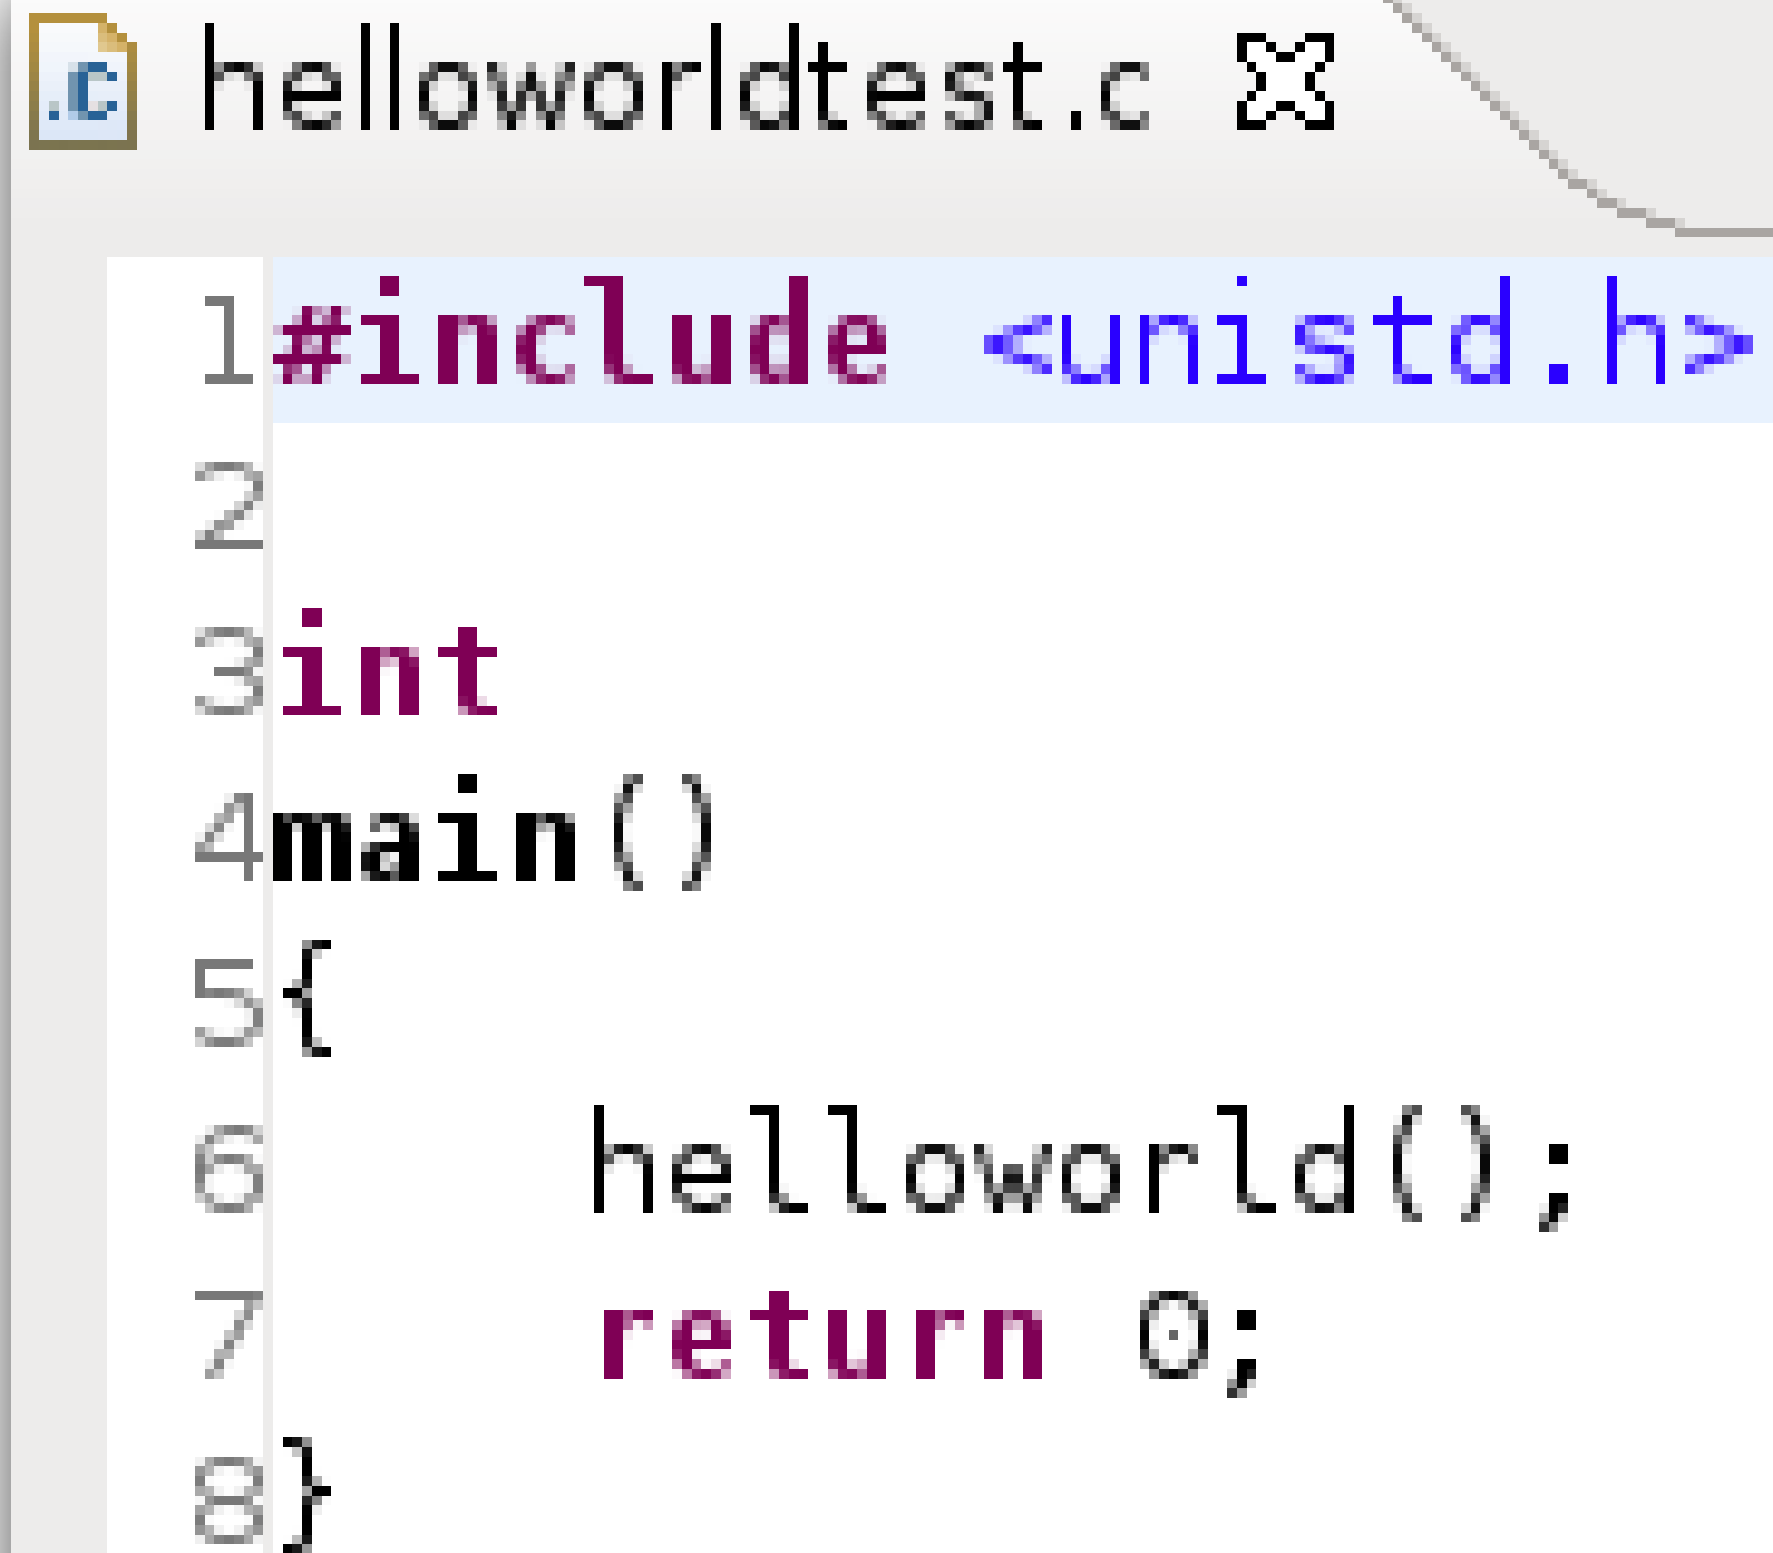
a sense
because,
these

```
109
110 #ifdef __GNUC__
111 /* GCC gets into a snit if _exit isn't declared to not return */
112 #define __DEAD __attribute__((__noreturn__))
113 #else
114 #define __DEAD
115 #endif
116
117 /* Required. */
118 __DEAD void _exit(int code);
119 int execv(const char *prog, char *const *args);
120 pid_t fork(void);
121 int waitpid(pid_t pid, int *returncode, int flags);
122 /*
123  * Open actually takes either two or three args: the optional third
124  * arg is the file
125  * security and p
126  */
127 int open(const char *pathname, int flags, mode_t mode);
128 int read(int filedes, void *buf, size_t size);
129 int write(int filehandle, const void *buf, size_t size);
130 int close(int filehandle);
131 int reboot(int code);
132 int sync(void);
133 /* mkdir - see sys/stat.h */
134 int mkdir(const char *dirname, mode_t mode);
135
136 /* Recommended. */
137 int getpid(void);
138 int ioctl(int filehandle, int code, void *buf);
```



```
int helloworld();
int printchar(char c);
```

2. Add the user-level test function. For this, create a new subdirectory directory `user/testbin/`



A code editor window titled "helloworldtest.c" with a C file icon. The code is as follows:

```
1#include <unistd.h>
2
3int
4main()
5{
6    helloworld();
7    return 0;
8}
```

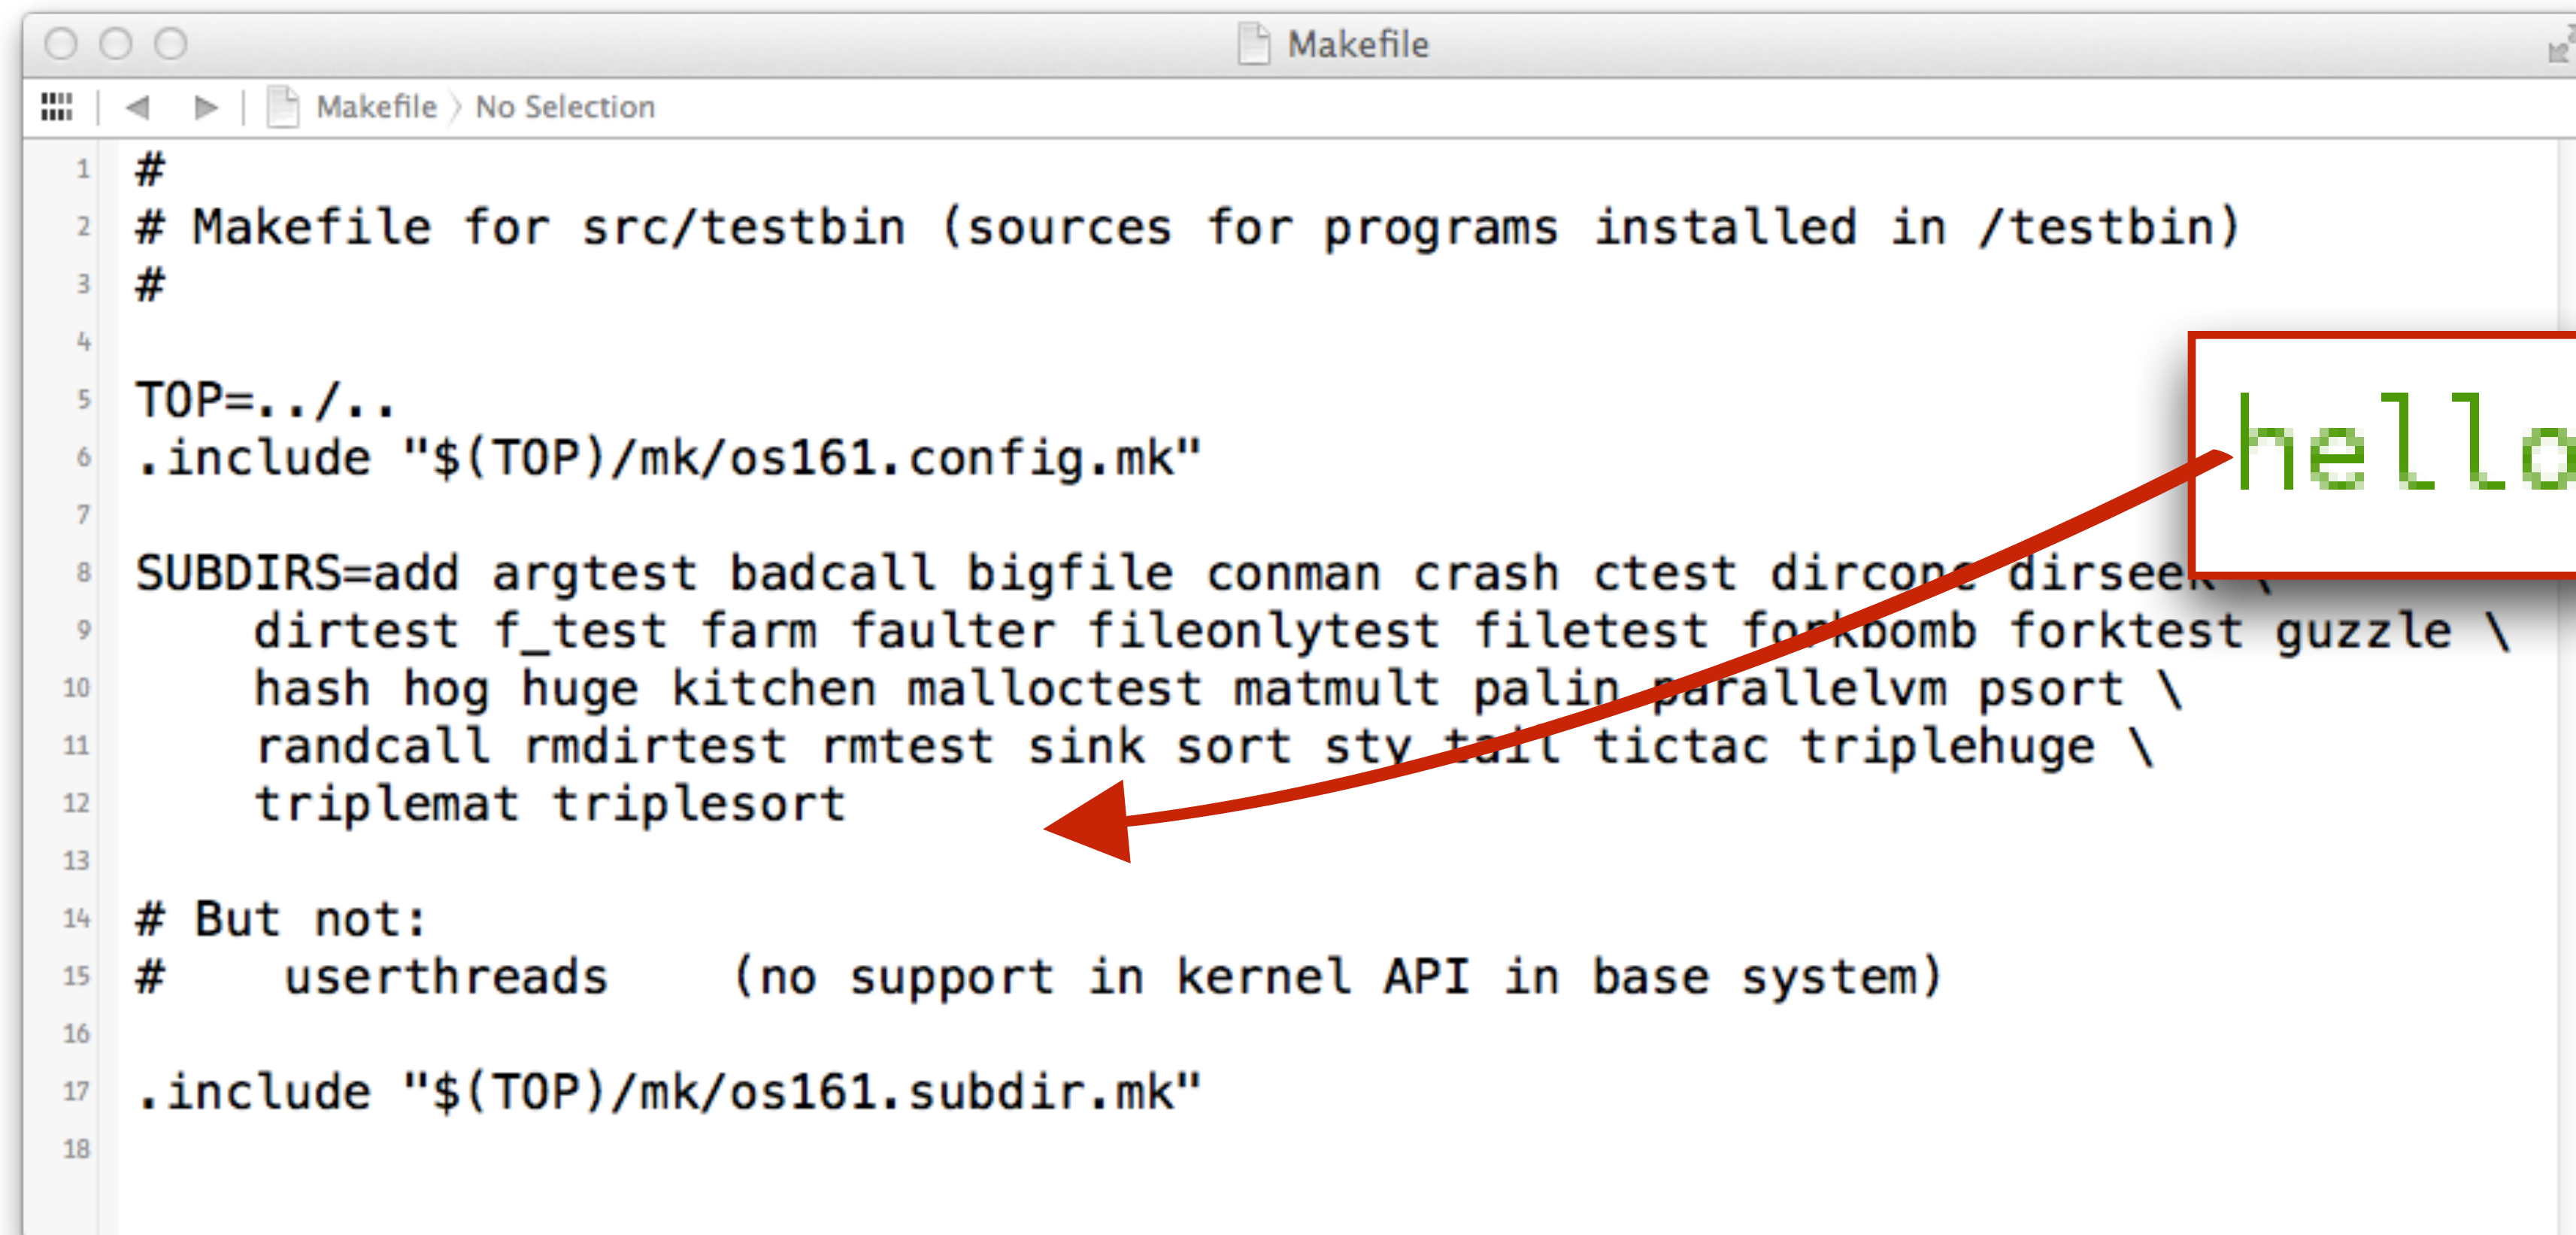
```
# Makefile for helloworldtest

TOP=../..
include "$(TOP)/mk/os161.config.mk"

PROG=helloworldtest
SRCS=helloworldtest.c
BINDIR=/testbin

include "$(TOP)/mk/os161.prog.mk"
```

3. Add an entry to the new function to the top-level *Makefile* in `user/testbin/` and inside it add the



```
1 #
2 # Makefile for src/testbin (sources for programs installed in /testbin)
3 #
4
5 TOP=../..
6 .include "$(TOP)/mk/os161.config.mk"
7
8 SUBDIRS=add argtest badcall bigfile conman crash ctest dircone dirseek \
9   dirtest f_test farm faulter fileonlytest filetest forkbomb forktest guzzle \
10  hash hog huge kitchen malloctest matmult palin parallelvm psort \
11  randcall rmdirtest rmtest sink sort sty tail tictac triplehuge \
12  triplemat triplesort
13
14 # But not:
15 #   userthreads      (no support in kernel API in base system)
16
17 .include "$(TOP)/mk/os161.subdir.mk"
18
```

helloworldtest

1. Kernel-level steps
2. User-level steps
- 3. Testing the new system call**

Testing User Programs:

- Inside the root folder, run the command “sys161 kernel”
- In the os161 terminal, run the command “p testbin/<name>” where name is the name of your program

Hellotest Program:

```
OS/161 kernel [? for menu]: p testbin/hellotest
Operation took 0.000145920 seconds
OS/161 kernel [? for menu]: syscall: #40, args 0 0 0 0
Hello World!
syscall: #3, args 0 0 0 0
Thread testbin/hellotest exiting due to 0 with value 0
```