

# Paging

CSE 4001

# Content

- Basic paging mechanism
- Limitations
- Protection
- Shared pages

# Paging

## **Basic problem with allocating contiguous blocks of memory for processes**


- Determining the size of memory blocks is difficult because different processes have different memory requirements.

**Paging:** physical address space is allowed to be non-contiguous

# Paging

## Basic paging method

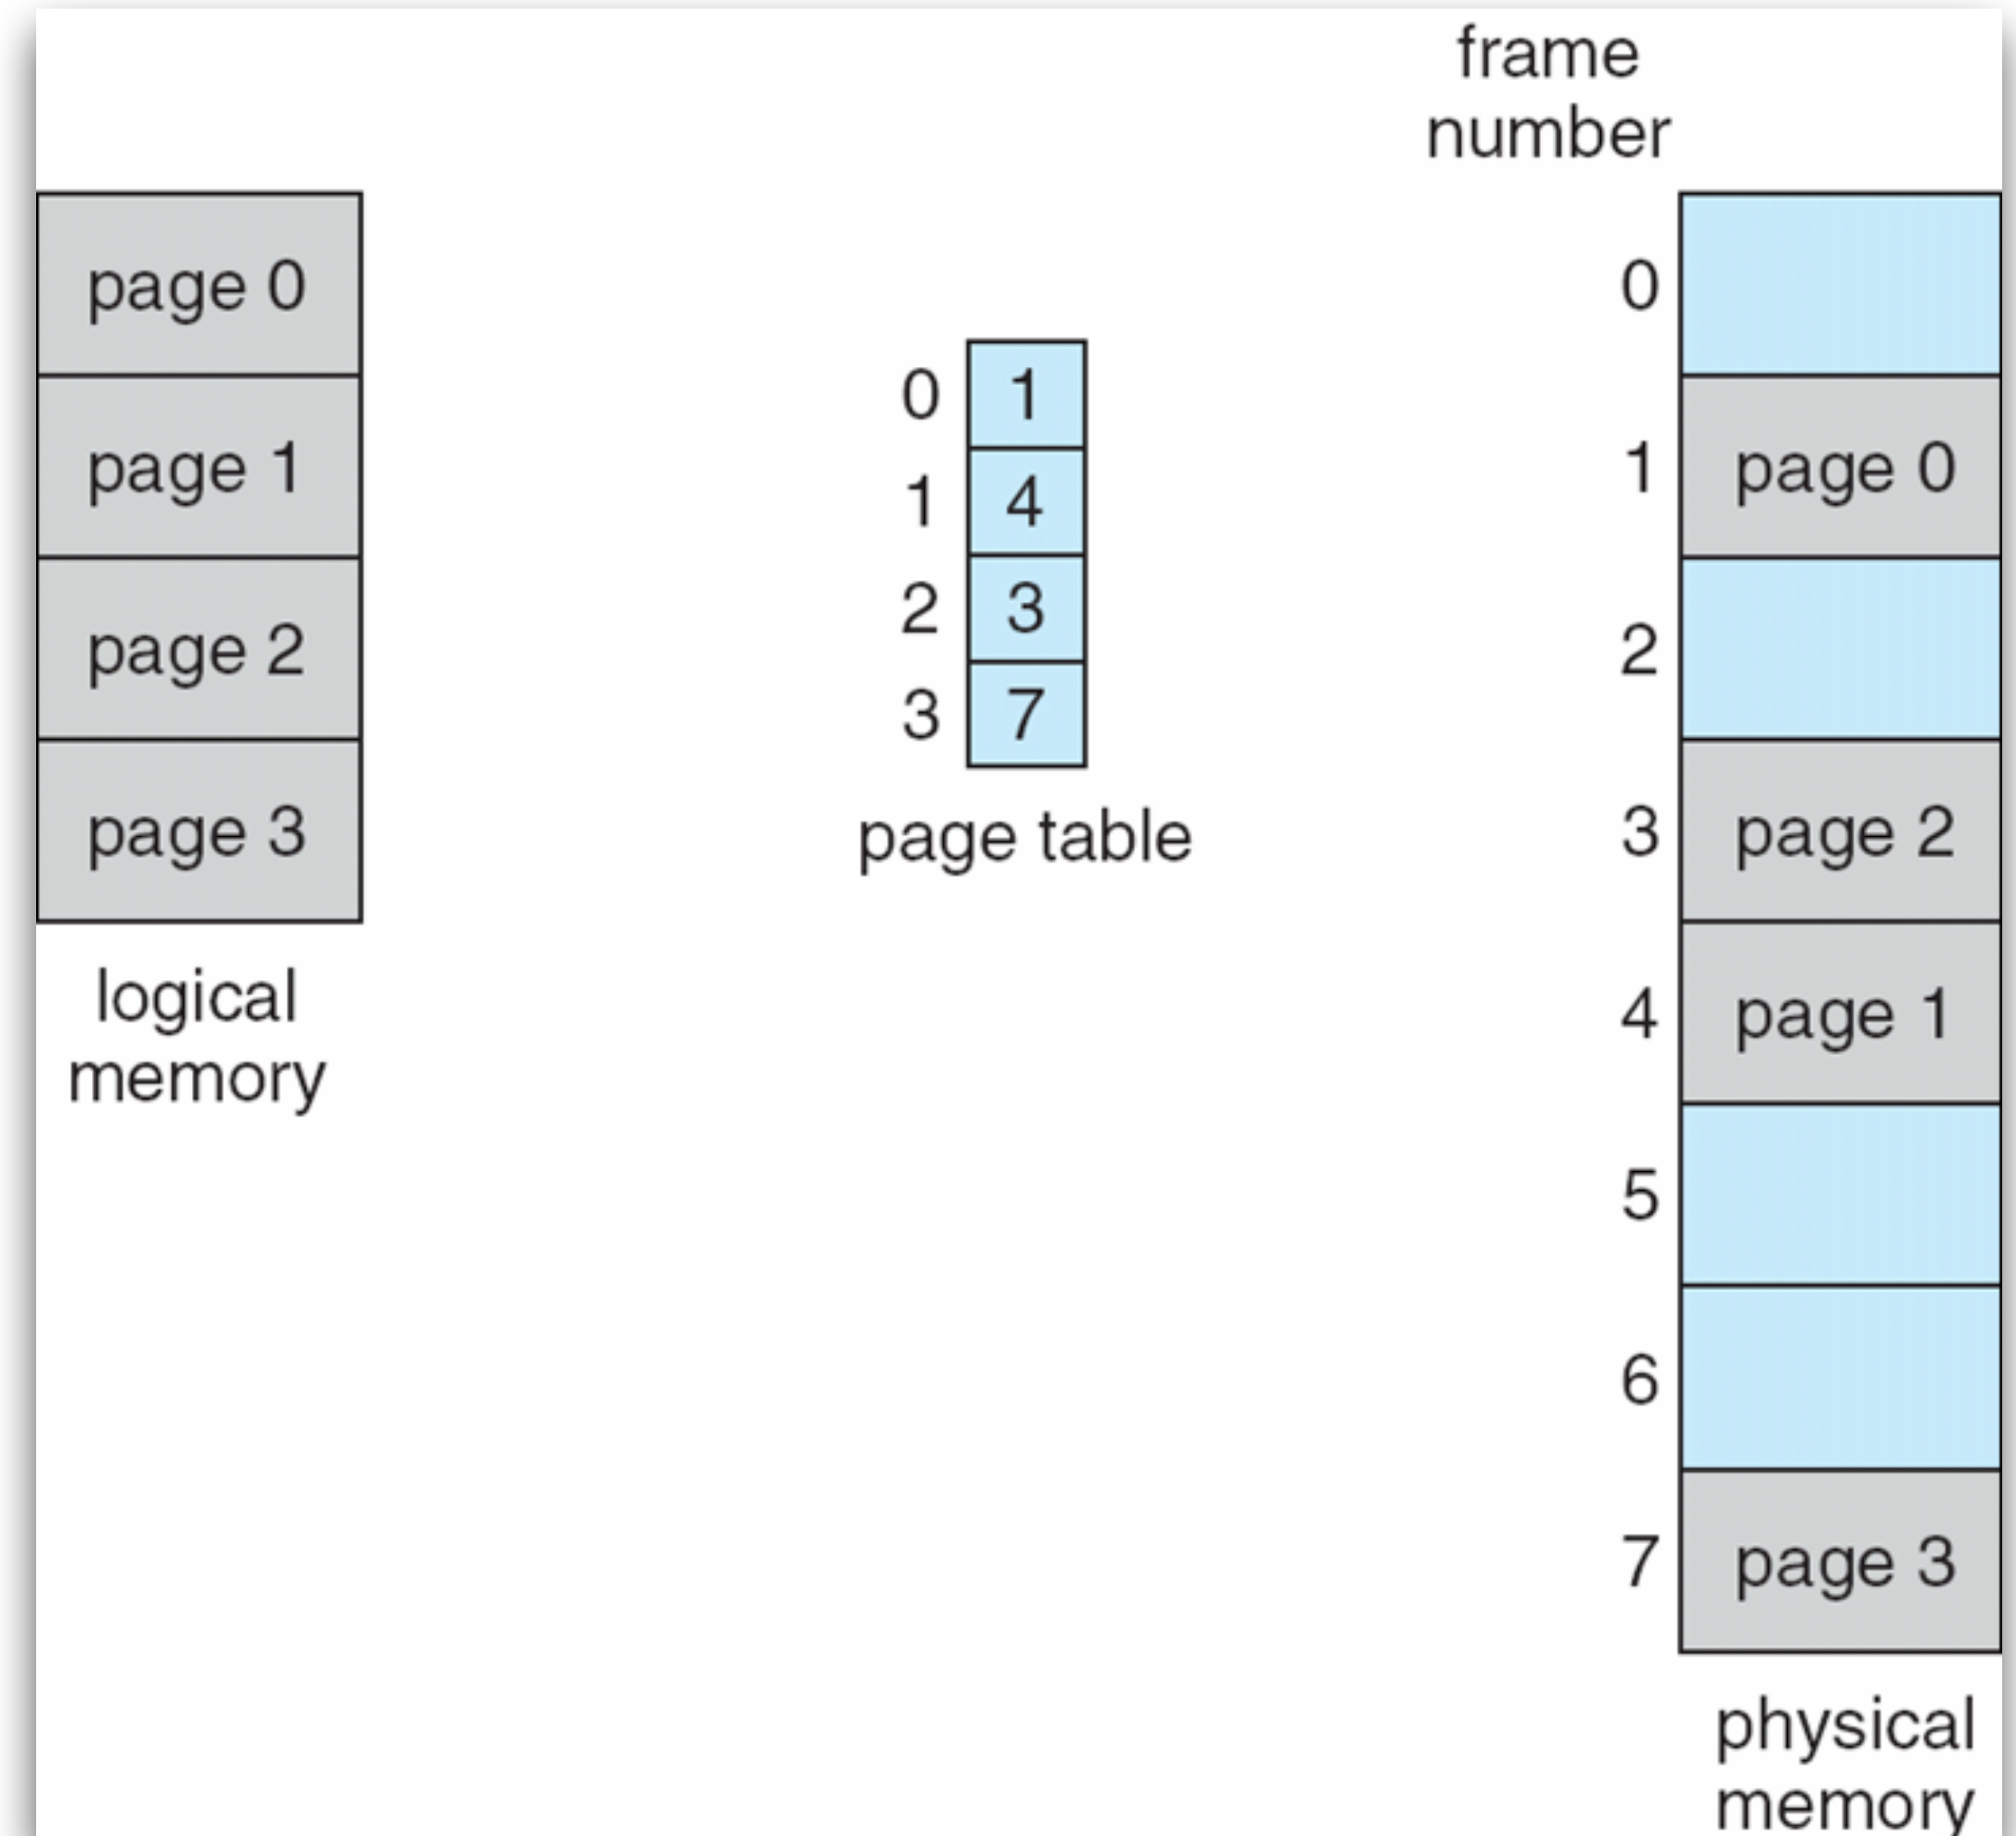
- Divide physical memory into fixed-sized blocks called **frames** (size is power of 2, between 512 bytes and 16 MB).
- Divide logical memory into blocks of same size called **pages**.




# Paging

## Basic paging method

- Any page can be assigned to any free page frame
- External fragmentation is eliminated
- Internal fragmentation is at most a part of one page per process




# Paging example: 32-byte memory and 4-byte pages

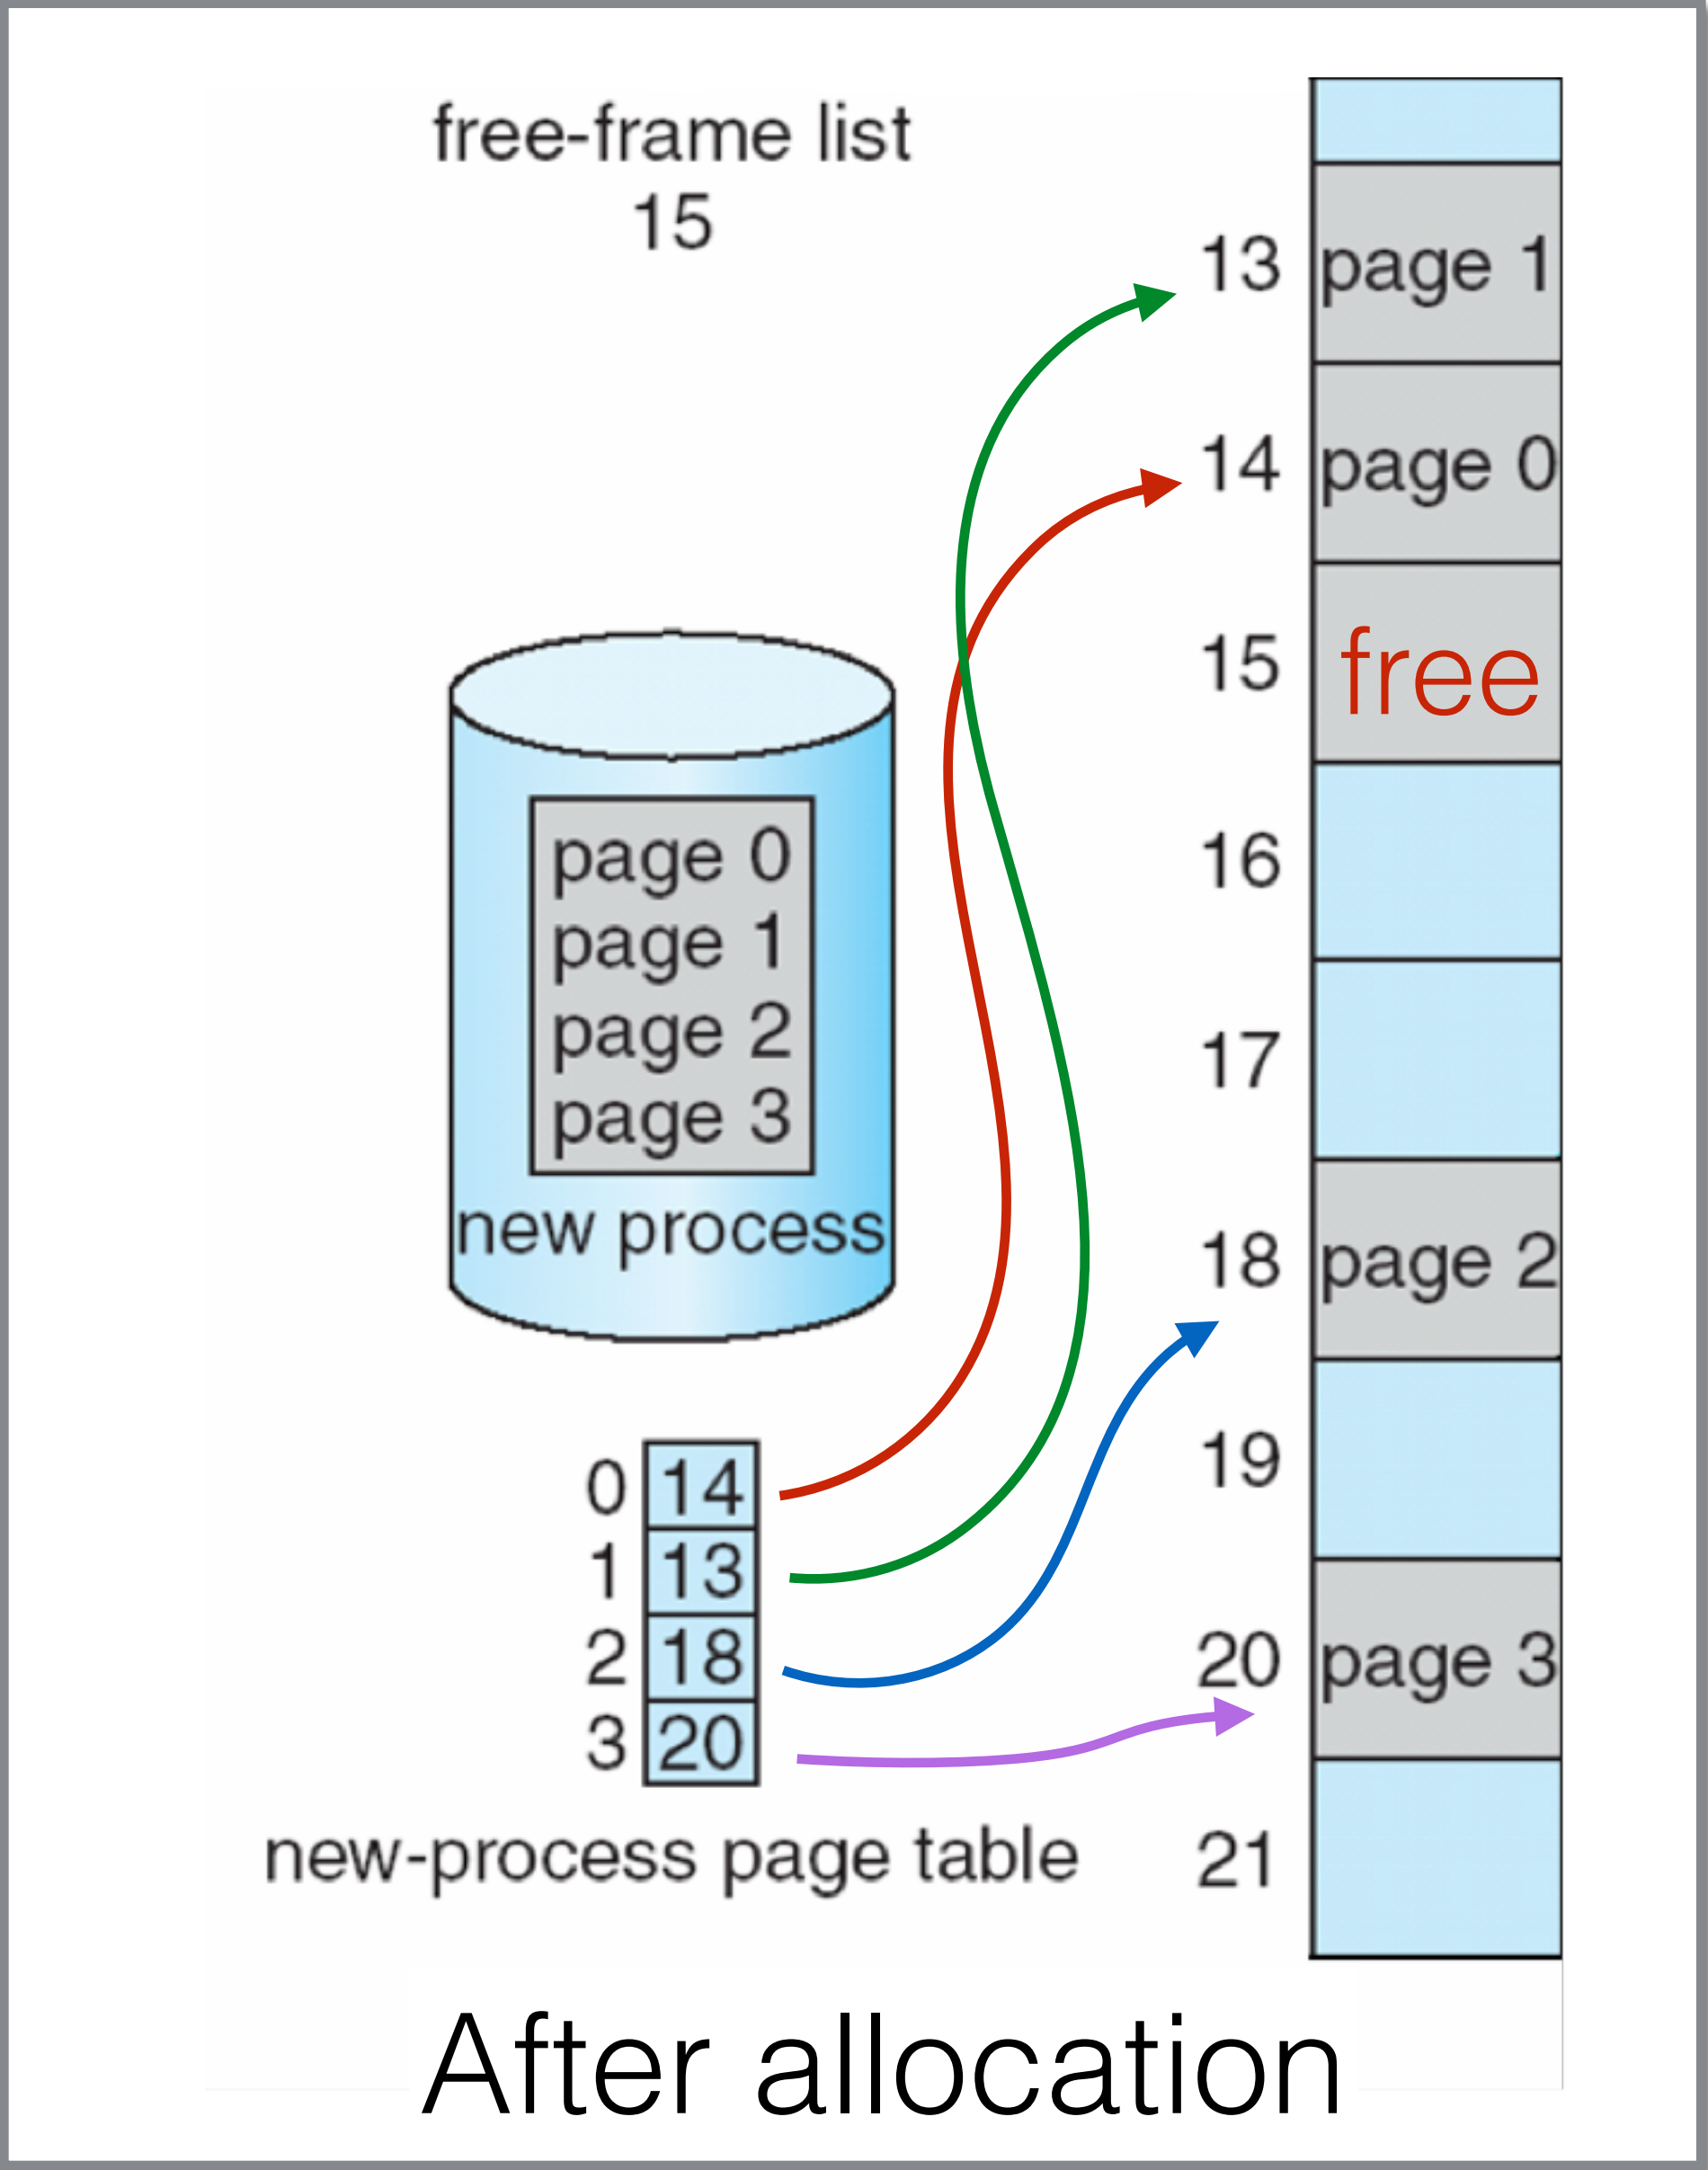
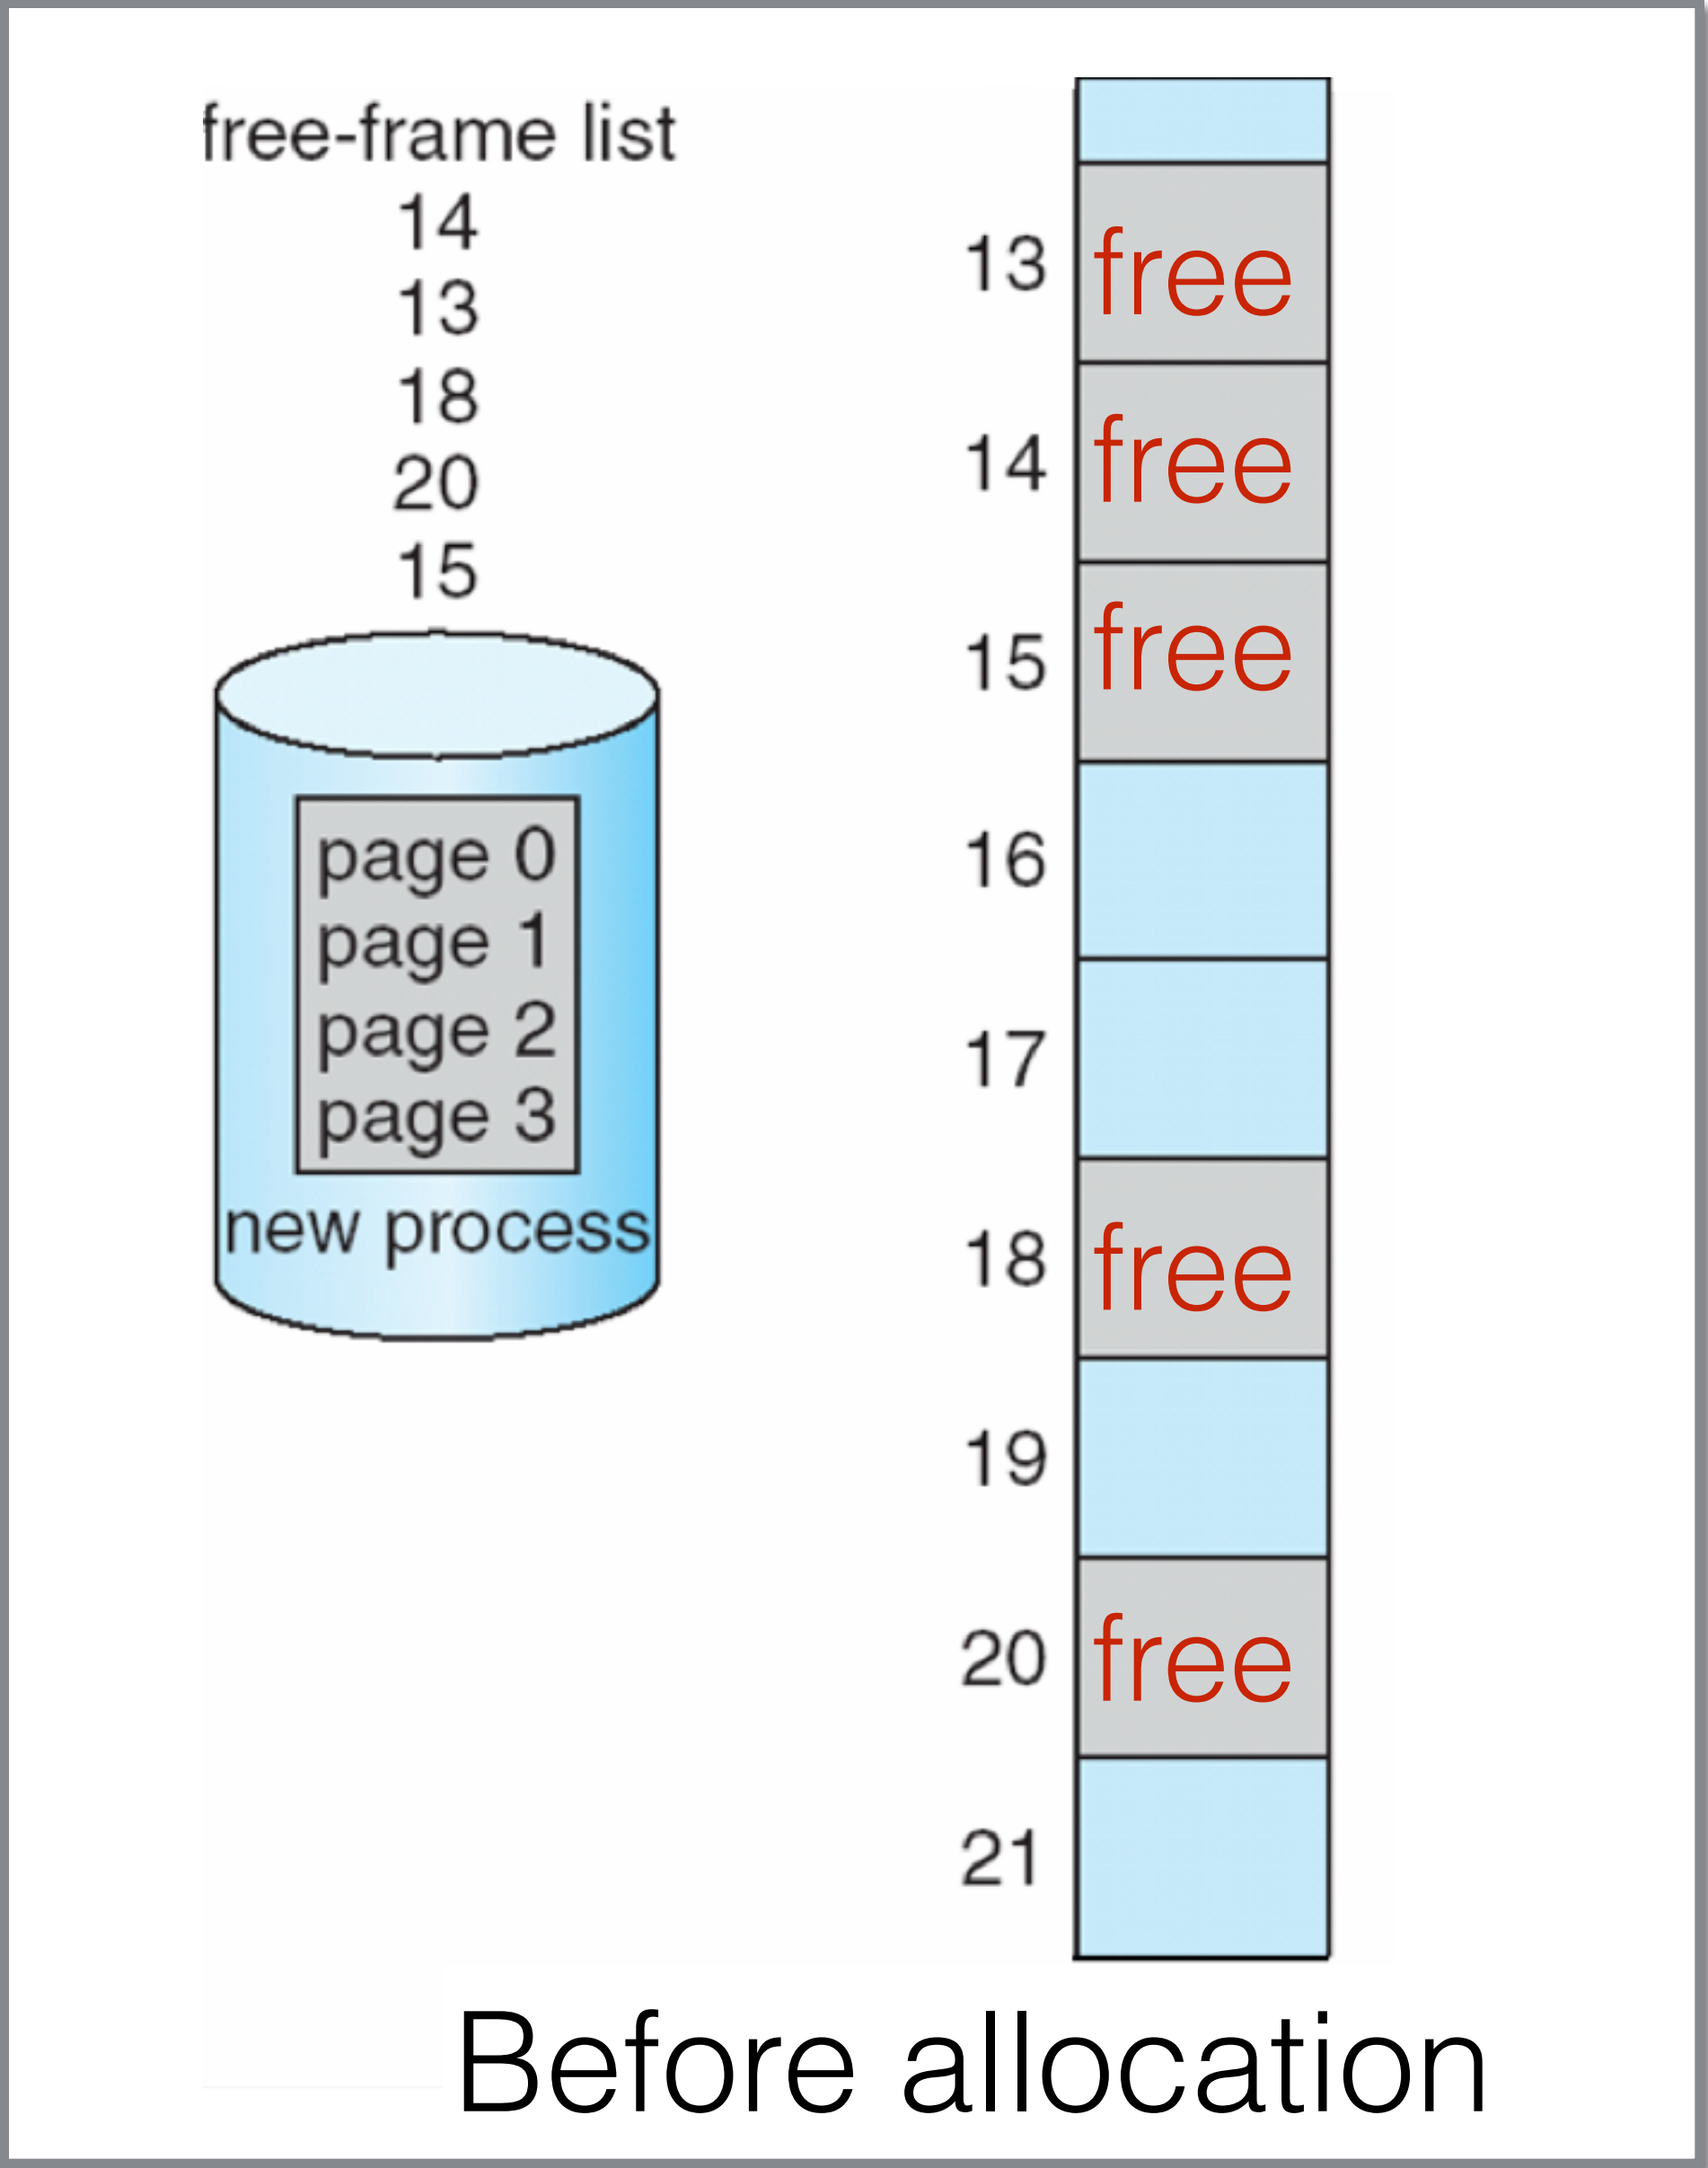




Paging example: 32-byte memory and 4-byte pages (with internal fragmentation)



# New process is executed: free frames before and after allocation






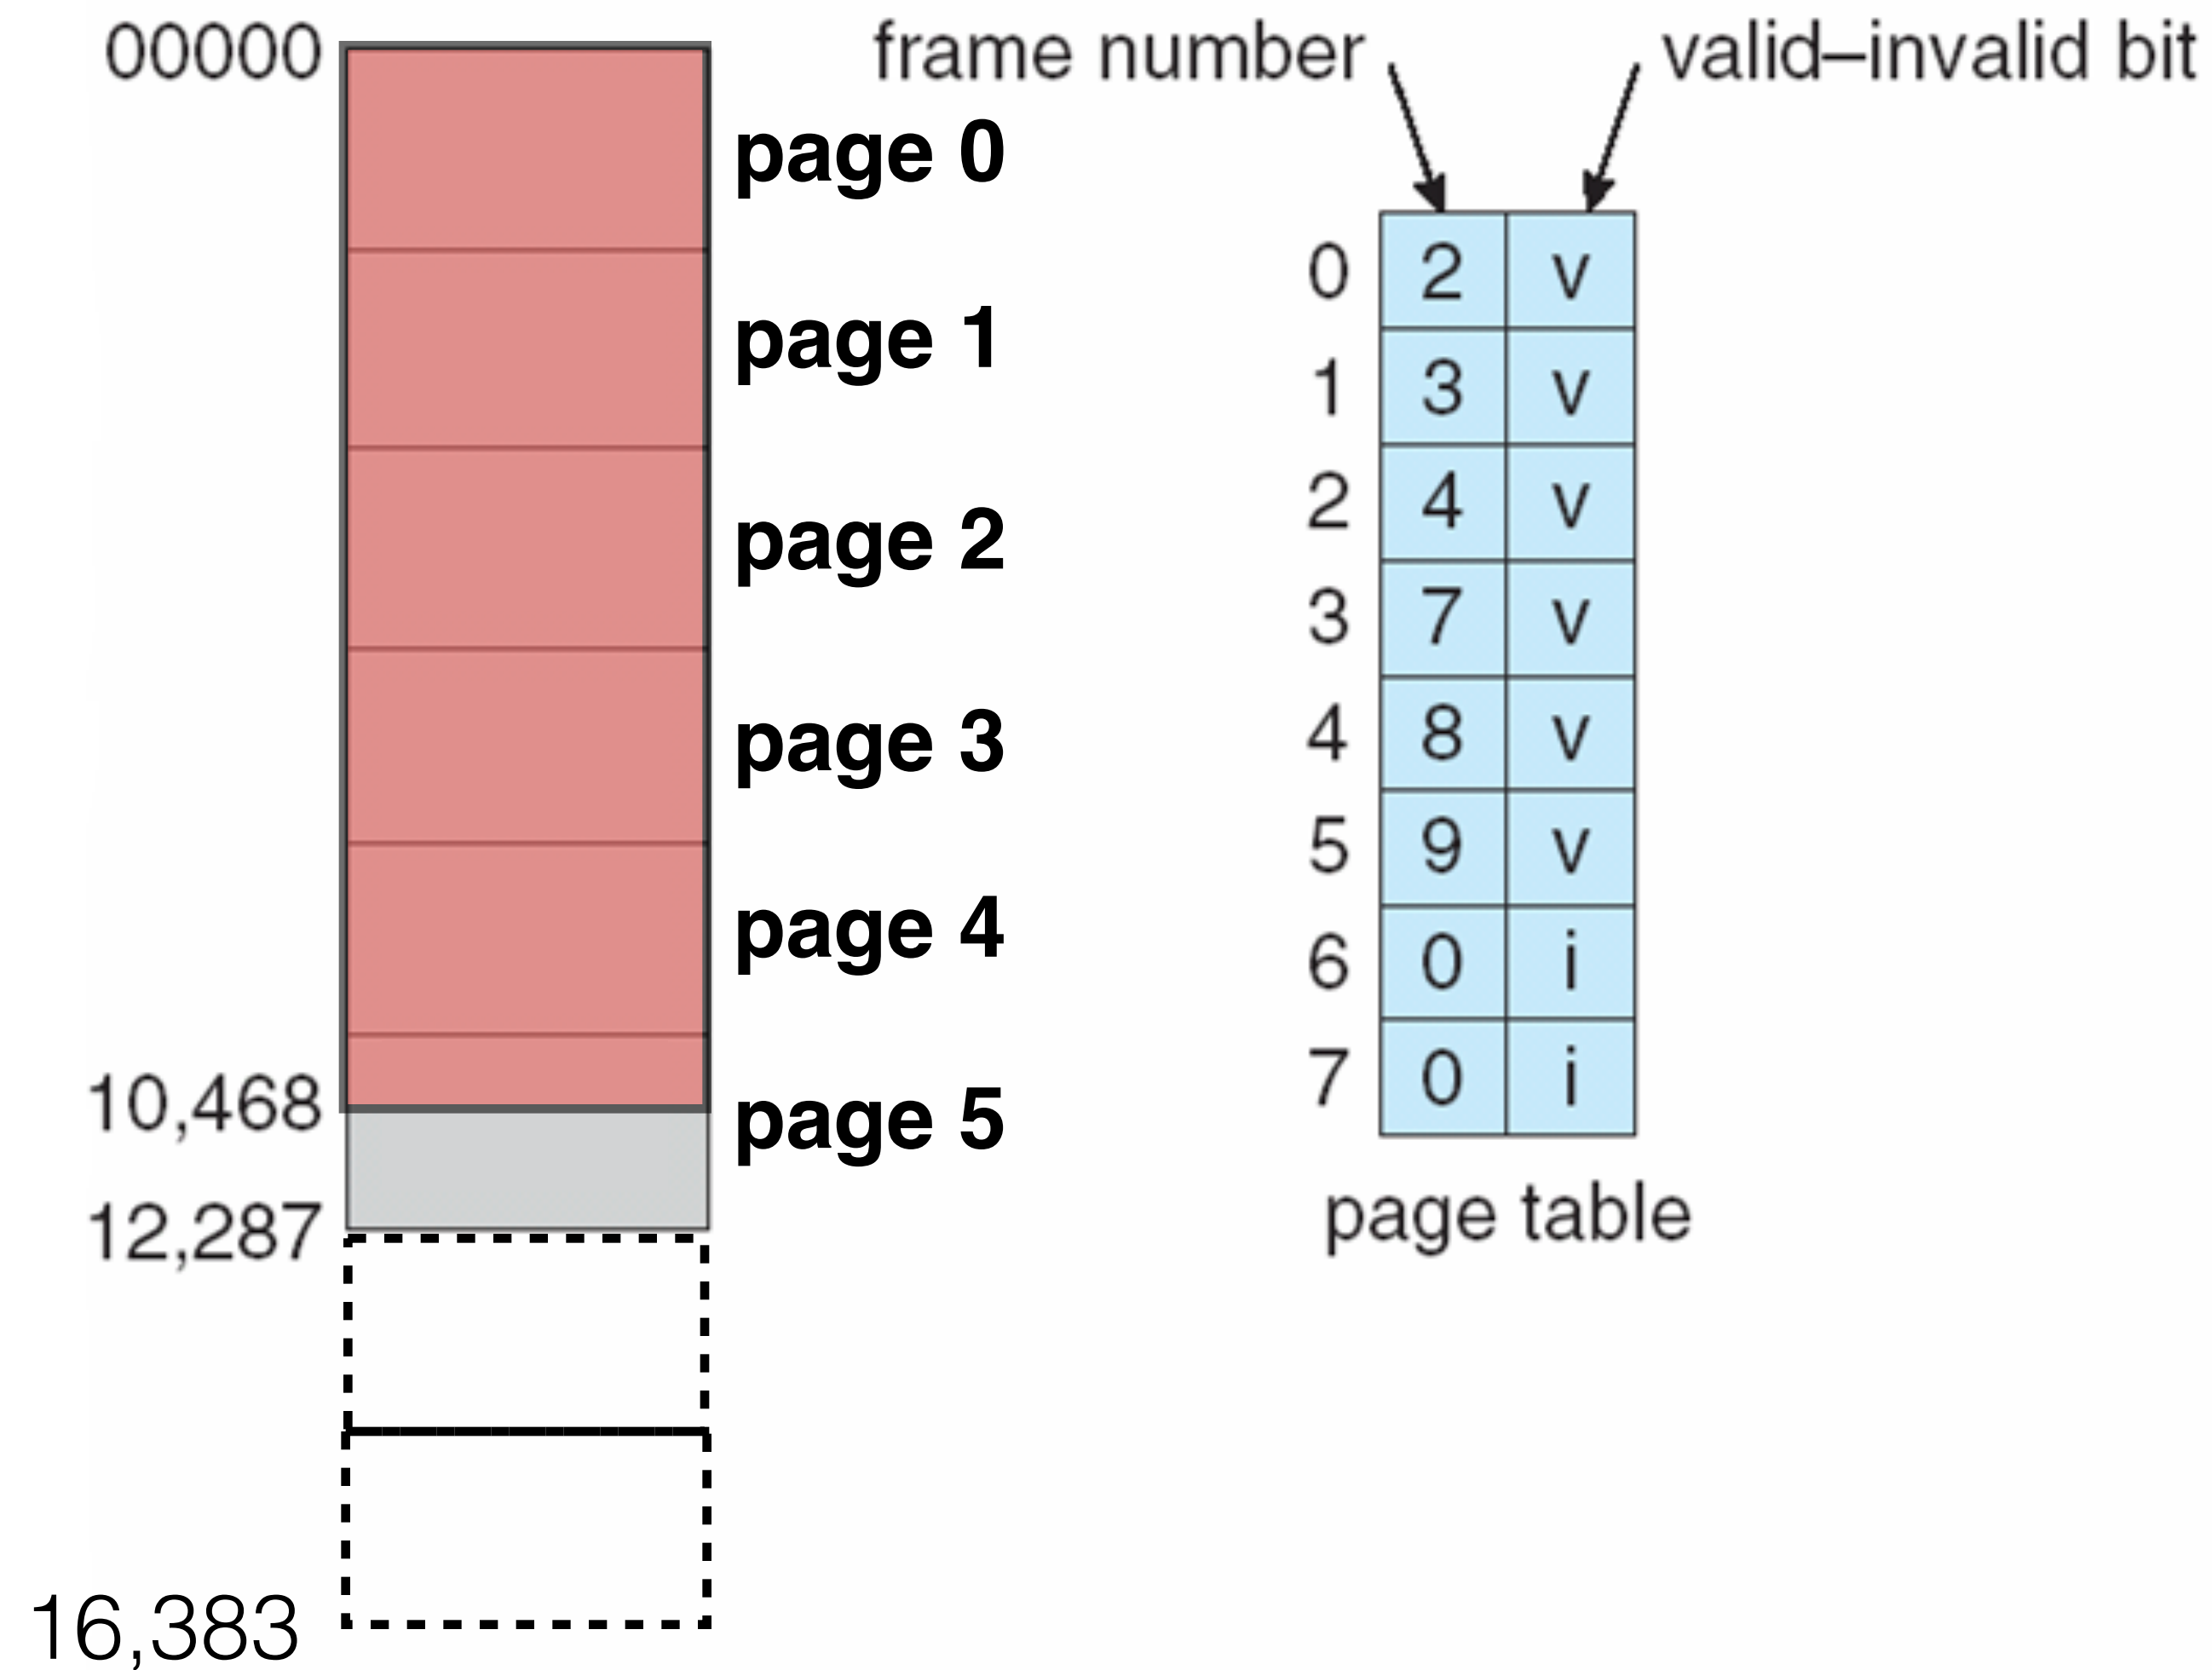
# Paging Limitations - Space

- Page table might need a lot of space
- **Registers** can be used to store page tables but they are only **feasible for small tables** (e.g., 256 entries).
- Modern computers have page tables of **1 million entries**.
- Such **large page tables are kept in main memory** and a page-table base register (PTBR) points to the table.

# Protection

- Memory protection: each frame has a **protection bit**.
- **Valid-invalid** bit for each entry in the page table:
  - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page.
  - “invalid” indicates that the page is not in the process’ logical address space.

# Protection



# A few more useful aspects of paging

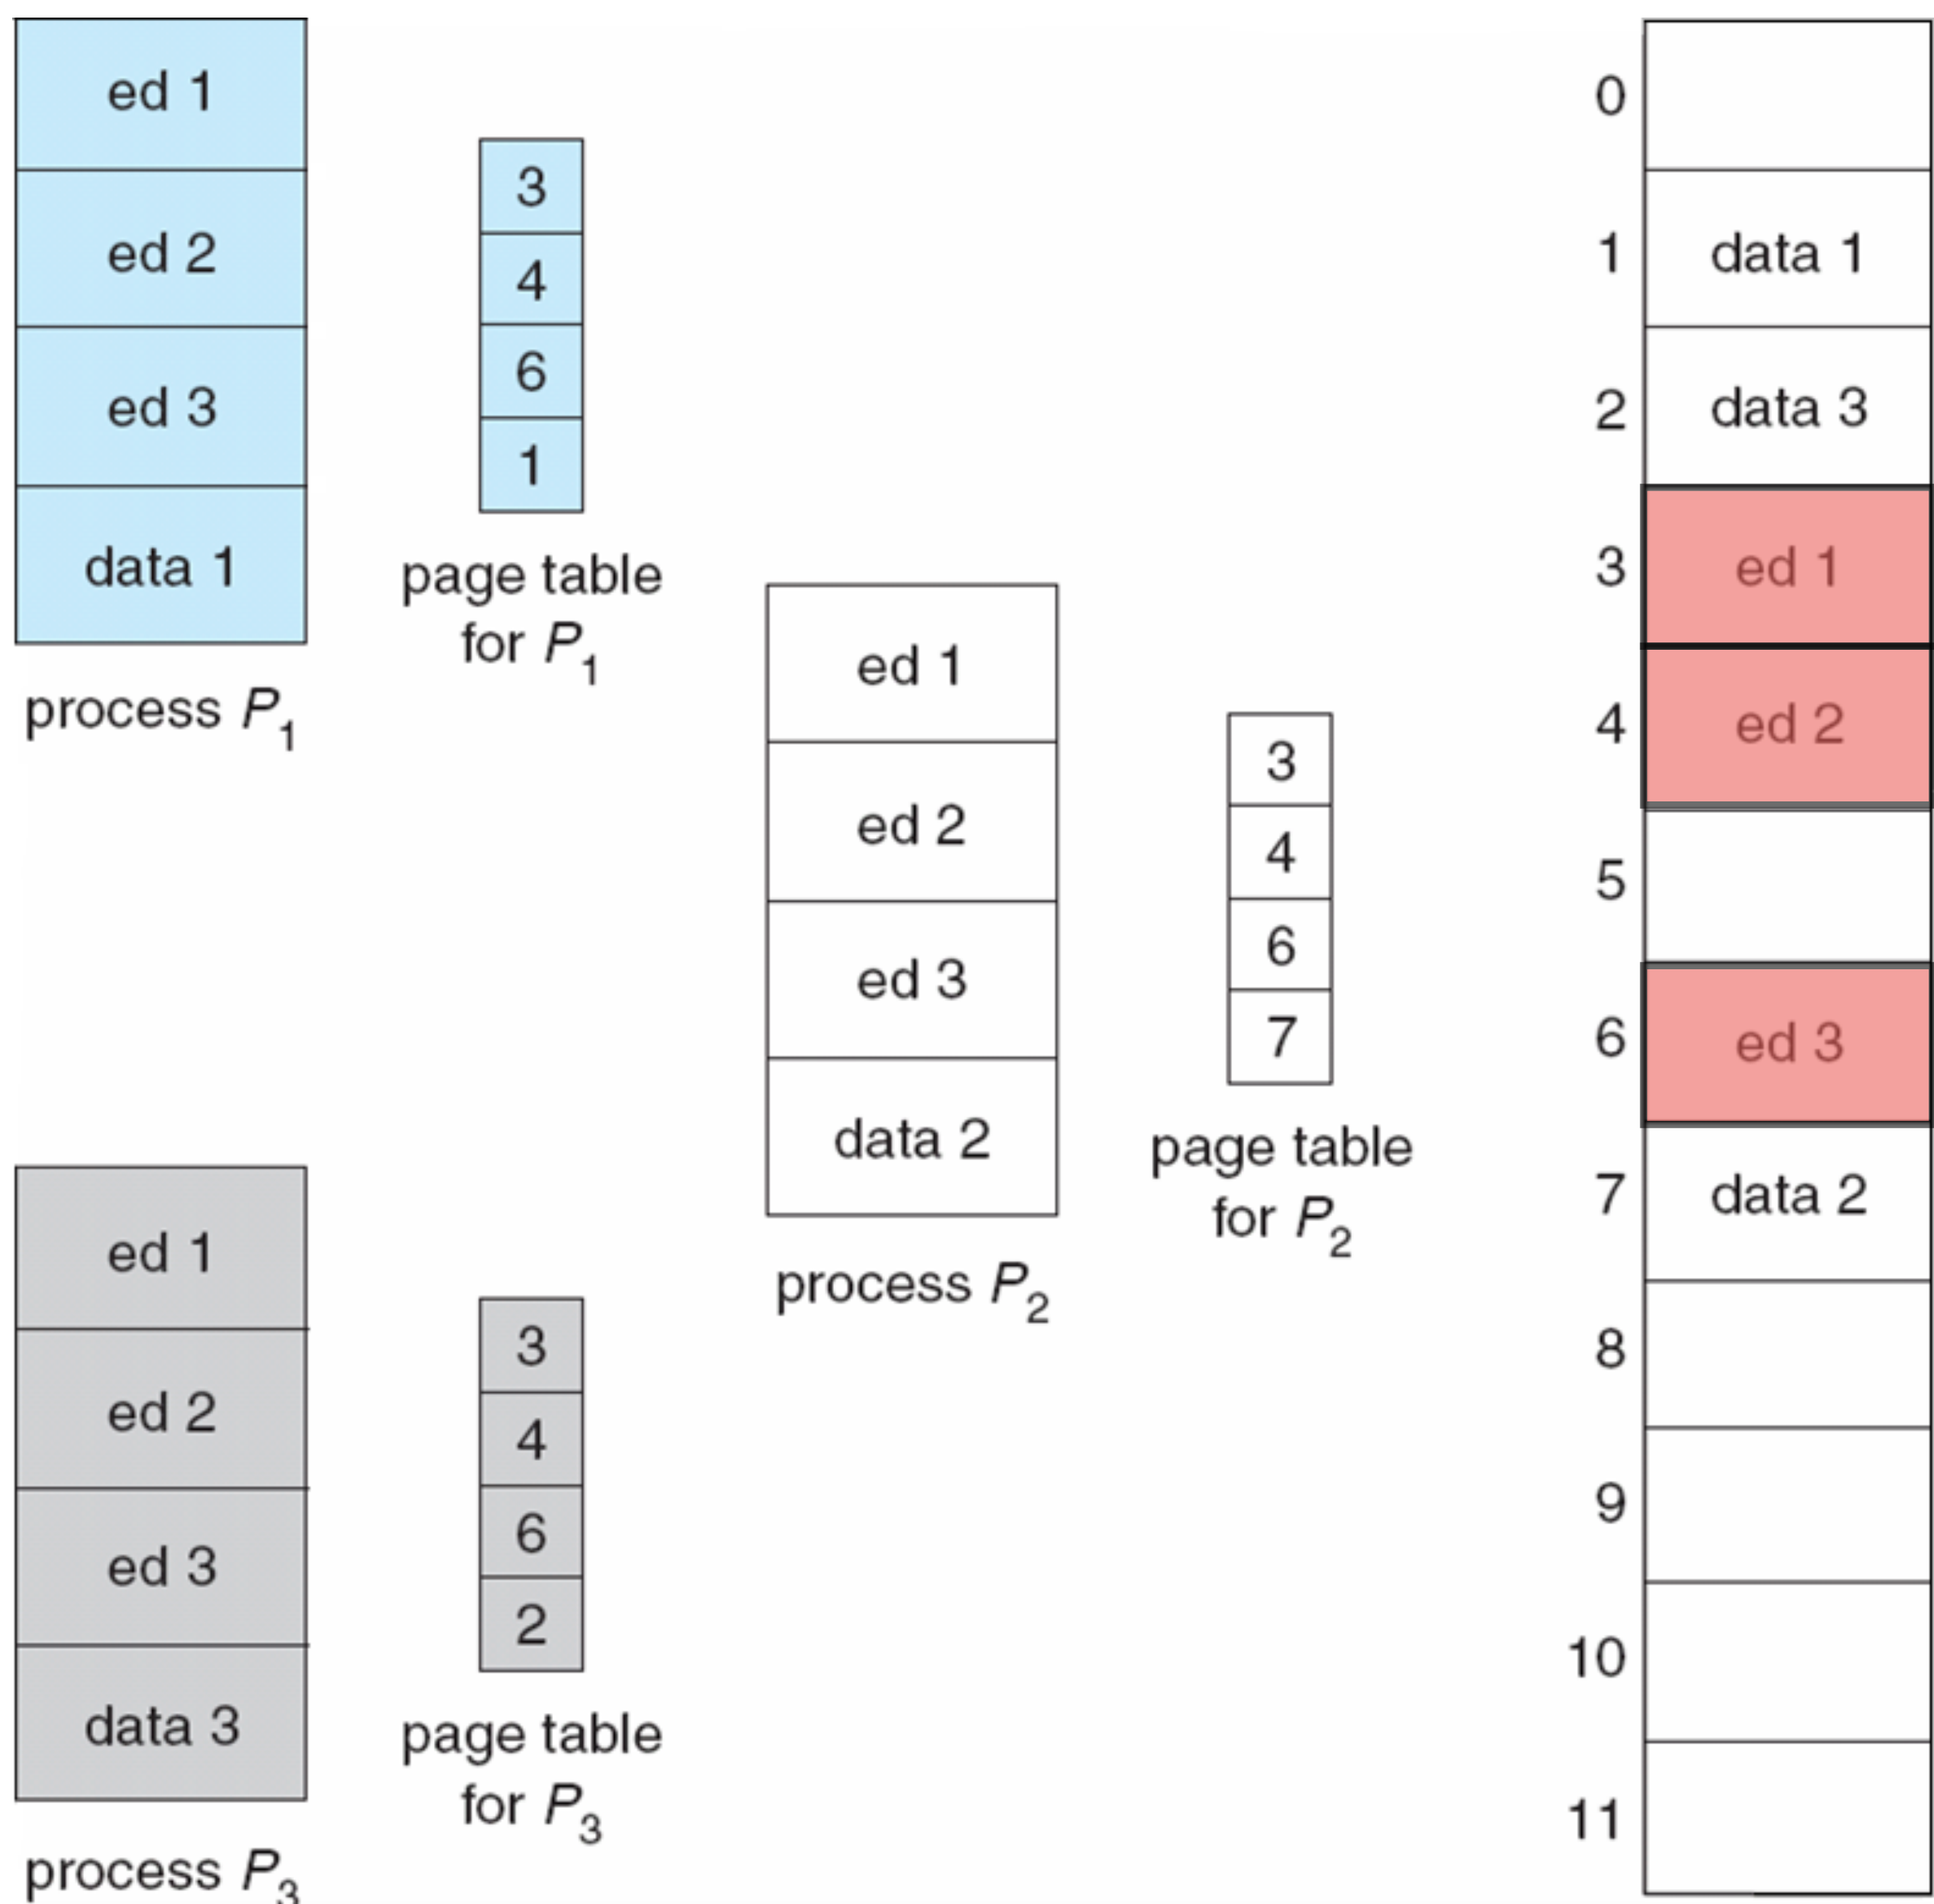
- Shared pages
- Copy-on-write
- Memory-mapped files

# Shared Pages

- Paging allows for the possibility of **sharing common code**.
- Sharing pages is useful in **time-sharing environments** (e.g., 40 users, each executing a text editor).
- OS can implement **shared-memory** (IPC) using shared pages.



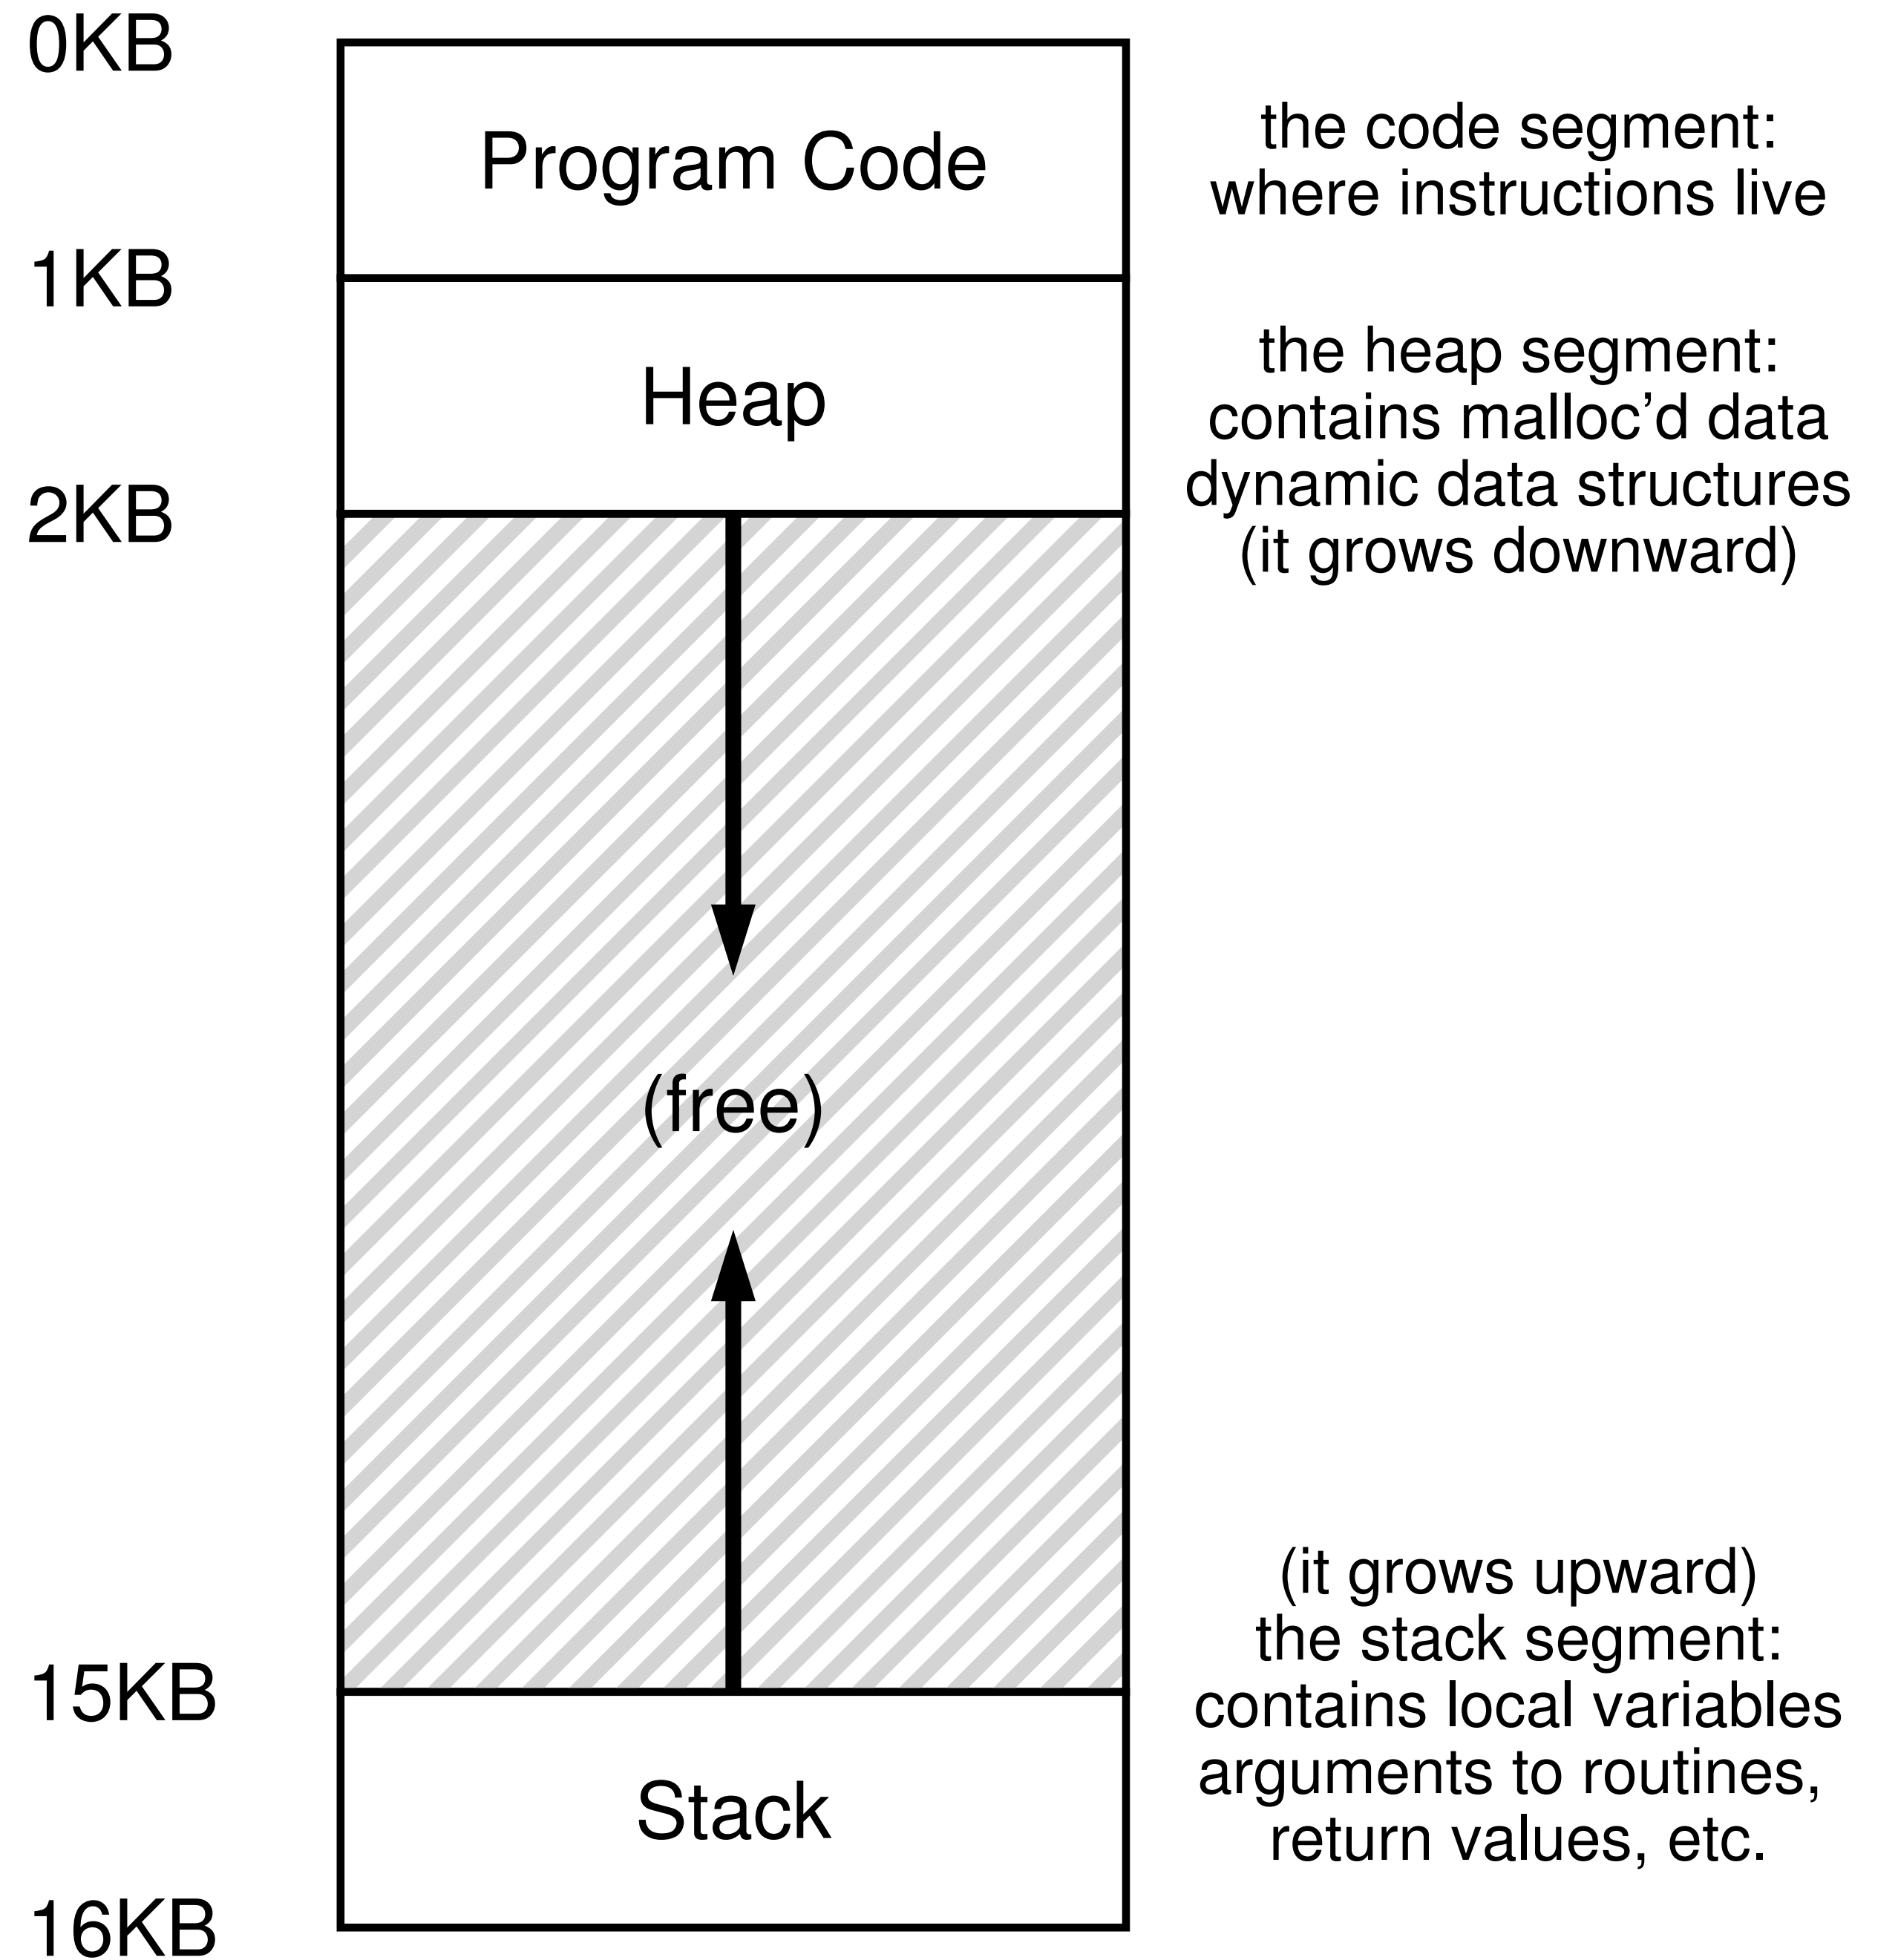
# Example of shared Pages



# Copy-on-Write (COW), e.g. on `fork( )`

- copy-on-write (COW), e.g., on `fork( )`
  - Instead of copying all pages, create shared mappings of parent pages in child address space
    - A. Make shared mappings read-only in child space
    - B. When child does a write, a protection fault occurs, OS takes over and can then copy the page and resume child.

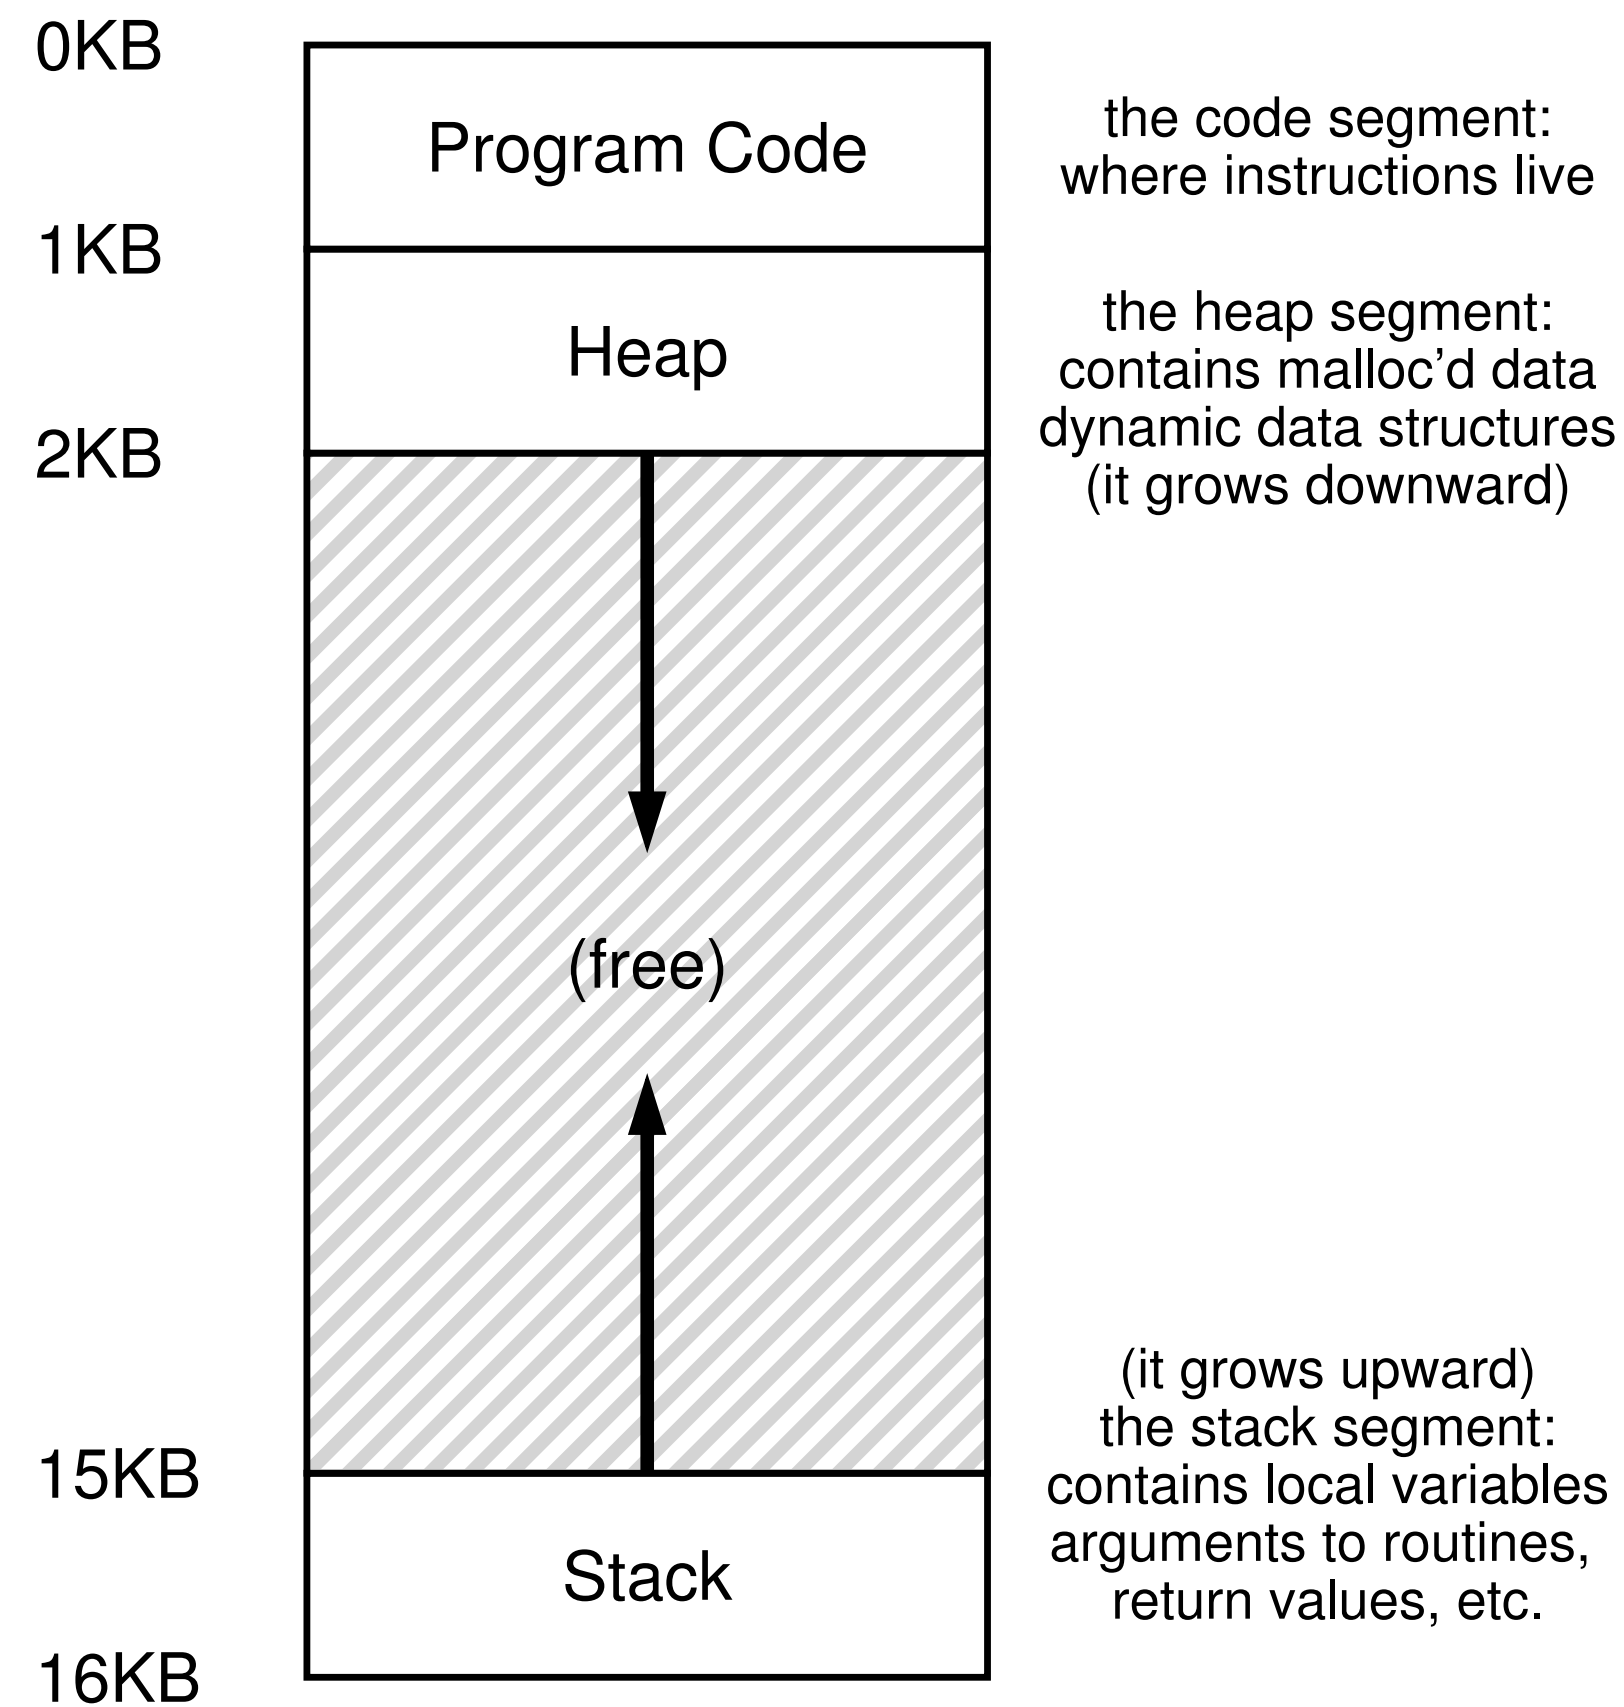
# Example address space



# Types of memory

**Stack:** Short-lived memory. Allocations and deallocations are managed *implicitly* (e.g., by the compiler), not by the programmer.

**Heap:** Long-lived memory. Allocations and deallocations are *explicitly* handled by the programmer.



# Examples

```
void func() {  
    int x;  
    ...  
}
```



# Examples

```
void func() {  
    int *x = (int *) malloc(sizeof(int));  
    ...  
}
```

# Every address you see is virtual

Here's a little program that prints out the locations of the `main()` routine (where code lives), the value of a heap-allocated value returned from `malloc()`, and the location of an integer on the stack:

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  int main(int argc, char *argv[]) {
4      printf("location of code   : %p\n", (void *) main);
5      printf("location of heap   : %p\n", (void *) malloc(1));
6      int x = 3;
7      printf("location of stack  : %p\n", (void *) &x);
8      return x;
9  }
```

When run on a 64-bit Mac OS X machine, we get the following output:


```
location of code   : 0x1095afe50
location of heap   : 0x1096008c0
location of stack  : 0x7fff691aea64
```

# Segmentation

- Memory-management scheme that supports the user's view of memory.
- View memory as a collection of variable-sized segments, with no necessary ordering among segments.

# Segmentation (a program)


- We think of a program as a main program, a stack, a math library, etc.
- Each module is referred to by name
- In this view of a program, we might not care whether the stack is stored before or after the `sqrt()` function.




# Logical view of segmentation

- For simplicity of implementation, each segment is addressed by a **segment number** and an **offset**:

<segment-number, offset>



user space



physical memory space



# Segmentation Hardware

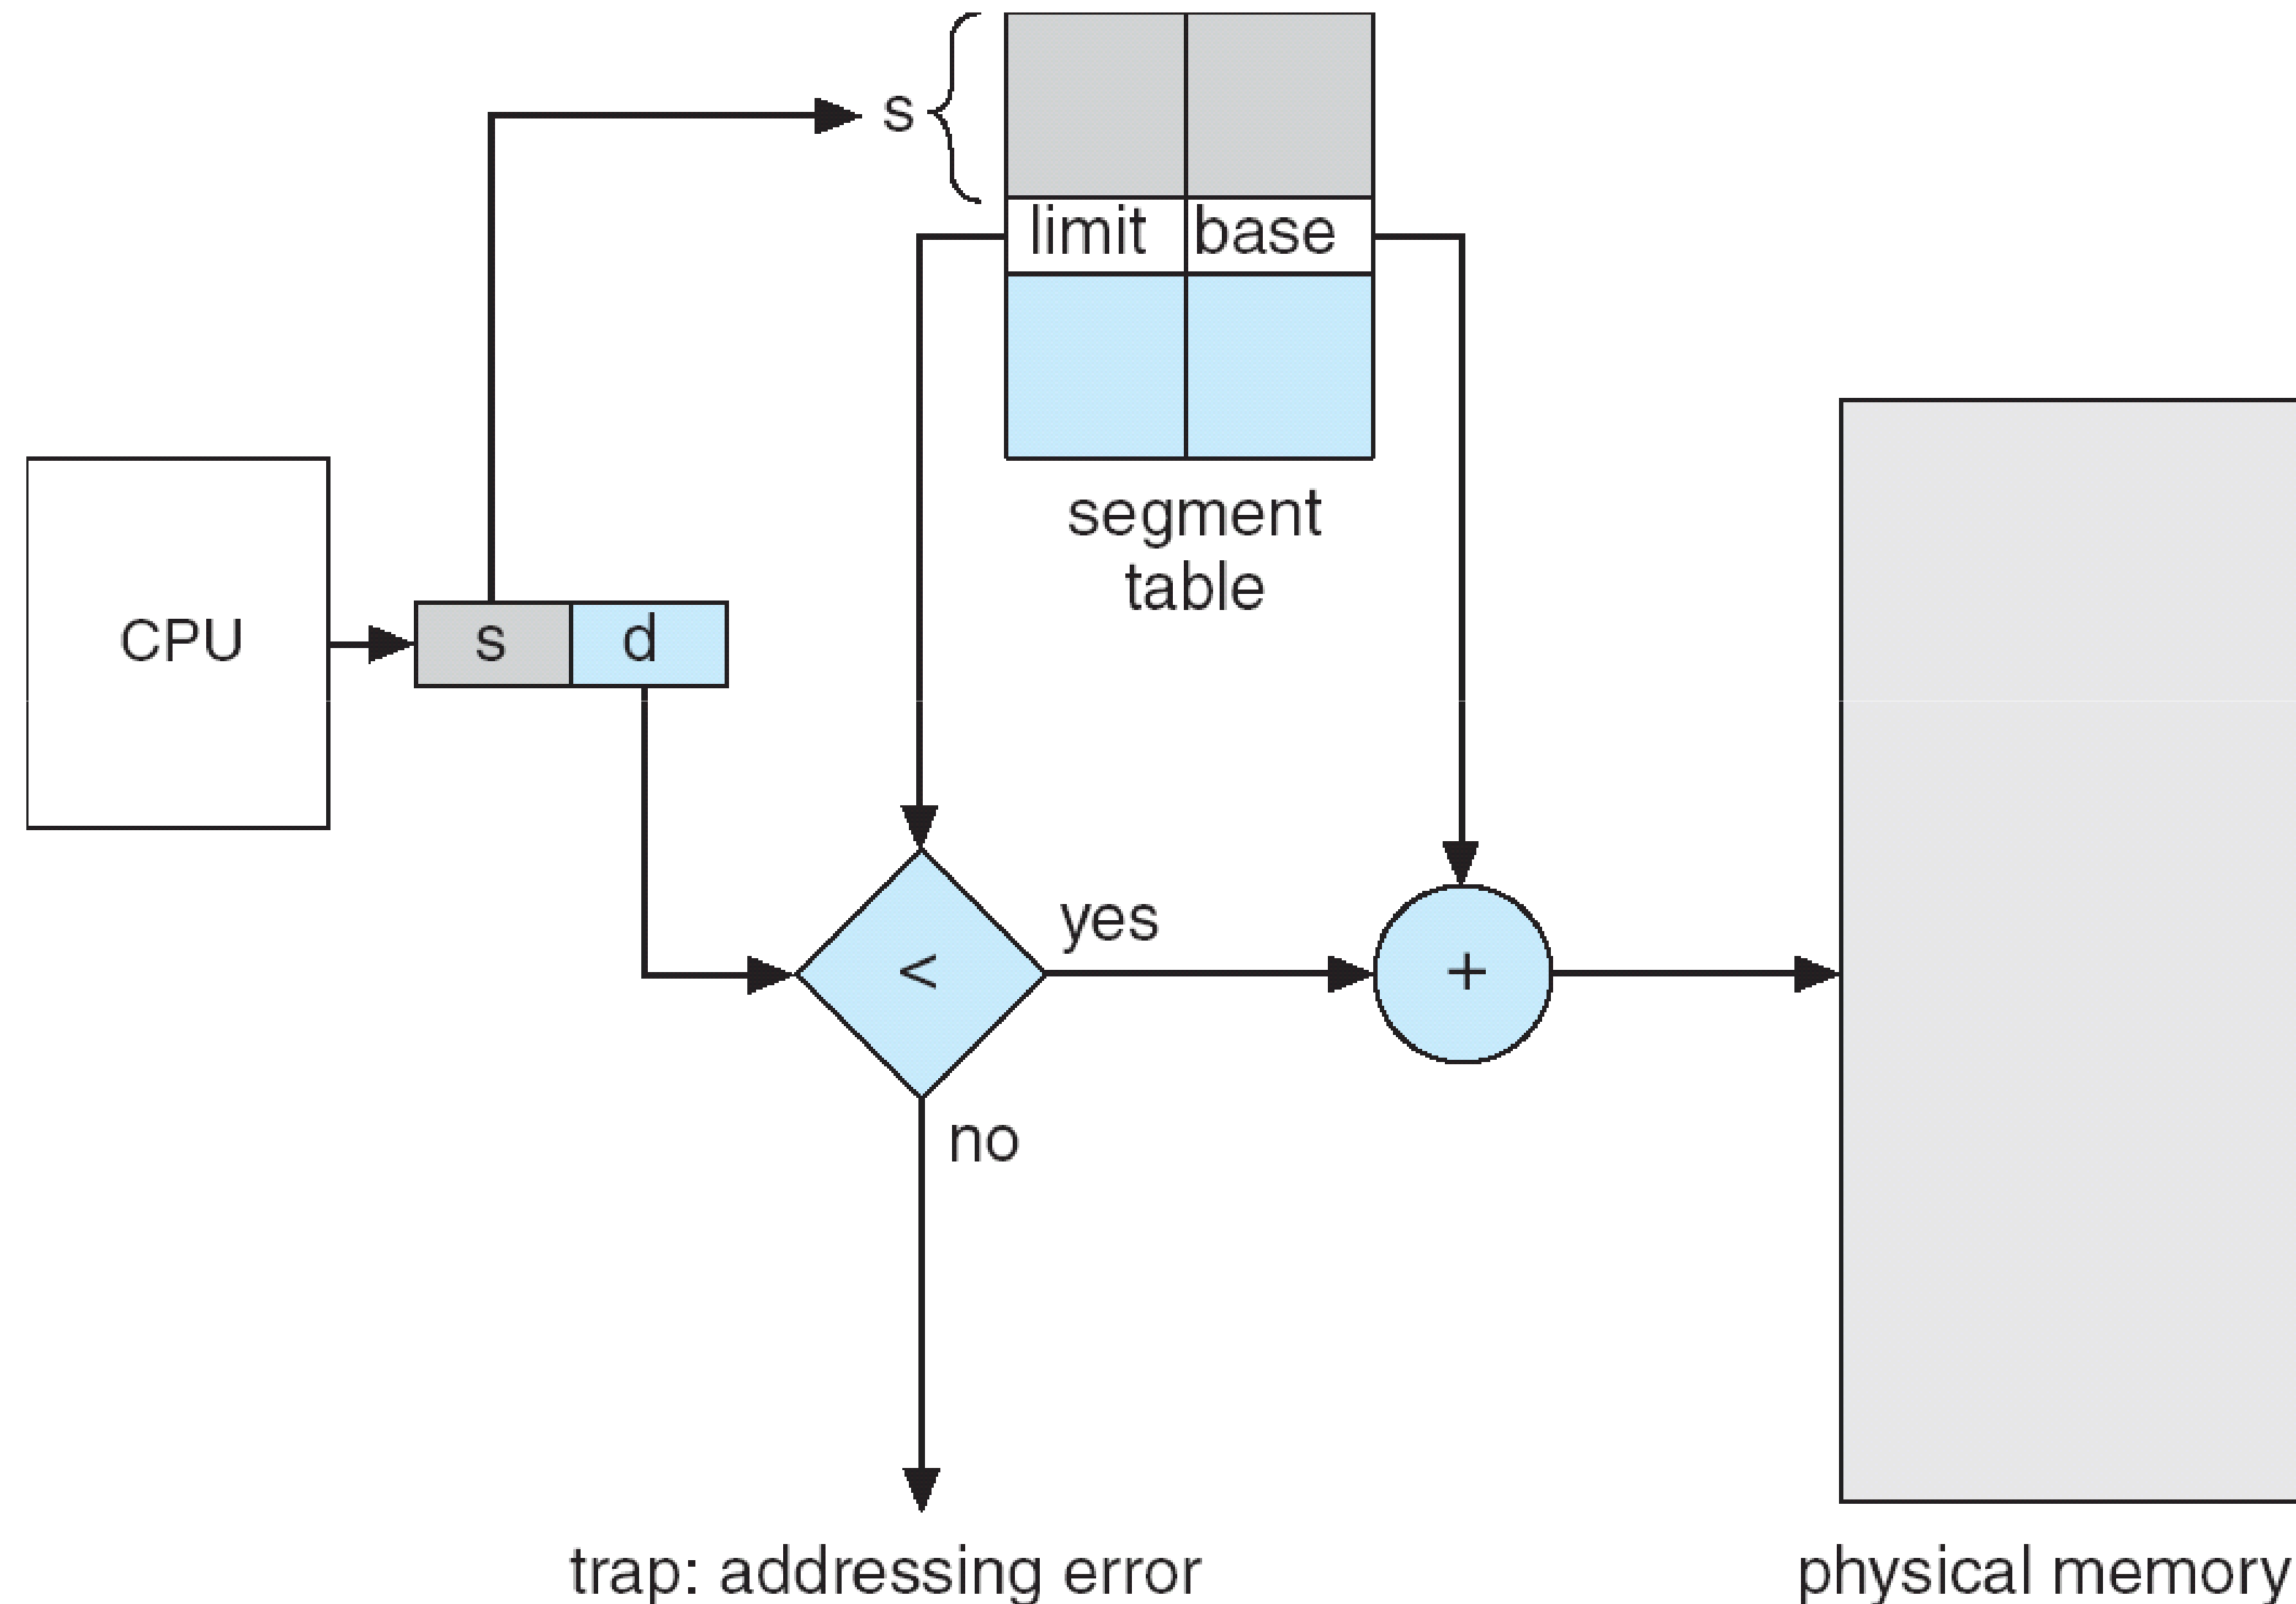
## Segment tables:

**Base:** starting address of the segment in physical memory.

**Limit:** length of the segment.

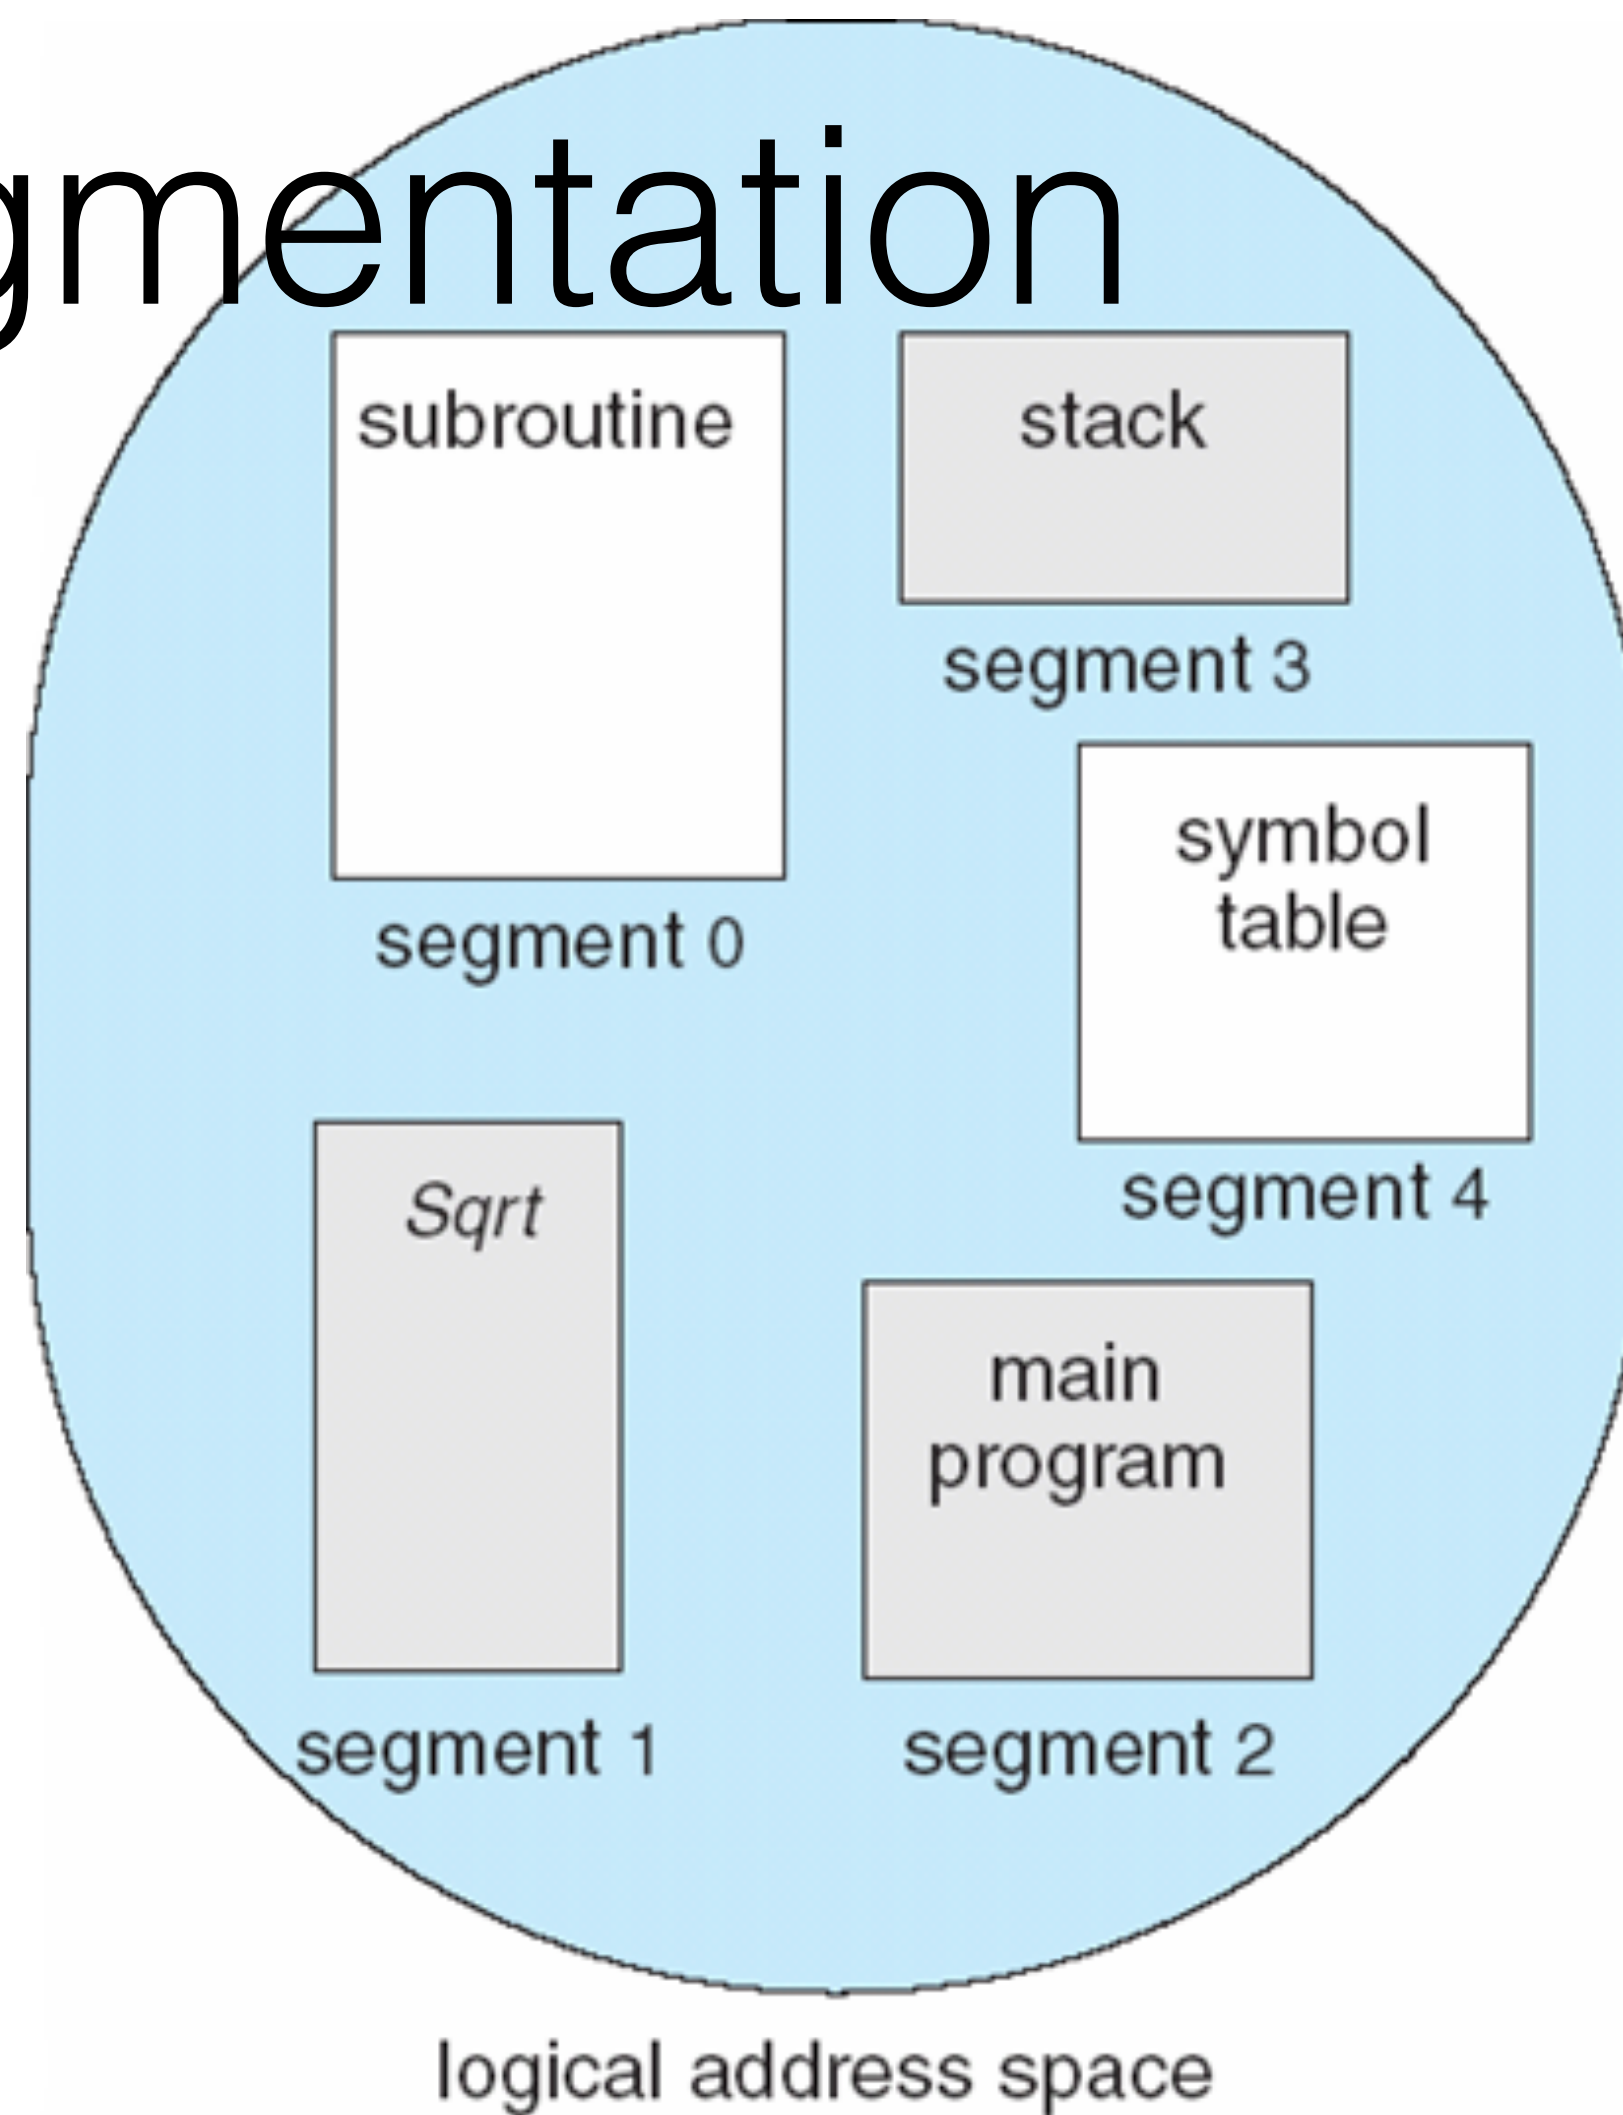
Additional metadata includes protection bits.

<segment-number, offset>




# Example of segmentation

- Logical memory divided into 5 segments.
- Segment 2 is 400 bytes long and begins at location 4,300.
- **Question:** What happens if there is a reference to byte 1,222 of segment 0?



	limit	base
0	1000	1400
1	400	6300
2	400	4300
3	1100	3200
4	1000	4700

segment table



# Some questions

How do paging and segmentation compare with respect to the following issues?

- External fragmentation
- Internal fragmentation
- Ability to share code across processes

# Some questions

Assuming a 1-KB page size, what are the page numbers and offset for the following address:

A. 2375

B. 256

# Some questions

Segment	Base	Length
0	219	600
1	2300	14
2	90	100
3	1327	580
4	1952	96

What are the physical addresses for the following logical addresses?

a. 0,430

b. 2,500