

# Virtual Memory


CSE 4001

# Content

- Demand paging

# Virtual Memory

- Separation of user logical memory from physical memory.
- Programs can be partially in memory for execution
- Logical address space can be much larger than physical address space



# Implementation


Virtual memory can be implemented via:

- Demand paging
- Demand segmentation

# Demand paging

Bring a page into memory only when it is needed:


- Less I/O needed
- Less memory needs
- Faster response
- More users





# Demand paging

- Demand paging is similar to a paging system with swapping, where processes reside in secondary memory (e.g., disk).
- **Lazy swapper**: only bring pages when they are needed.
- In the context of demand paging, we use the term **pager** instead of **swapper**.





# Valid-Invalid Bit


- **Hardware support** is needed to distinguish between the pages that are in memory and the ones that are on the disk.
- We can re-use the support provided by the **valid-invalid bit** in the page table.
  - **Bit == valid** then page is in memory (and is valid).
  - **Bit == invalid** then page is either not a valid one for that process or is valid but is currently in disk (pager needs to bring it to main memory).

# Valid-Invalid Bit

- Marking a page invalid has no effect if the process never attempts to access that page.
- Pages that are in memory are called **memory resident**.




physical memory

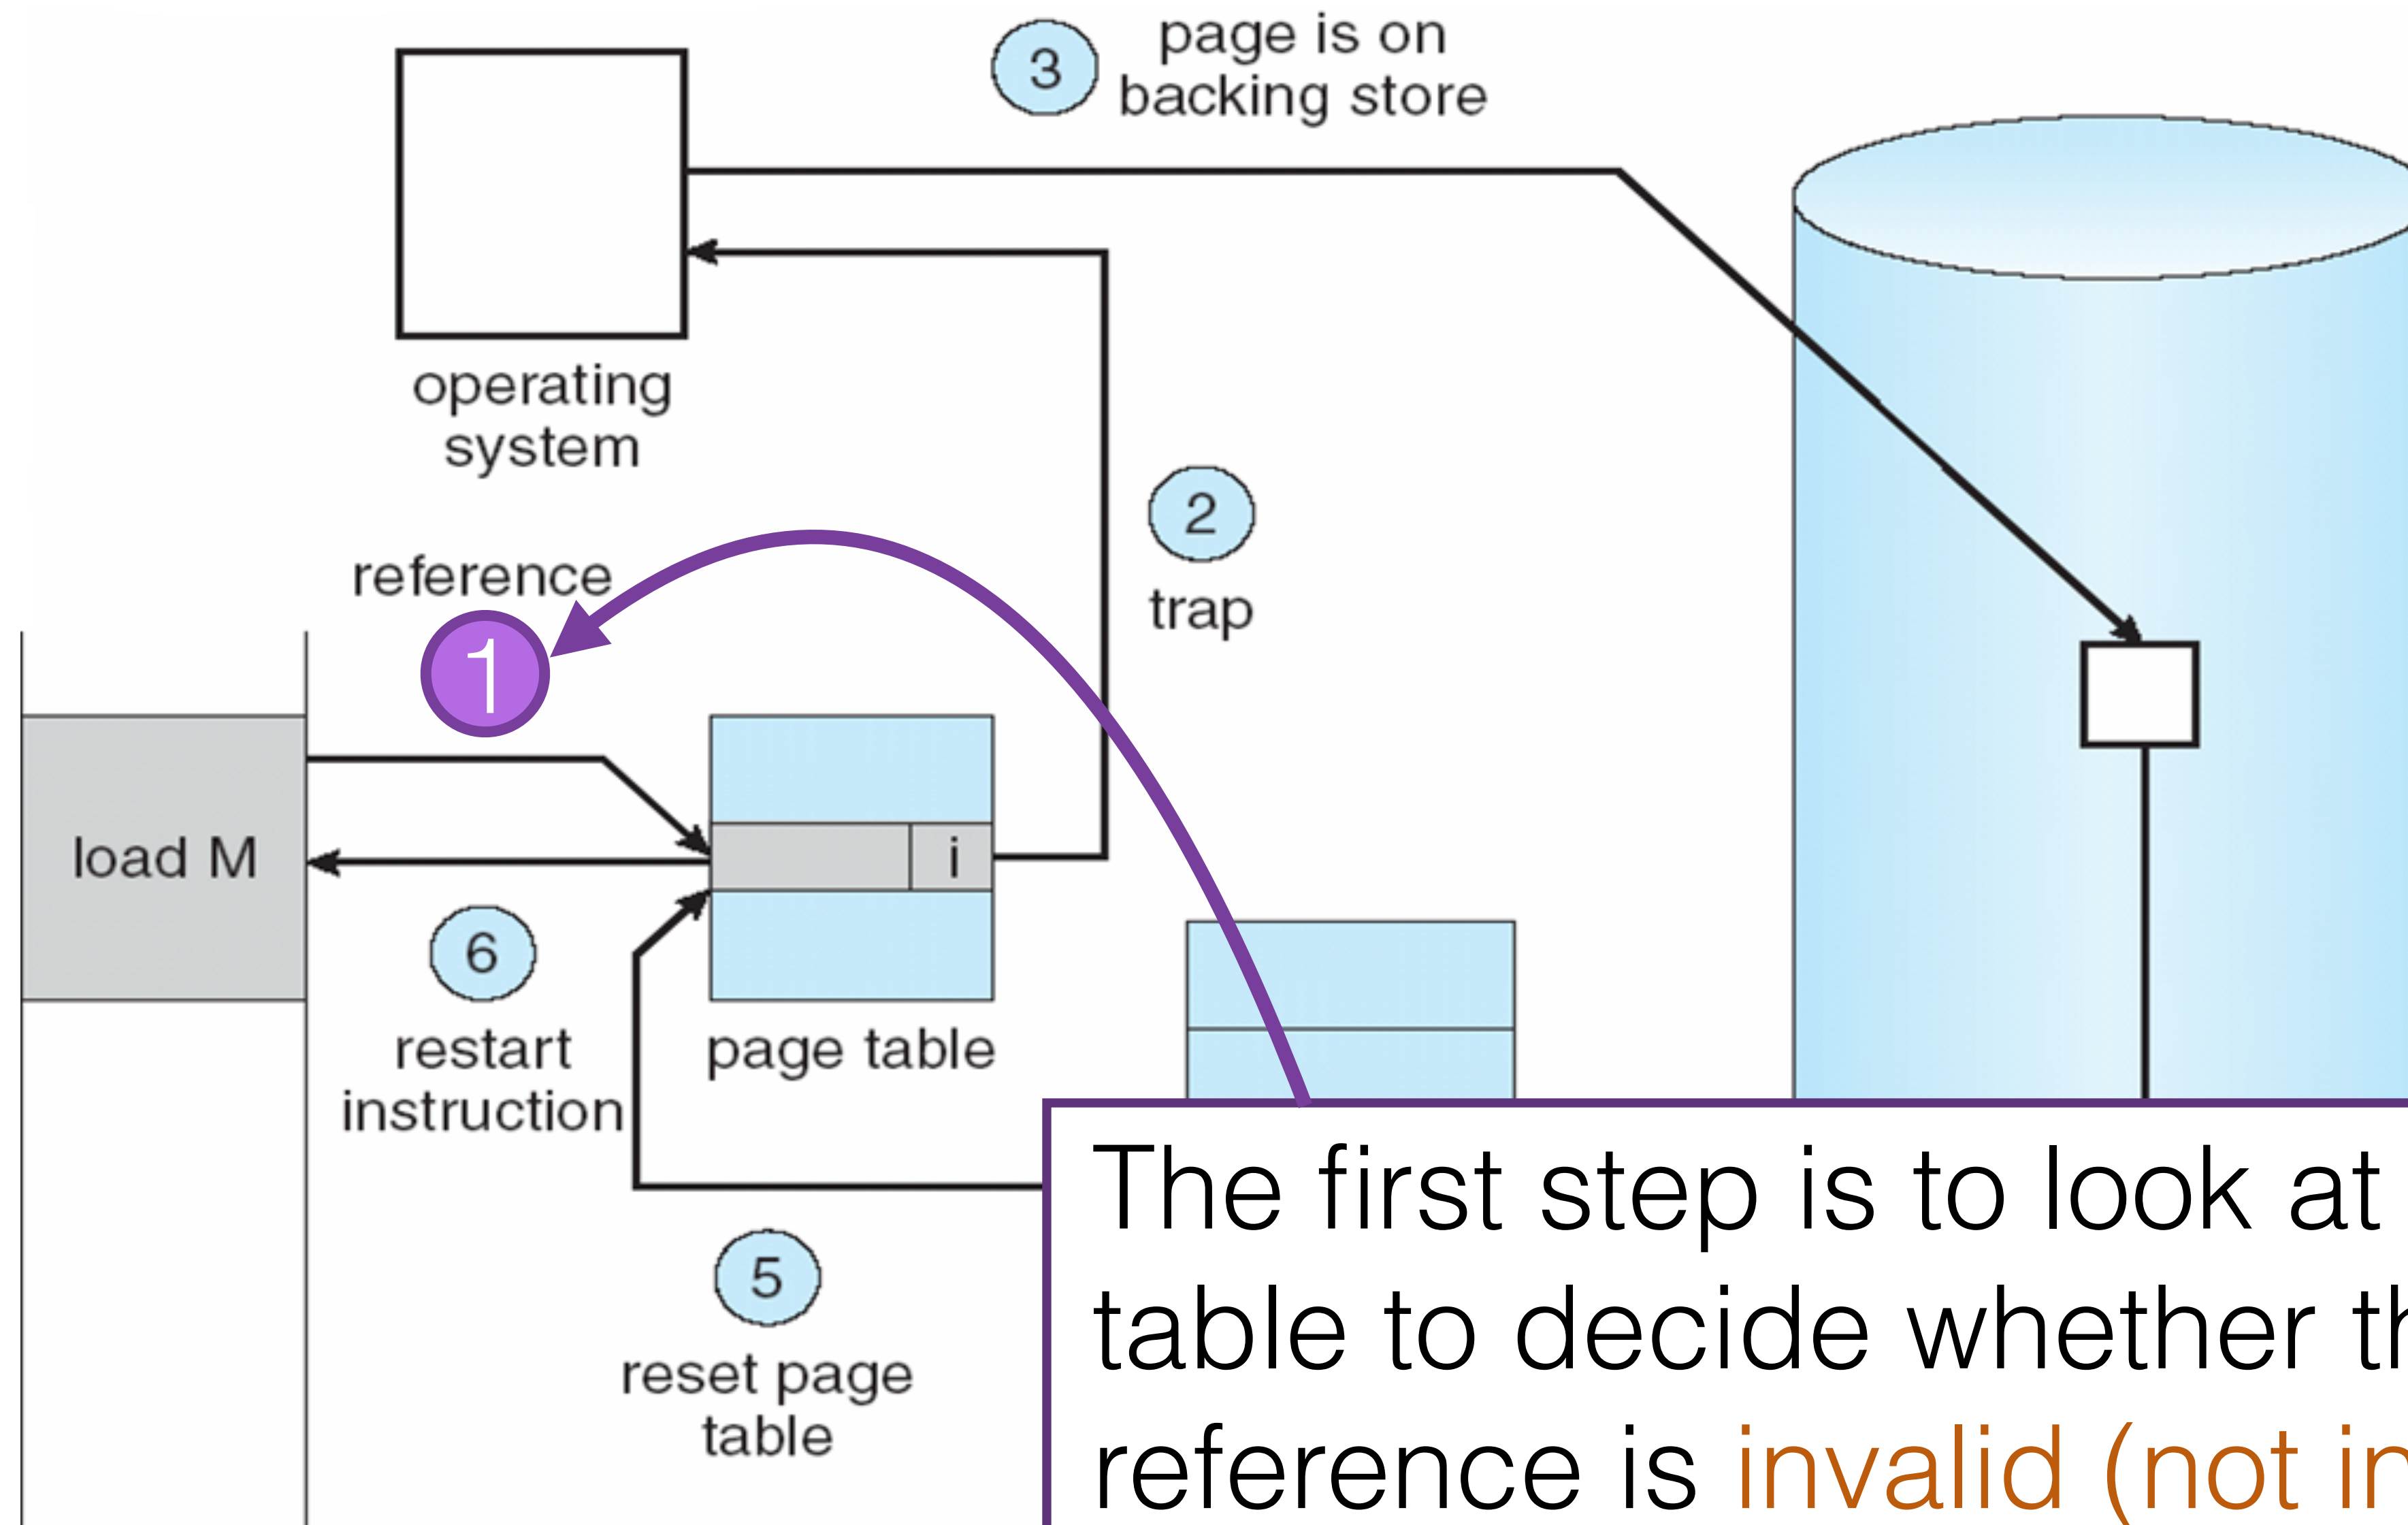




# Page Faults

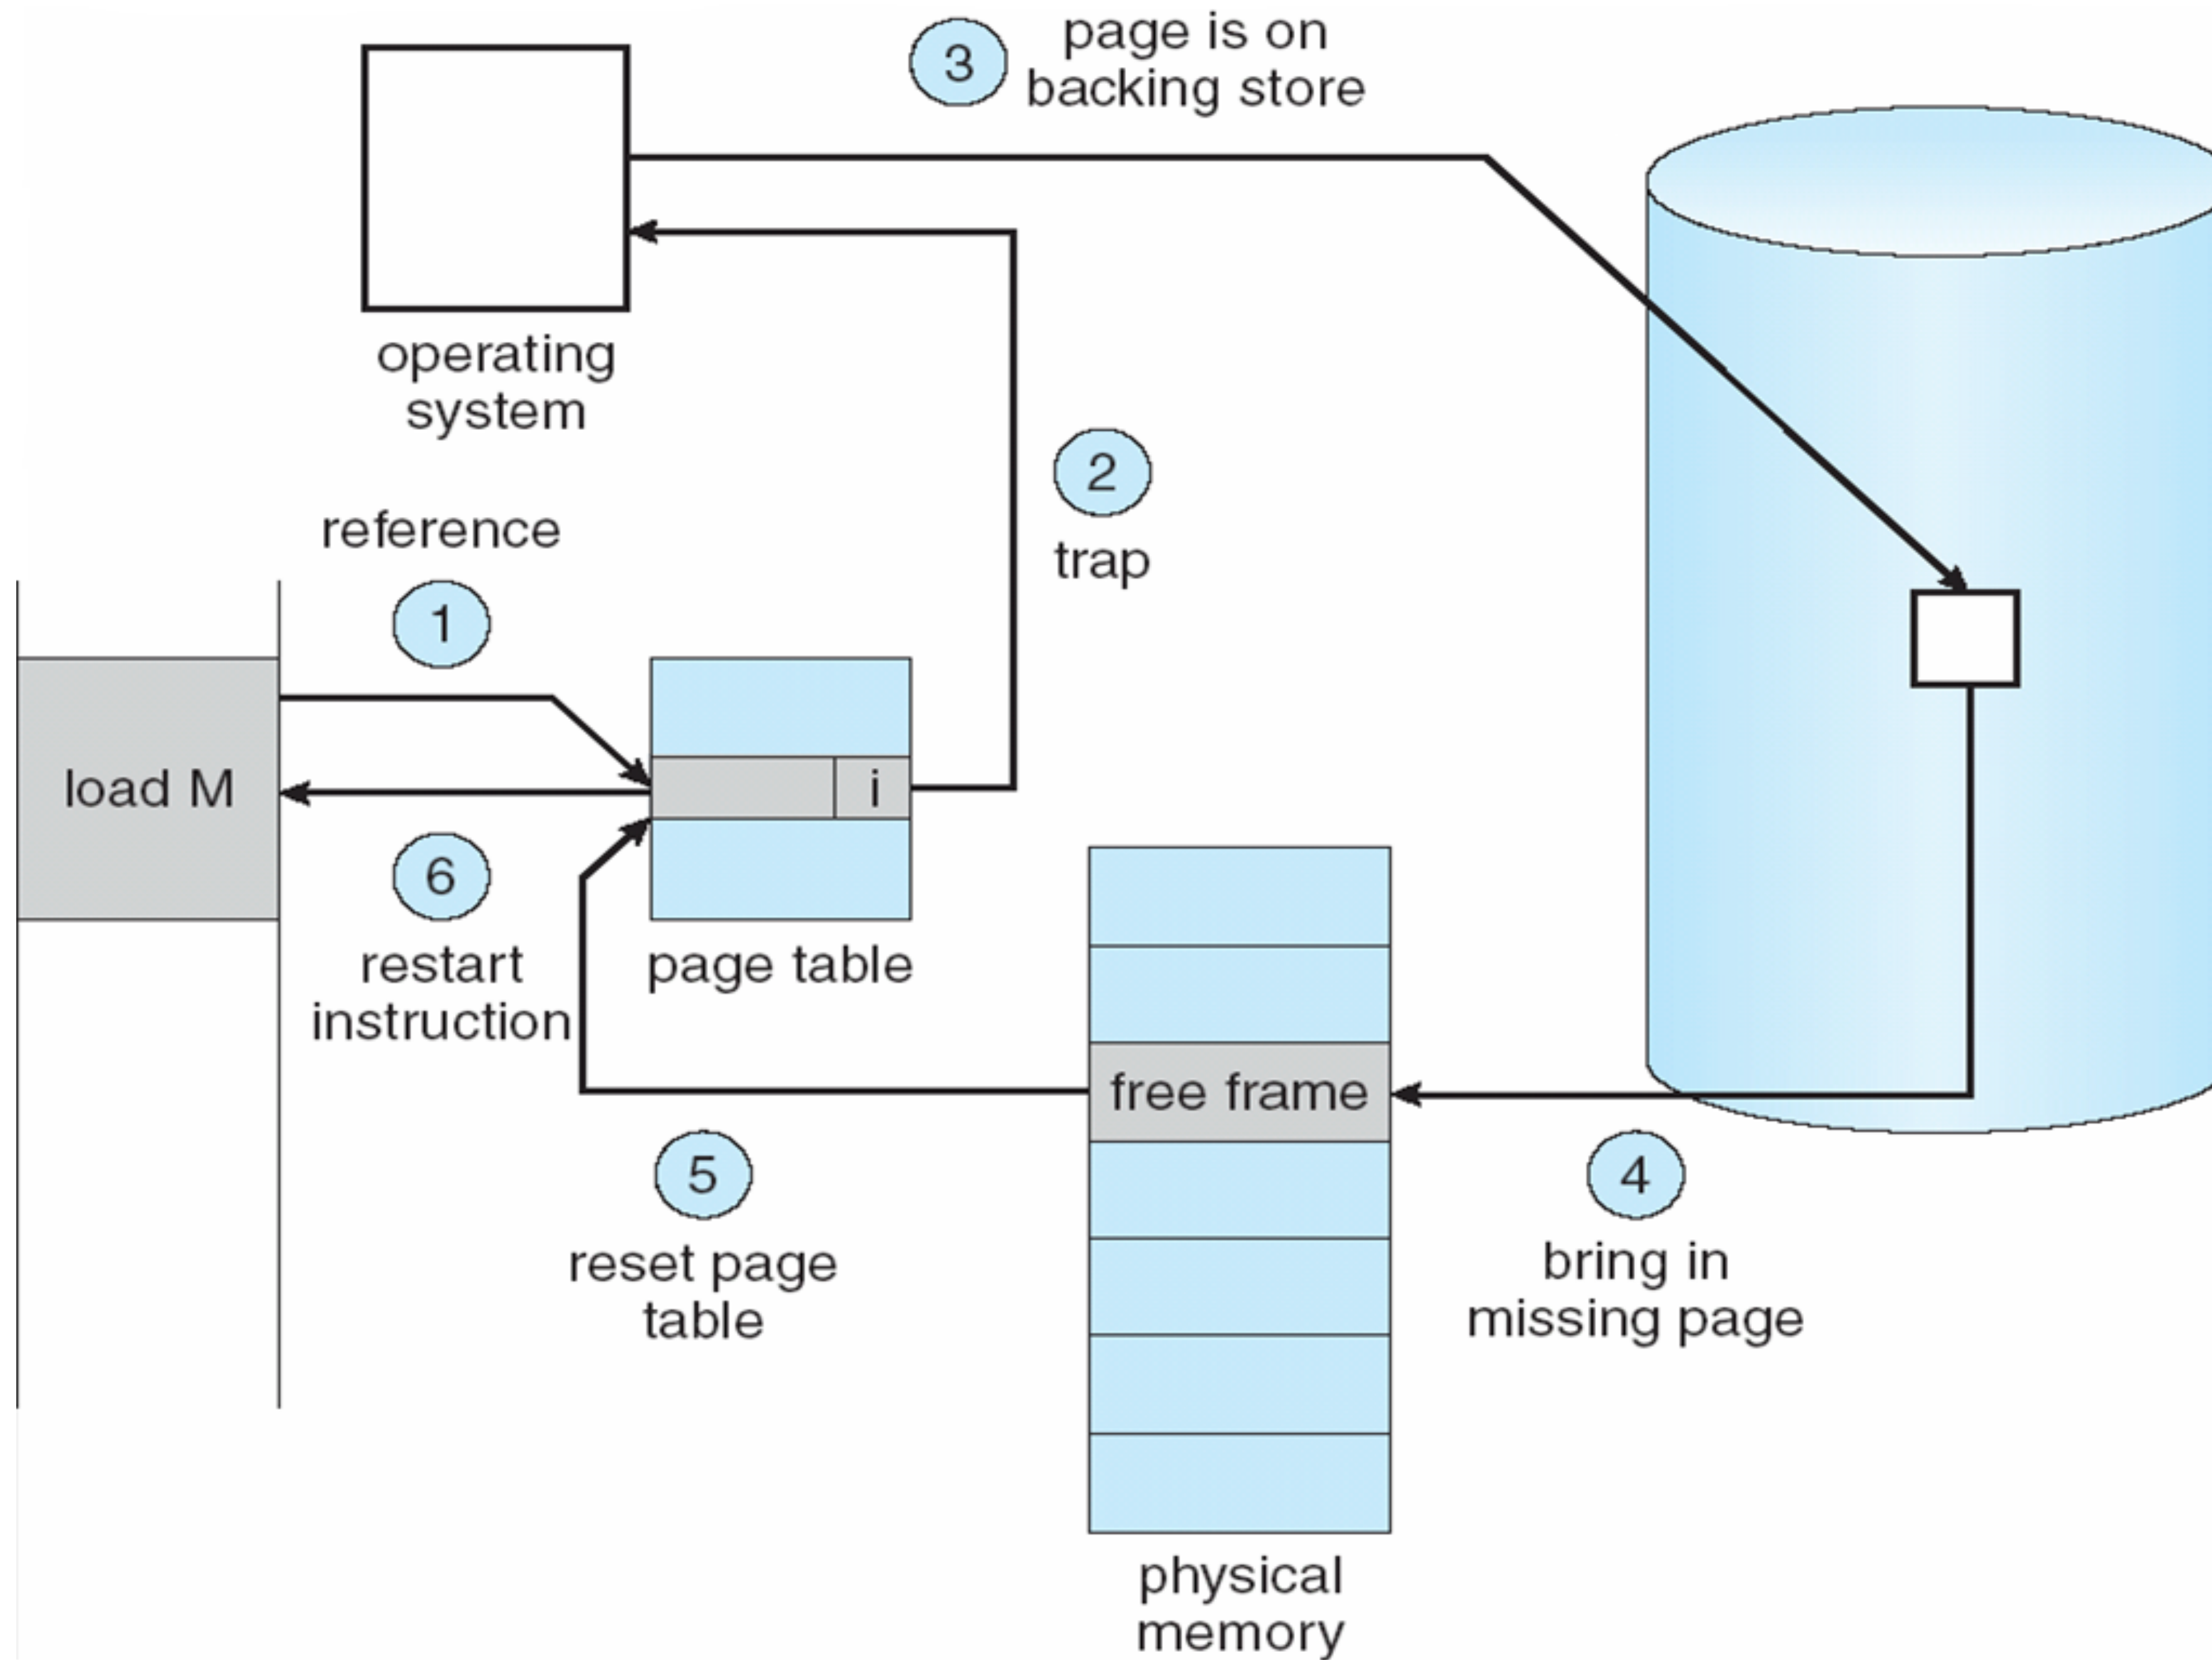
- What happens when a process tries to access non-resident pages?
- **Page Fault:** A trap that results because the OS's failed to bring the desired page into memory.



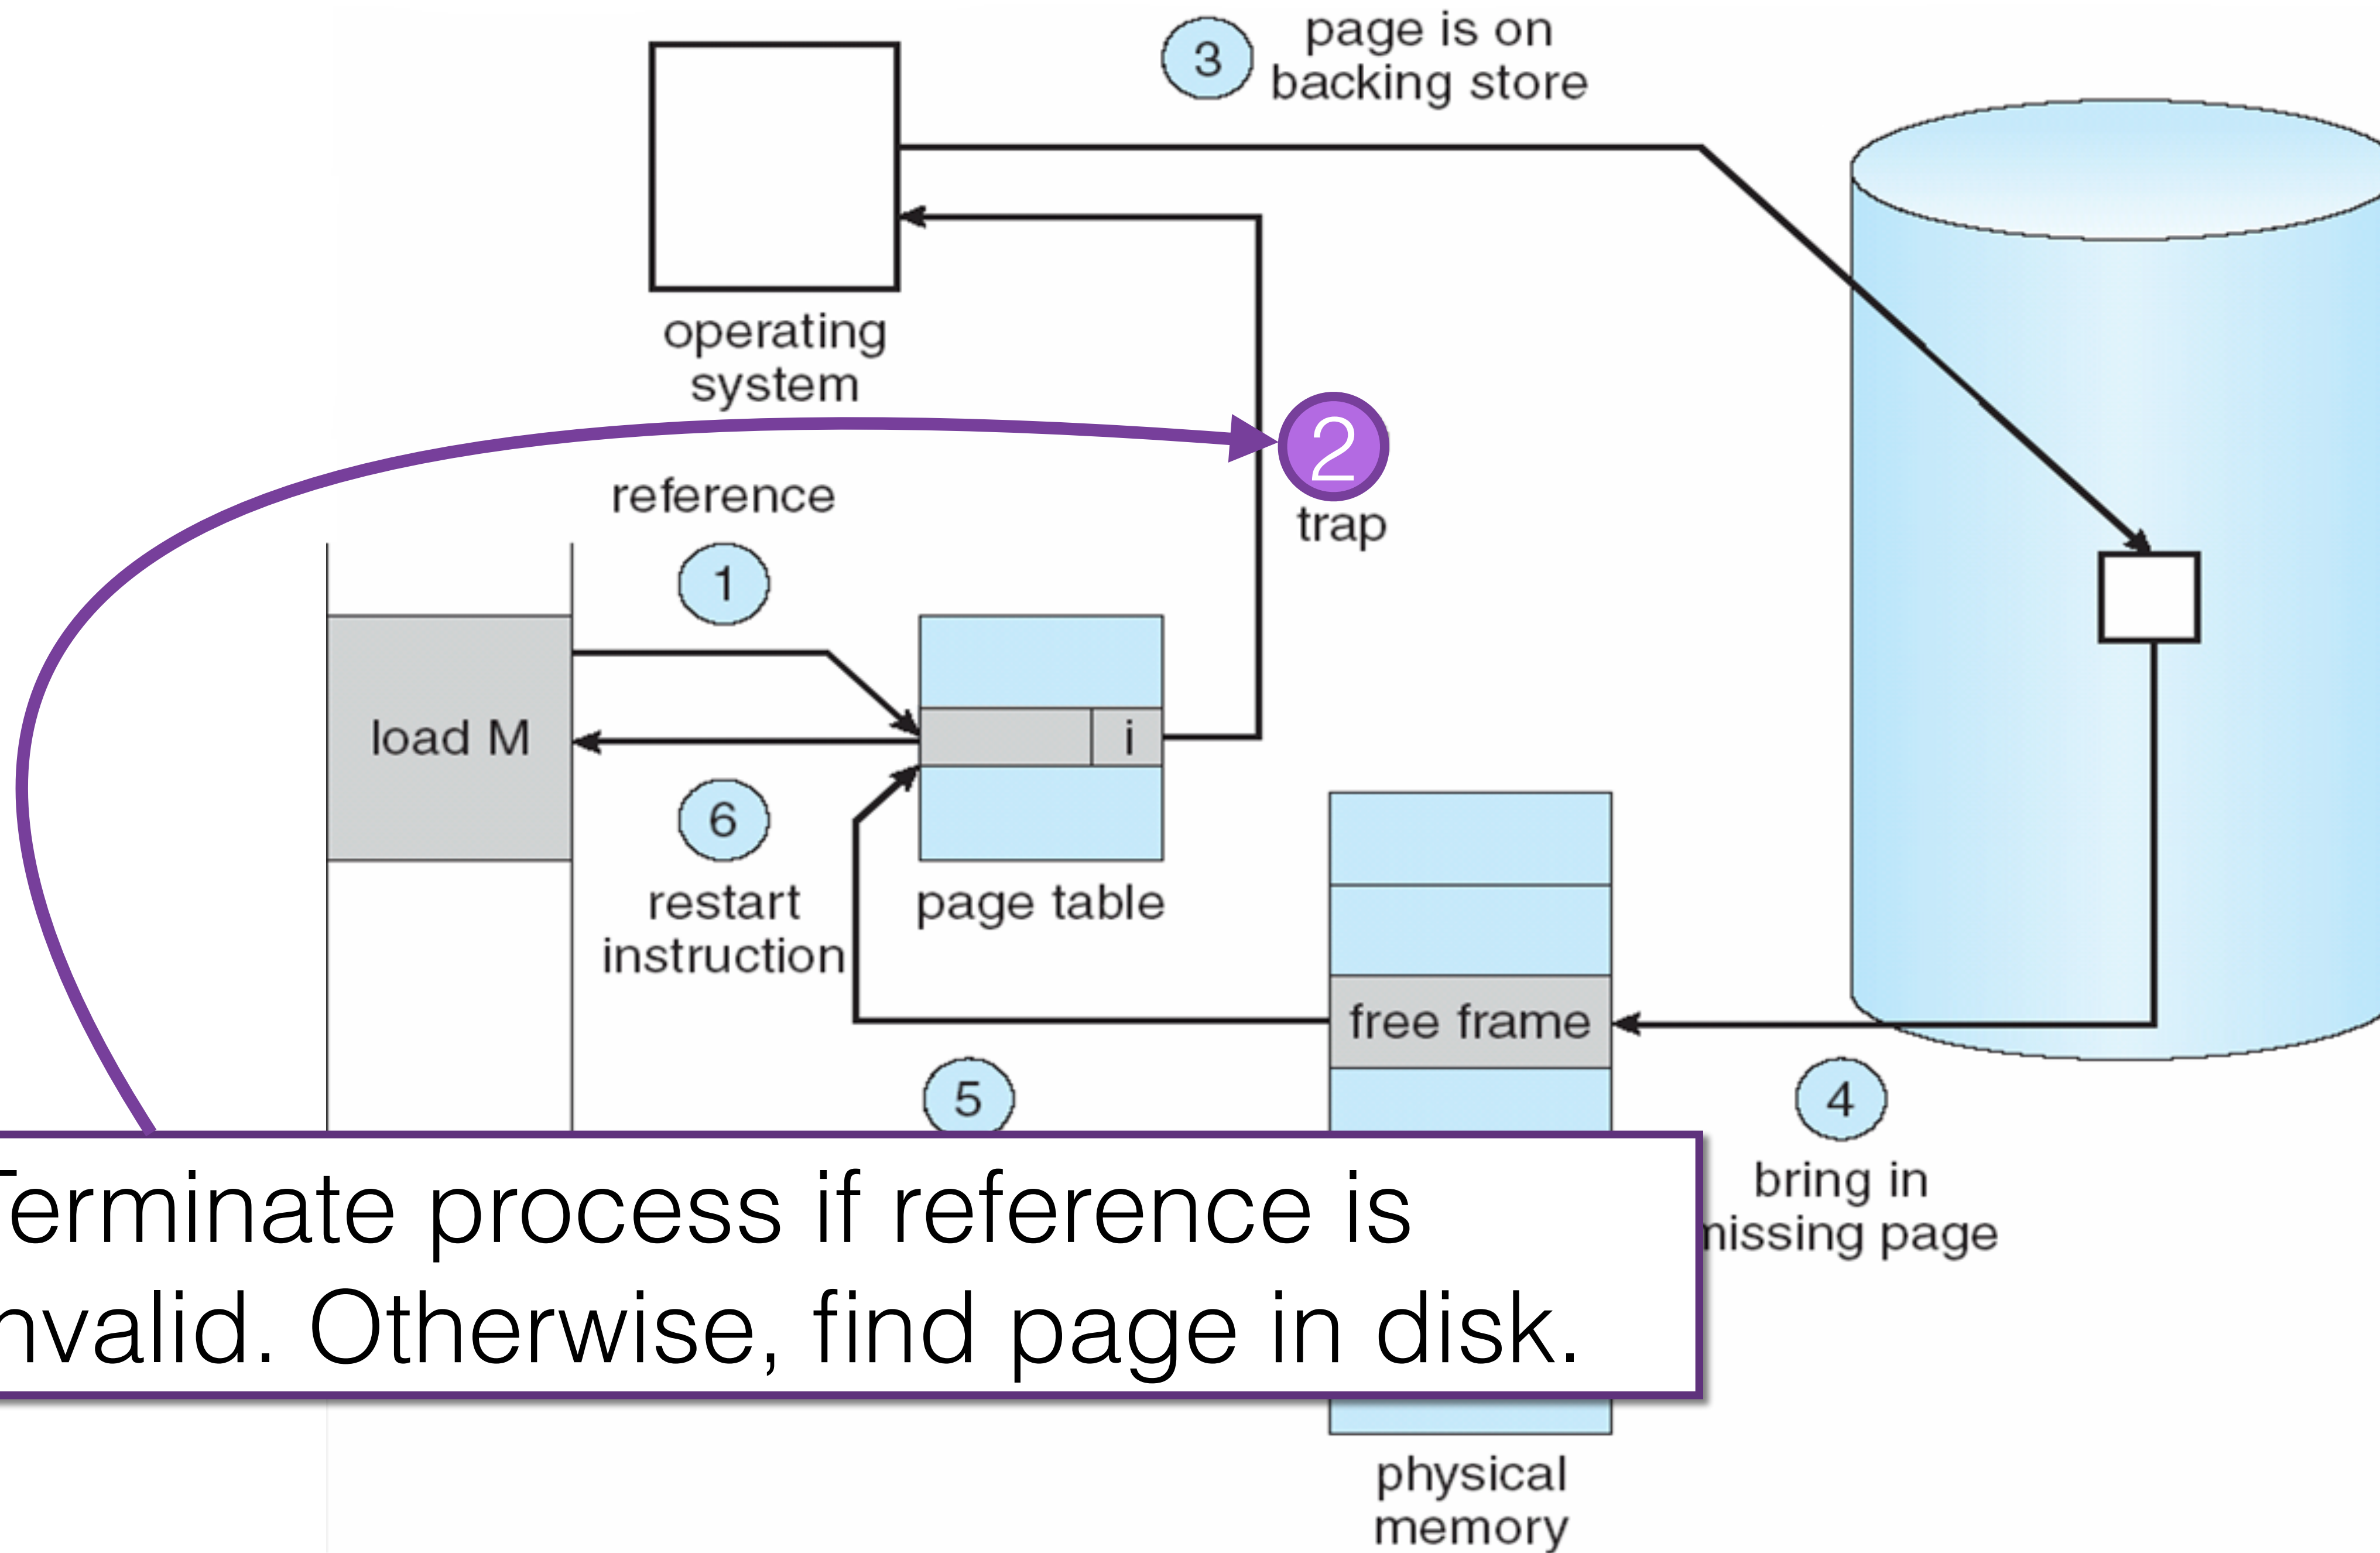


The first step is to look at another table to decide whether the actual reference is **invalid** (not in the process address space) or is simply **not in memory**.






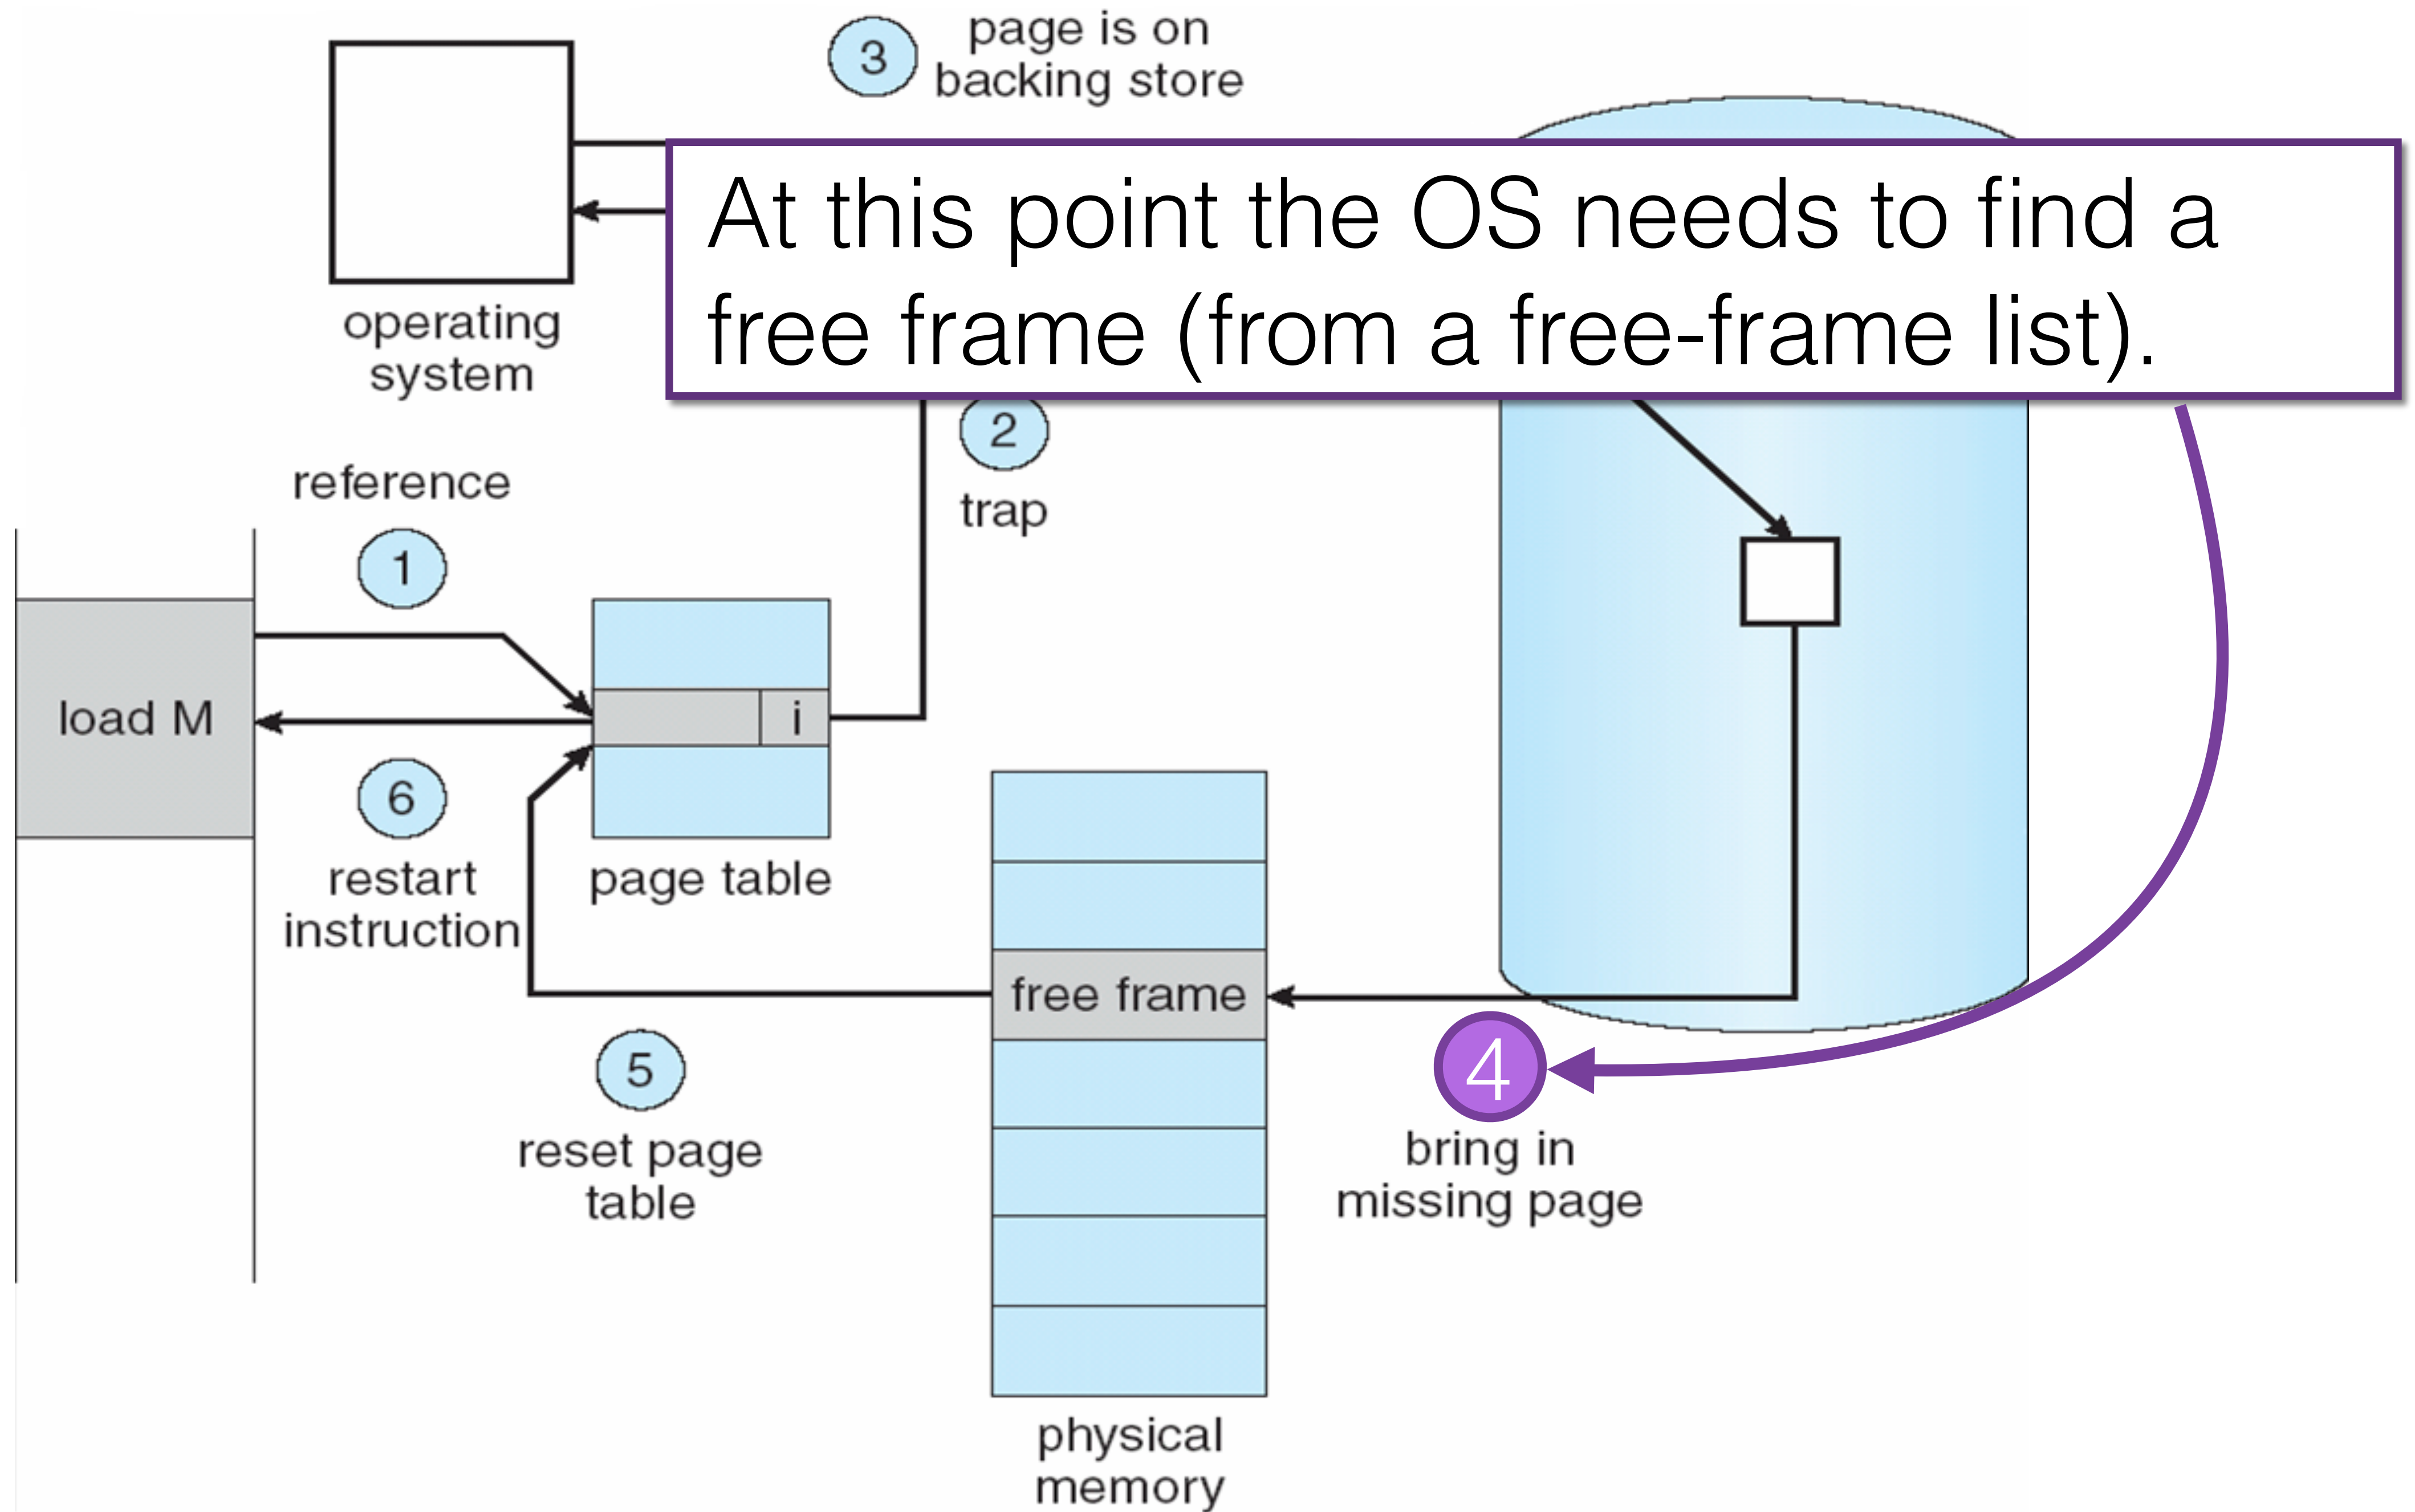





Terminate process if reference is invalid. Otherwise, find page in disk.












# Locality in memory-reference pattern

- Theoretically, some programs could access several new pages with a single instruction.
- In this case, system performance could be seriously degraded.
- Luckily, this behavior is unlikely.



# Writing code with demand-paging in mind...

## ■ Program structure

- `Int[128,128] data;`
- Each row is stored in one page

### ● Program 1

```
for (j = 0; j < 128; j++)  
    for (i = 0; i < 128; i++)  
        data[i,j] = 0;
```

128 x 128 = 16,384 page faults

### ● Program 2

```
for (i = 0; i < 128; i++)  
    for (j = 0; j < 128; j++)  
        data[i,j] = 0;
```

128 page faults


# Thrashing

- The process does not have “enough” pages, the page-fault rate is very high and CPU becomes sub-utilized.
- The OS wants to maximize CPU utilization. As a result, it decides that it is a good idea to increase the degree of multiprogramming by adding new processes to the system.
- **Thrashing**: A process is spending more time paging than executing.



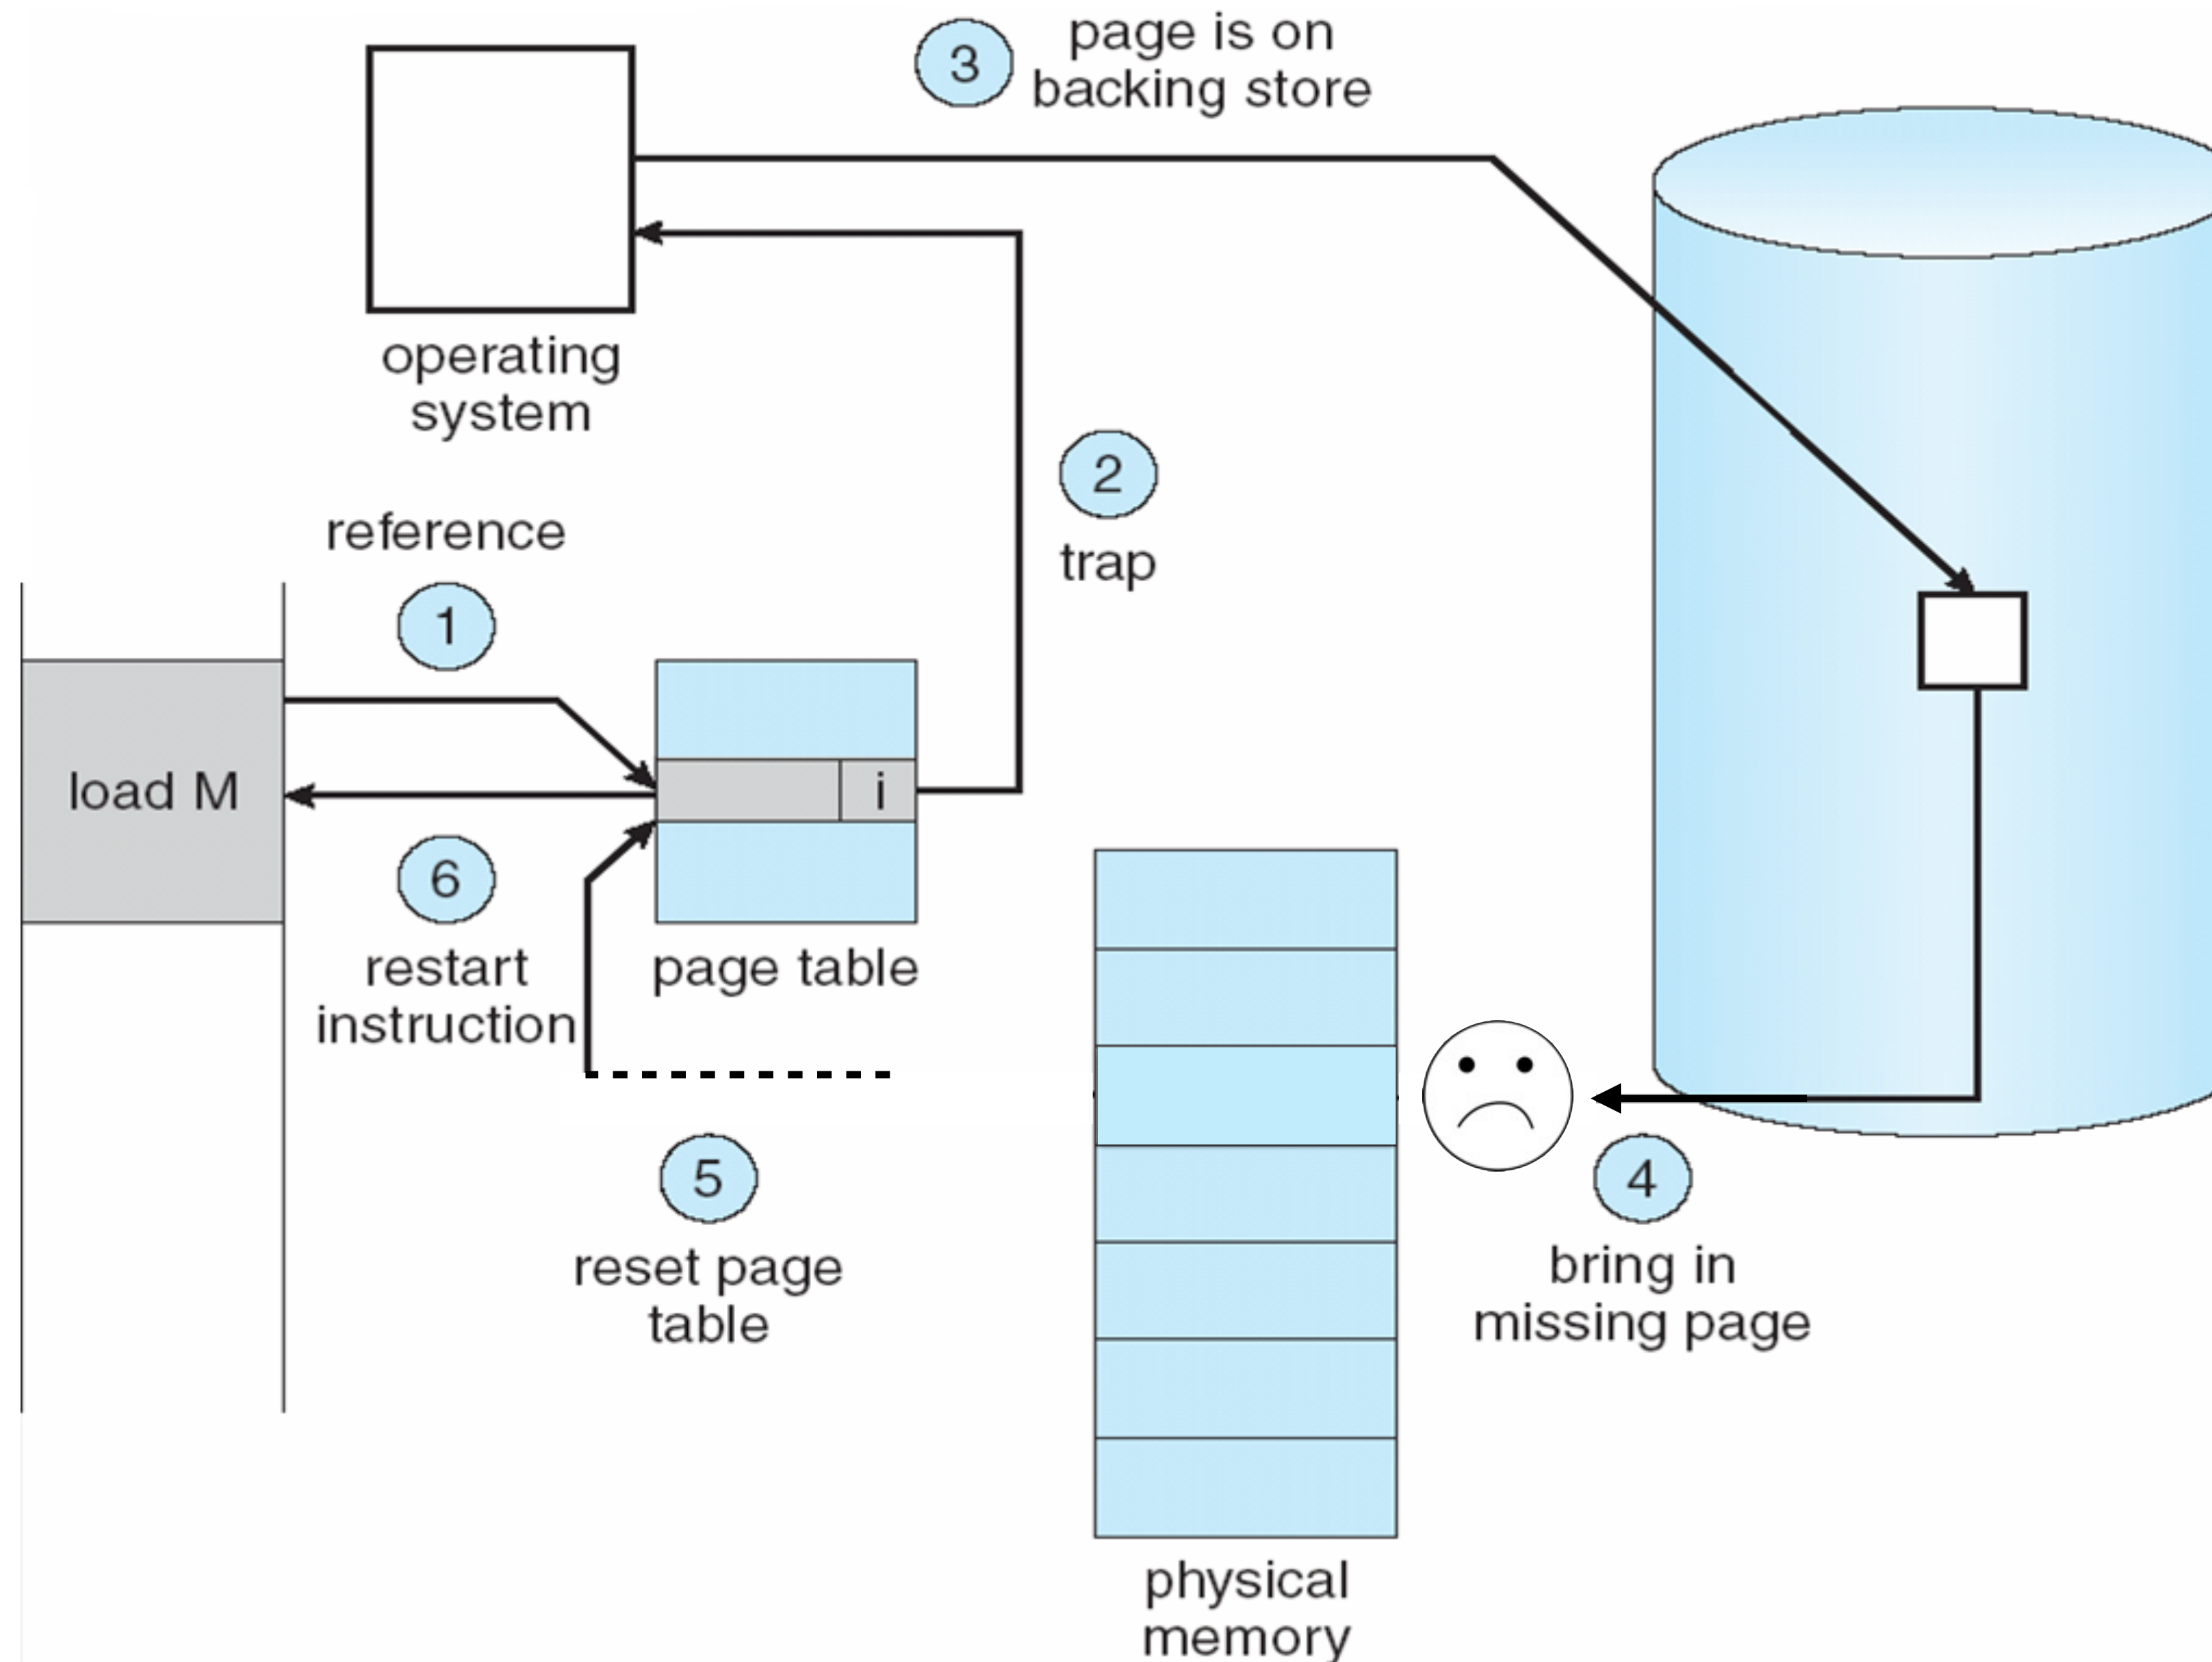
# Thrashing

- The OS wants to maximize CPU utilization. As a result, it decides that it is a good idea to increase the degree of multiprogramming by adding new processes to the system.

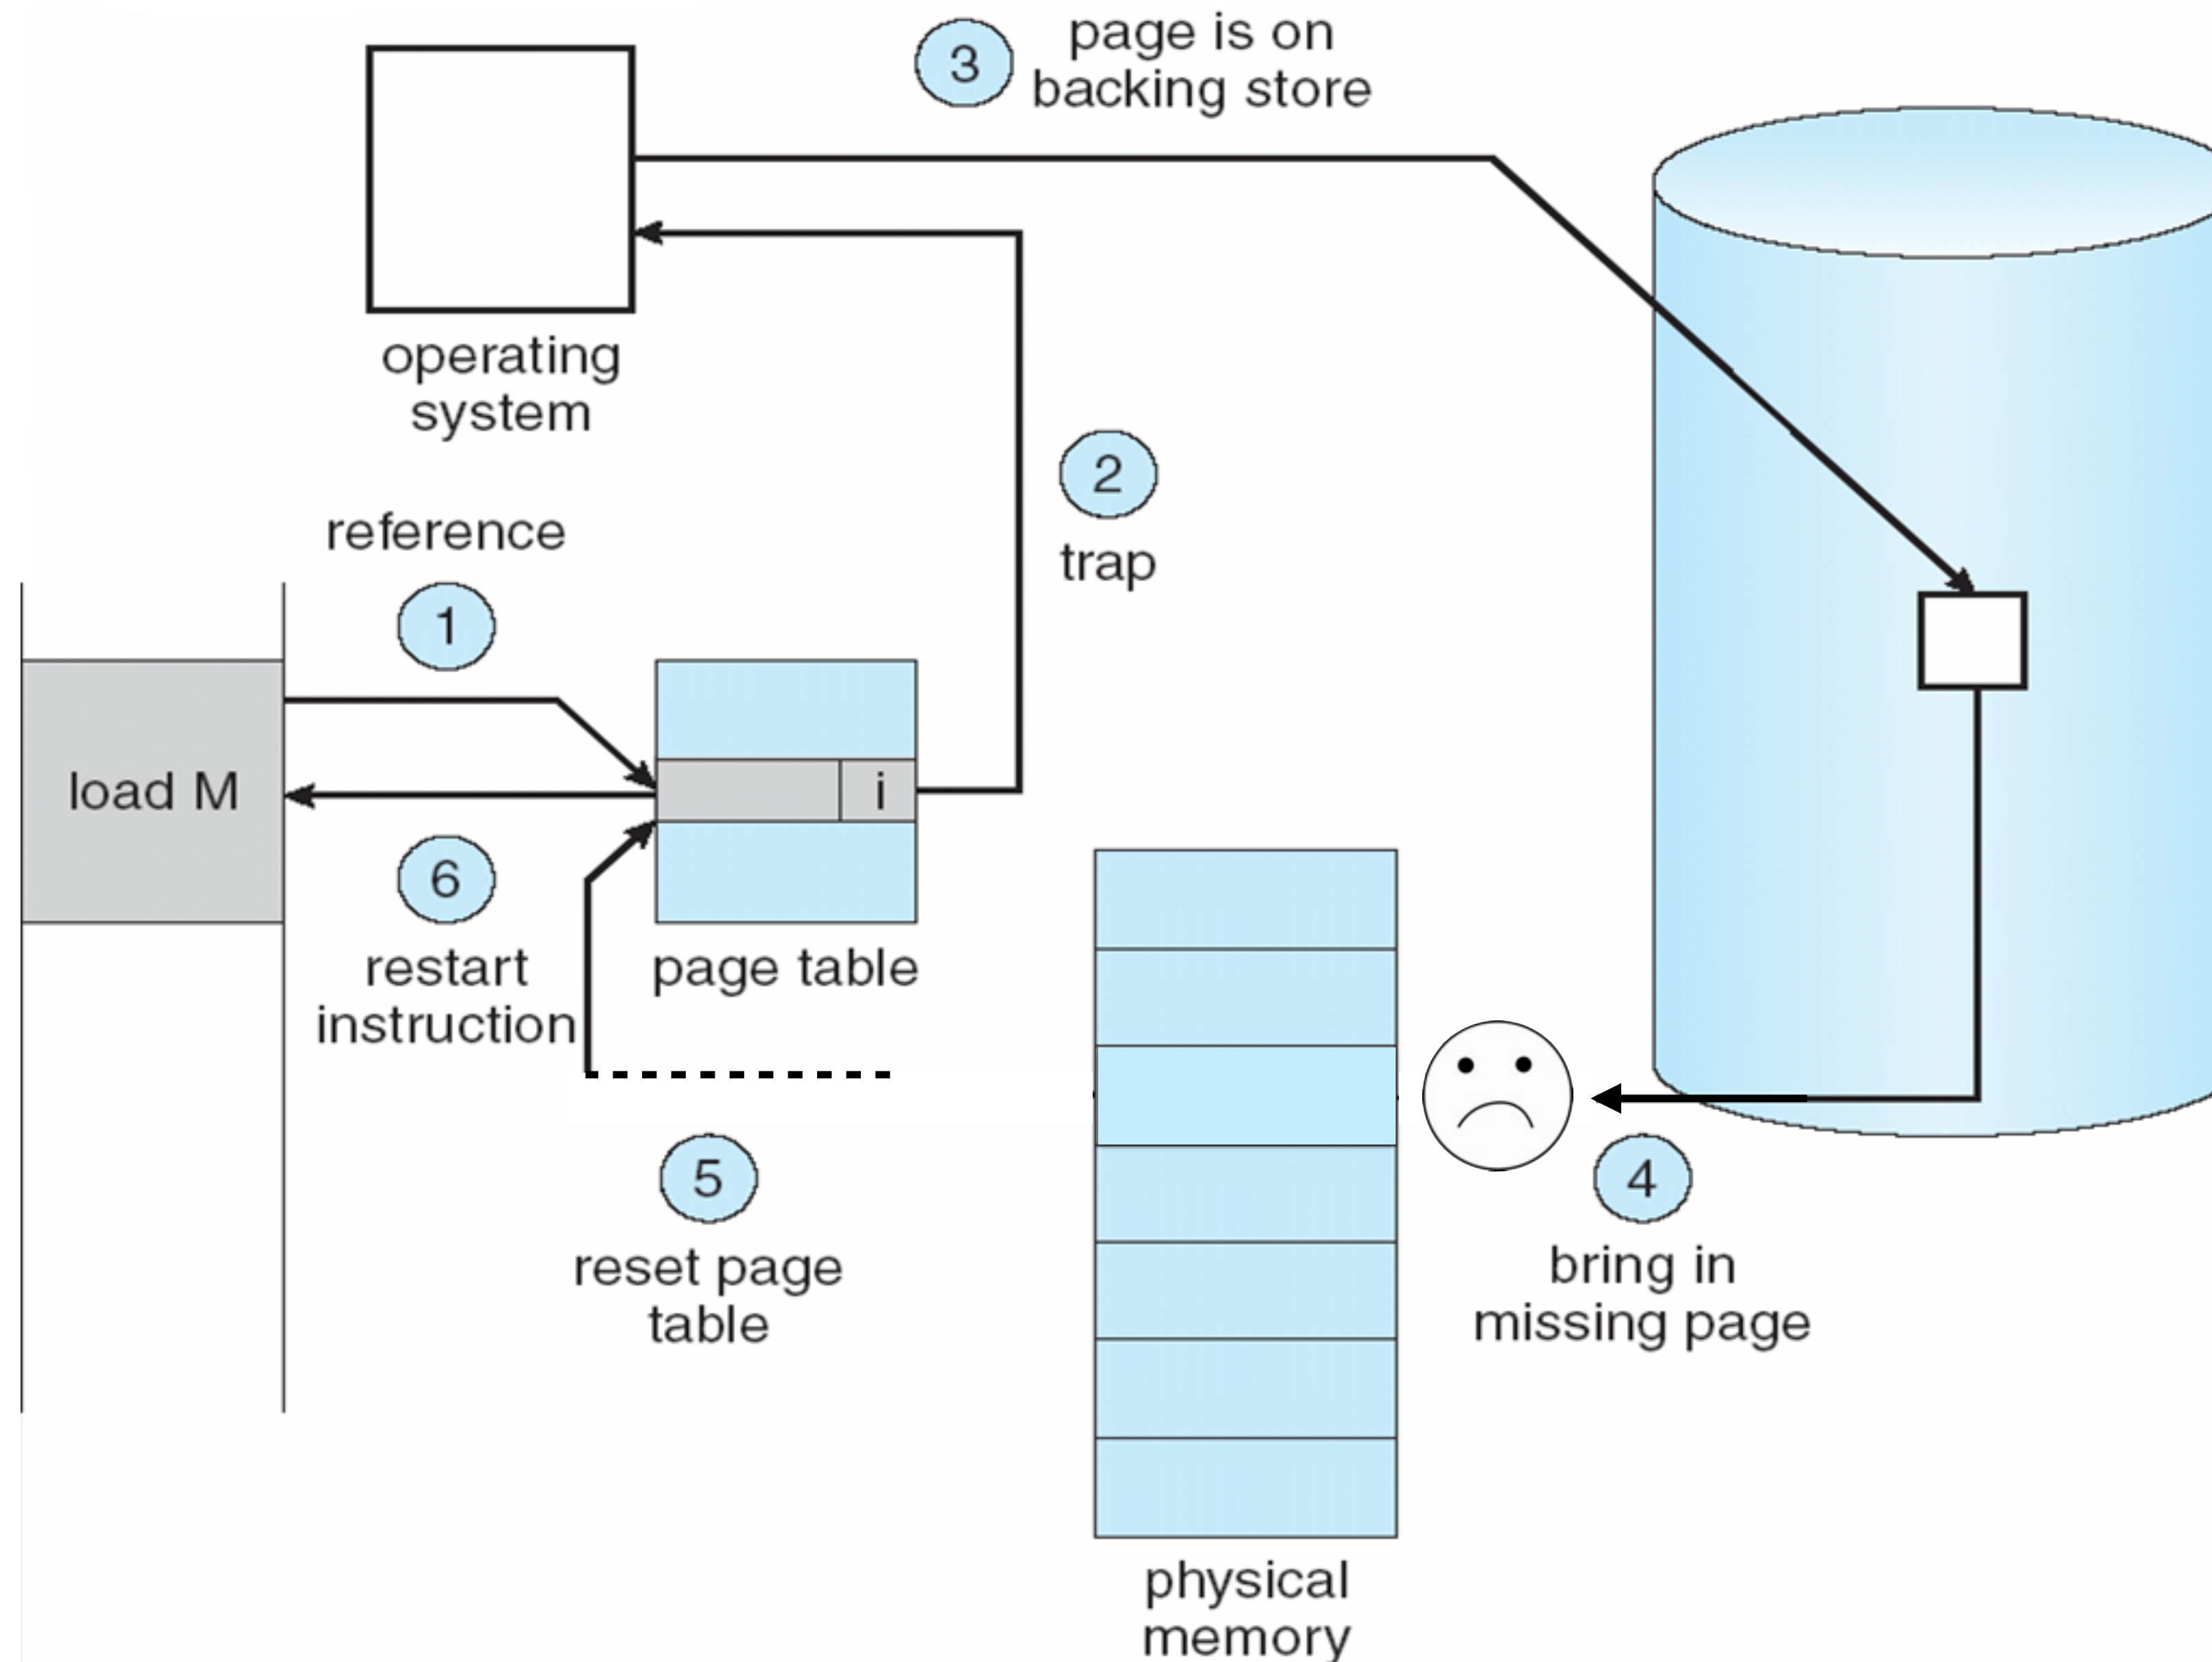






# What happens if there is no free frame?



# What happens if there is no free frame?







# Page-Replacement Algorithms


- FIFO algorithm
- Optimal page-replacement algorithm
- Least-recently used (LRU) algorithm
- Second-chance algorithm (clock)



# FIFO Algorithms

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1




page frames

# FIFO Algorithms

Access	Hit/Miss?	Evict	Resulting Cache State	
0	Miss		First-in→	0
1	Miss		First-in→	0, 1
2	Miss		First-in→	0, 1, 2
0	Hit		First-in→	0, 1, 2
1	Hit		First-in→	0, 1, 2
3	Miss	0	First-in→	1, 2, 3
0	Miss	1	First-in→	2, 3, 0
3	Hit		First-in→	2, 3, 0
1	Miss	2	First-in→	3, 0, 1
2	Miss	3	First-in→	0, 1, 2
1	Hit		First-in→	0, 1, 2

# Belady's anomaly

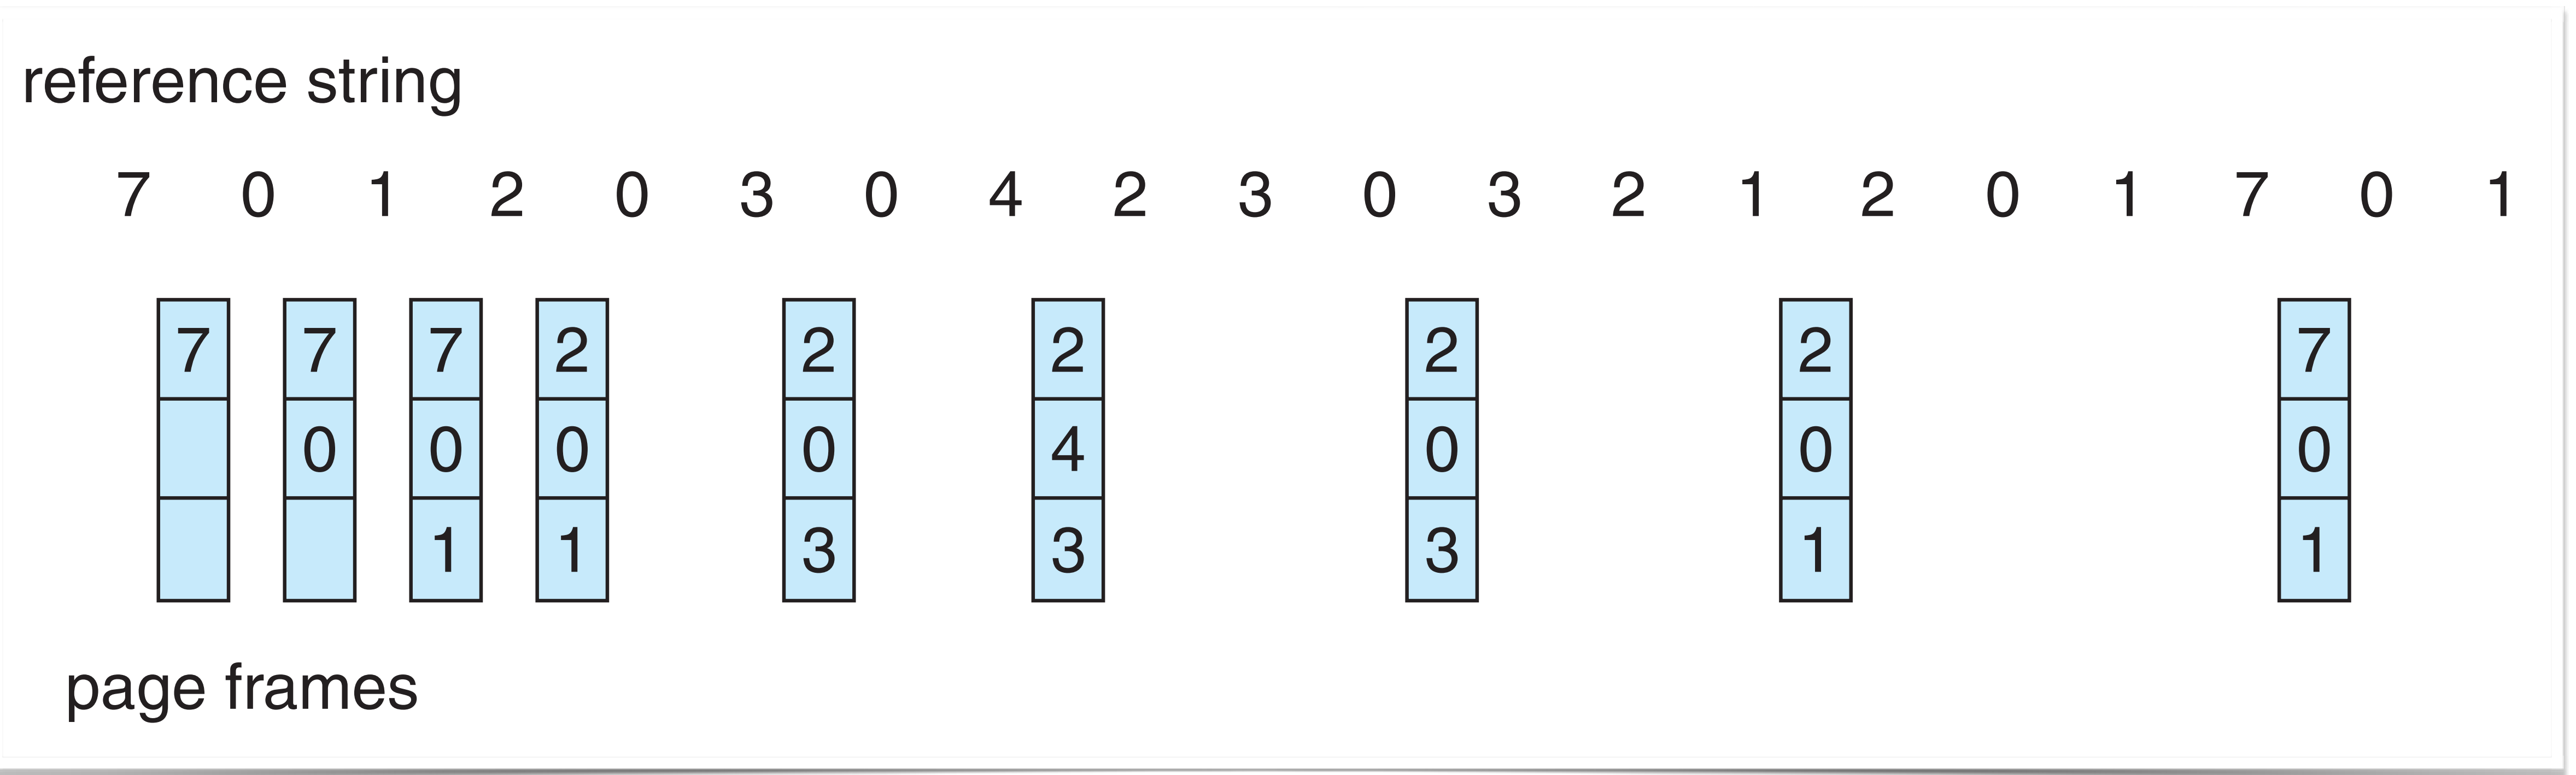


**Paper:** L. A. Belady, R. A. Nelson, G. S. Shedler, An anomaly in space-time characteristics of certain programs running in paging machine, Comm. ACM , 12, 1 (1969) 349–353.



# Optimal Algorithm

**Policy:** Replace the page that will not be used for the longest period of time.



# Optimal Algorithm


Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0, 1
2	Miss		0, 1, 2
0	Hit		0, 1, 2
1	Hit		0, 1, 2
3	Miss	2	0, 1, 3
0	Hit		0, 1, 3
3	Hit		0, 1, 3
1	Hit		0, 1, 3
2	Miss	3	0, 1, 2
1	Hit		0, 1, 2

# LRU Algorithm

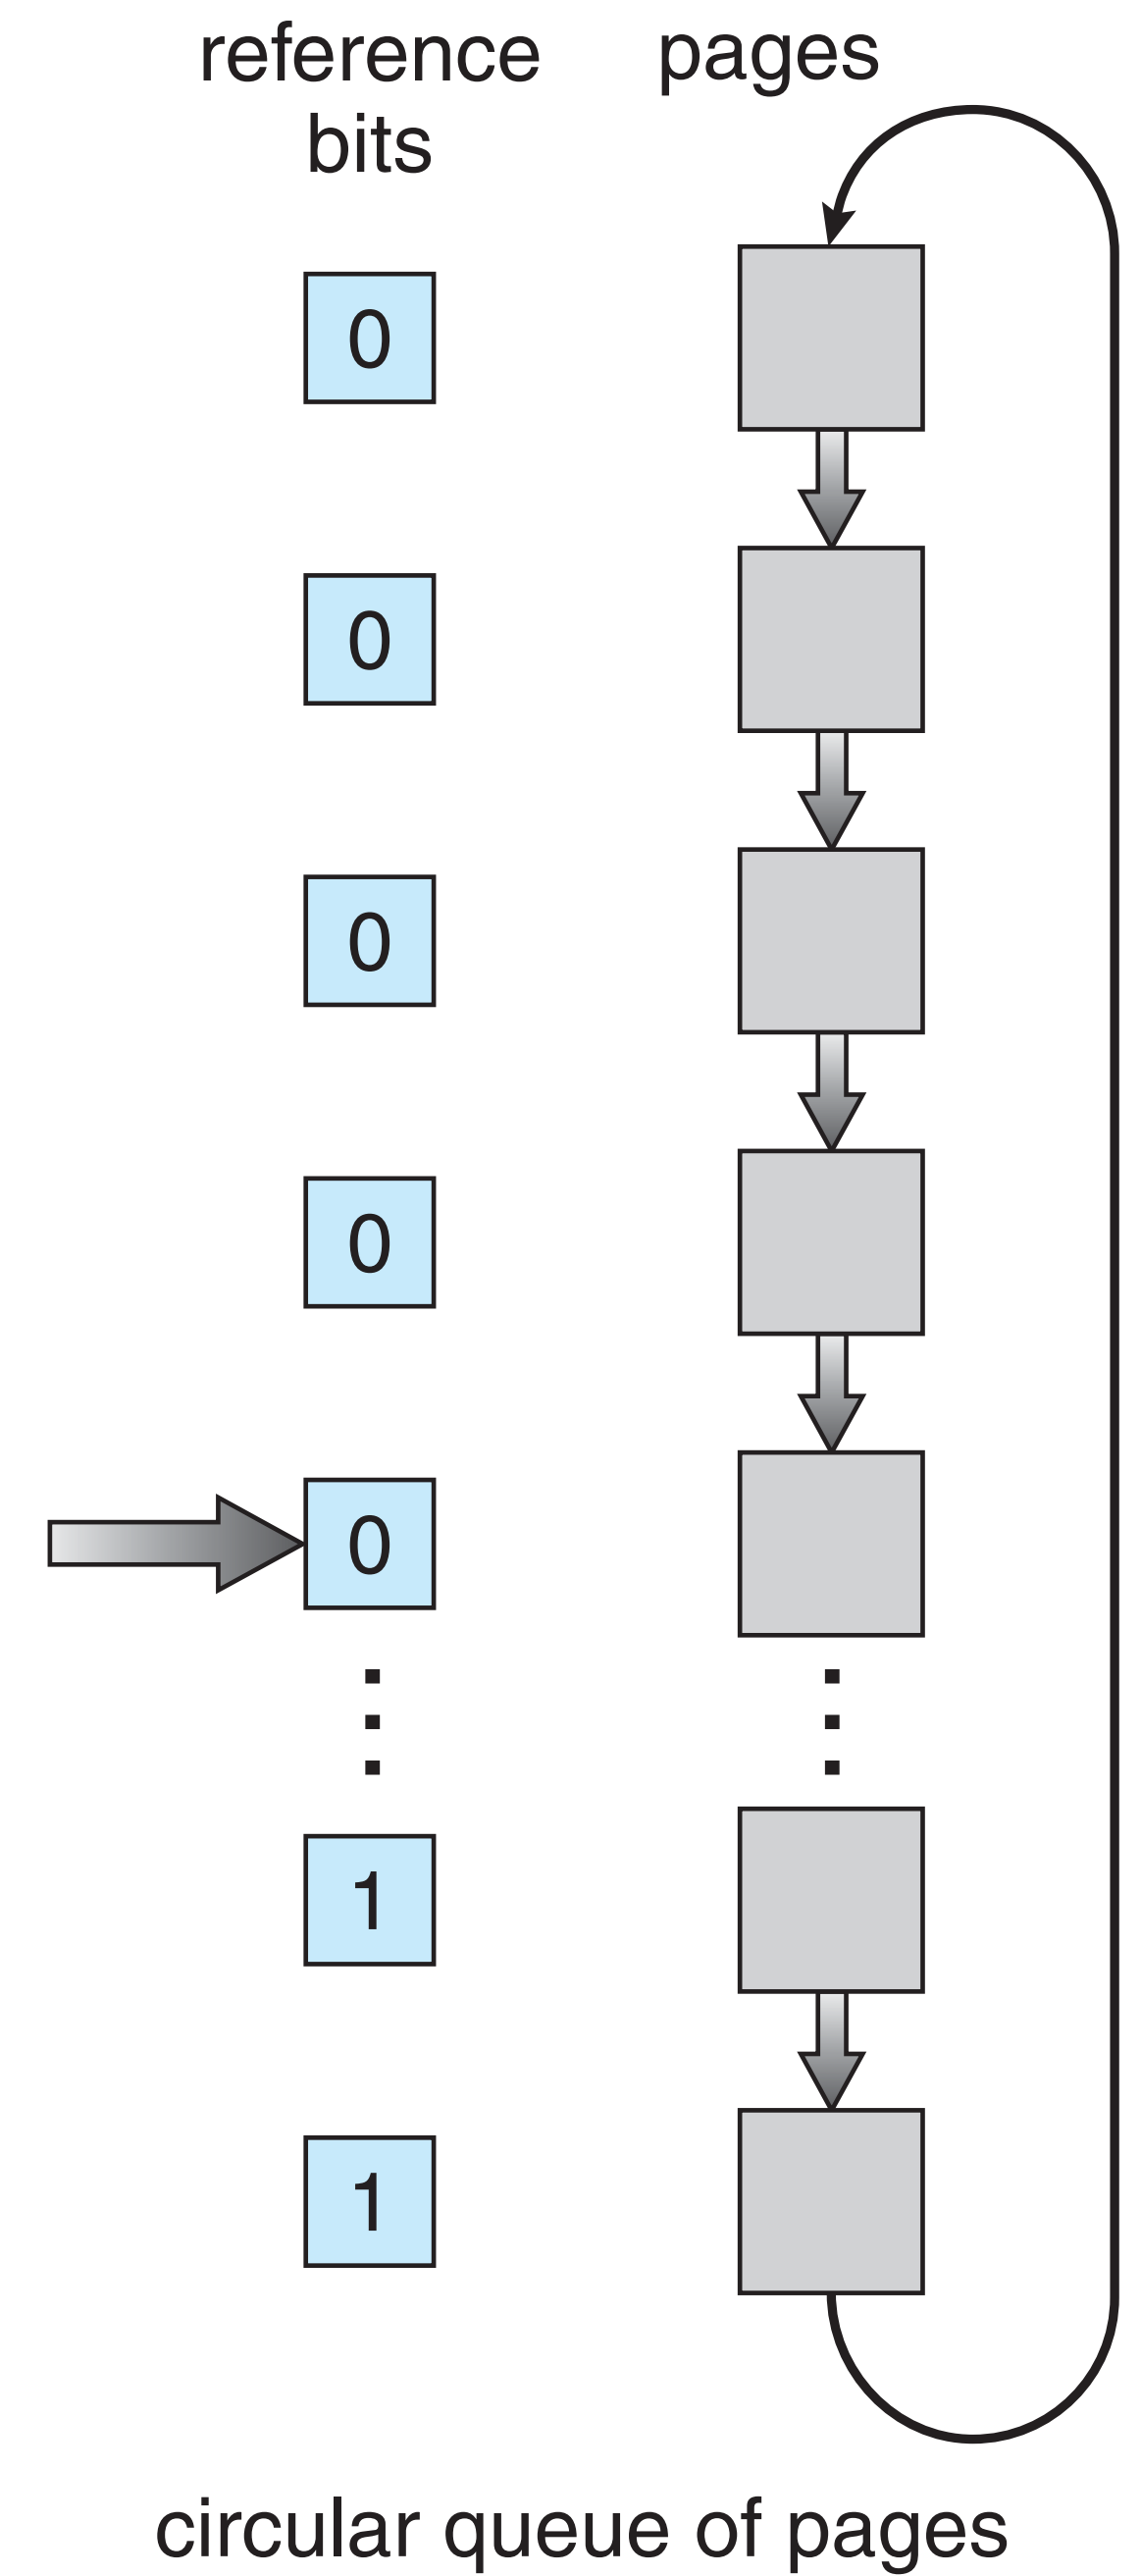
Access	Hit/Miss?	Evict	Resulting Cache State	
0	Miss		LRU→	0
1	Miss		LRU→	0, 1
2	Miss		LRU→	0, 1, 2
0	Hit		LRU→	1, 2, 0
1	Hit		LRU→	2, 0, 1
3	Miss	2	LRU→	0, 1, 3
0	Hit		LRU→	1, 3, 0
3	Hit		LRU→	1, 0, 3
1	Hit		LRU→	0, 3, 1
2	Miss	0	LRU→	3, 1, 2
1	Hit		LRU→	3, 2, 1



# Second-chance Algorithm



(a)




(b)


# Working-set Model

page reference table

. . . 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 . . .



$$WS(t_1) = \{1, 2, 5, 6, 7\}$$



$$WS(t_2) = \{3, 4\}$$

# Working Sets and Page-fault frequency

